

COMPARISON_MediaPipe+CNN+LSTM

June 21, 2025

```
[1]: from modules.SignLanguageProcessor import load_and_preprocess_data, parse_frame
import os
```

```
[2]: ROOT_PATH = ''
sequences, labels, label_map = load_and_preprocess_data(os.path.
    ↪join(ROOT_PATH, 'data'))
```

```
[3]: num_classes = len(label_map)
```

```
[4]: len(labels)
```

```
[4]: 3488
```

```
[5]: sequences.shape
```

```
[5]: (3488, 3, 61, 3)
```

```
[6]: from sklearn.model_selection import train_test_split

X_train, X_temp, y_train, y_temp = train_test_split(
    sequences, labels, test_size=0.4, stratify=labels, random_state=42
)

X_val, X_test, y_val, y_test = train_test_split(
    X_temp, y_temp, test_size=0.5, stratify=y_temp, random_state=42
)
```

```
[7]: import numpy as np
def normalize_landmark_data(X):
    """
    Normalize the landmark features (x, y) to have zero mean and unit variance_
    ↪across the training set.
    Assumes X shape is (N, F, L, T), where F=3 (x, y, vis).
    """
    X = X.copy()
    # Flatten across all samples, landmarks, and frames
    x_vals = X[:, 0, :, :].flatten()
    y_vals = X[:, 1, :, :].flatten()
```

```

    # Compute mean and std
    x_mean, x_std = np.mean(x_vals), np.std(x_vals)
    y_mean, y_std = np.mean(y_vals), np.std(y_vals)

    # Normalize
    X[:, 0, :, :] = (X[:, 0, :, :] - x_mean) / x_std
    X[:, 1, :, :] = (X[:, 1, :, :] - y_mean) / y_std

    return X, (x_mean, x_std), (y_mean, y_std)

def apply_normalization(X, x_mean, x_std, y_mean, y_std):
    X = X.copy()
    X[:, 0, :, :] = (X[:, 0, :, :] - x_mean) / x_std
    X[:, 1, :, :] = (X[:, 1, :, :] - y_mean) / y_std
    return X

```

```

[8]: def reshape_frames_for_cnn(X, y):
      X = X.transpose(0, 3, 2, 1) # (N, T, L, F)
      X = X[..., np.newaxis]      # (N, T, L, F, 1)
      return X,y

```

```

[9]: X_train_norm, (x_mean, x_std), (y_mean, y_std) = ↵
      ↪normalize_landmark_data(X_train)
      X_val_norm = apply_normalization(X_val, x_mean, x_std, y_mean, y_std)
      X_test_norm = apply_normalization(X_test, x_mean, x_std, y_mean, y_std)

      X_train_cnn, y_train_cnn = reshape_frames_for_cnn(X_train_norm, y_train)
      X_val_cnn, y_val_cnn = reshape_frames_for_cnn(X_val_norm, y_val)
      X_test_cnn, y_test_cnn = reshape_frames_for_cnn(X_test_norm, y_test)

      print(X_train_cnn.shape)
      print(y_train_cnn.shape)

```

```

(2092, 3, 61, 3, 1)
(2092,)

```

```

[10]: input_shape = X_train_cnn.shape[1:]
      print(input_shape)

```

```

(3, 61, 3, 1)

```

```

[11]: import tensorflow as tf

      train_ds = tf.data.Dataset.from_tensor_slices((X_train_cnn, y_train_cnn))
      train_ds = train_ds.shuffle(buffer_size=1000).batch(64).prefetch(tf.data.
      ↪AUTOTUNE)

```

```

val_ds = tf.data.Dataset.from_tensor_slices((X_val_cnn, y_val_cnn))
val_ds = val_ds.batch(64).prefetch(tf.data.AUTOTUNE)

test_ds = tf.data.Dataset.from_tensor_slices((X_test_cnn, y_test_cnn))
test_ds = test_ds.batch(64).prefetch(tf.data.AUTOTUNE)

```

```

[12]: from tensorflow.keras.models import Sequential
      from tensorflow.keras.layers import TimeDistributed, Conv2D, MaxPooling2D, Flatten, Input
      from tensorflow.keras.layers import LSTM, Dropout, Dense, BatchNormalization

model = Sequential([
    Input((3, 61, 3, 1)),
    TimeDistributed(Conv2D(32, (3, 2), activation='relu', padding='same')),
    TimeDistributed(BatchNormalization()),
    TimeDistributed(MaxPooling2D(pool_size=(2, 1))),
    TimeDistributed(Dropout(0.25)),

    TimeDistributed(Conv2D(64, (3, 2), activation='relu', padding='same')),
    TimeDistributed(BatchNormalization()),
    TimeDistributed(MaxPooling2D(pool_size=(2, 1))),
    TimeDistributed(Flatten()),

    LSTM(128, return_sequences=False),
    Dropout(0.5),
    Dense(num_classes, activation='softmax')
])

model.compile(optimizer='adam', loss='sparse_categorical_crossentropy',
             metrics=['accuracy'])

```

```

[13]: history = model.fit(train_ds, validation_data=val_ds, epochs=50, batch_size=64)

```

```

Epoch 1/50
33/33          6s 66ms/step -
accuracy: 0.0526 - loss: 3.3014 - val_accuracy: 0.0745 - val_loss: 3.1521
Epoch 2/50
33/33          2s 50ms/step -
accuracy: 0.0682 - loss: 3.1579 - val_accuracy: 0.0774 - val_loss: 3.0845
Epoch 3/50
33/33          2s 50ms/step -
accuracy: 0.0759 - loss: 3.0987 - val_accuracy: 0.1089 - val_loss: 3.0402
Epoch 4/50
33/33          2s 48ms/step -
accuracy: 0.0840 - loss: 3.0622 - val_accuracy: 0.0960 - val_loss: 2.9985
Epoch 5/50
33/33          2s 51ms/step -
accuracy: 0.0886 - loss: 3.0486 - val_accuracy: 0.1418 - val_loss: 2.9539

```

Epoch 6/50
 33/33 2s 61ms/step -
 accuracy: 0.0963 - loss: 3.0170 - val_accuracy: 0.0960 - val_loss: 2.9128
 Epoch 7/50
 33/33 2s 64ms/step -
 accuracy: 0.1267 - loss: 2.9517 - val_accuracy: 0.1605 - val_loss: 2.8491
 Epoch 8/50
 33/33 2s 58ms/step -
 accuracy: 0.1392 - loss: 2.8937 - val_accuracy: 0.1834 - val_loss: 2.8173
 Epoch 9/50
 33/33 2s 53ms/step -
 accuracy: 0.1743 - loss: 2.8339 - val_accuracy: 0.1848 - val_loss: 2.7654
 Epoch 10/50
 33/33 2s 58ms/step -
 accuracy: 0.2060 - loss: 2.7474 - val_accuracy: 0.2464 - val_loss: 2.6593
 Epoch 11/50
 33/33 2s 59ms/step -
 accuracy: 0.2264 - loss: 2.6703 - val_accuracy: 0.3195 - val_loss: 2.5625
 Epoch 12/50
 33/33 2s 57ms/step -
 accuracy: 0.2485 - loss: 2.5978 - val_accuracy: 0.3467 - val_loss: 2.4301
 Epoch 13/50
 33/33 2s 56ms/step -
 accuracy: 0.3020 - loss: 2.4781 - val_accuracy: 0.3496 - val_loss: 2.4377
 Epoch 14/50
 33/33 2s 49ms/step -
 accuracy: 0.3139 - loss: 2.3855 - val_accuracy: 0.3840 - val_loss: 2.3150
 Epoch 15/50
 33/33 2s 48ms/step -
 accuracy: 0.3704 - loss: 2.2743 - val_accuracy: 0.4513 - val_loss: 2.1860
 Epoch 16/50
 33/33 2s 51ms/step -
 accuracy: 0.4008 - loss: 2.1787 - val_accuracy: 0.3653 - val_loss: 2.2596
 Epoch 17/50
 33/33 2s 52ms/step -
 accuracy: 0.3991 - loss: 2.0795 - val_accuracy: 0.4513 - val_loss: 2.0517
 Epoch 18/50
 33/33 2s 50ms/step -
 accuracy: 0.4269 - loss: 2.0340 - val_accuracy: 0.4398 - val_loss: 2.0470
 Epoch 19/50
 33/33 2s 49ms/step -
 accuracy: 0.4710 - loss: 1.9426 - val_accuracy: 0.5000 - val_loss: 1.9191
 Epoch 20/50
 33/33 2s 49ms/step -
 accuracy: 0.4595 - loss: 1.8732 - val_accuracy: 0.4771 - val_loss: 1.9459
 Epoch 21/50
 33/33 2s 48ms/step -
 accuracy: 0.4952 - loss: 1.7686 - val_accuracy: 0.5244 - val_loss: 1.8199

Epoch 22/50
 33/33 2s 49ms/step -
 accuracy: 0.5095 - loss: 1.6887 - val_accuracy: 0.5086 - val_loss: 1.7860

Epoch 23/50
 33/33 2s 50ms/step -
 accuracy: 0.5202 - loss: 1.6639 - val_accuracy: 0.4871 - val_loss: 1.9429

Epoch 24/50
 33/33 2s 49ms/step -
 accuracy: 0.5427 - loss: 1.6006 - val_accuracy: 0.5029 - val_loss: 1.8247

Epoch 25/50
 33/33 2s 49ms/step -
 accuracy: 0.5362 - loss: 1.6011 - val_accuracy: 0.5688 - val_loss: 1.6381

Epoch 26/50
 33/33 2s 48ms/step -
 accuracy: 0.5516 - loss: 1.5341 - val_accuracy: 0.5458 - val_loss: 1.6210

Epoch 27/50
 33/33 2s 48ms/step -
 accuracy: 0.5697 - loss: 1.4711 - val_accuracy: 0.6418 - val_loss: 1.4740

Epoch 28/50
 33/33 2s 50ms/step -
 accuracy: 0.6006 - loss: 1.4071 - val_accuracy: 0.5989 - val_loss: 1.5208

Epoch 29/50
 33/33 2s 49ms/step -
 accuracy: 0.5919 - loss: 1.3827 - val_accuracy: 0.5287 - val_loss: 1.5904

Epoch 30/50
 33/33 2s 49ms/step -
 accuracy: 0.6250 - loss: 1.3623 - val_accuracy: 0.6218 - val_loss: 1.4242

Epoch 31/50
 33/33 2s 48ms/step -
 accuracy: 0.6179 - loss: 1.3073 - val_accuracy: 0.6490 - val_loss: 1.3192

Epoch 32/50
 33/33 2s 50ms/step -
 accuracy: 0.6405 - loss: 1.2929 - val_accuracy: 0.6289 - val_loss: 1.3961

Epoch 33/50
 33/33 2s 54ms/step -
 accuracy: 0.6491 - loss: 1.2397 - val_accuracy: 0.5487 - val_loss: 1.5479

Epoch 34/50
 33/33 2s 53ms/step -
 accuracy: 0.6258 - loss: 1.2620 - val_accuracy: 0.5888 - val_loss: 1.4296

Epoch 35/50
 33/33 2s 51ms/step -
 accuracy: 0.6500 - loss: 1.2119 - val_accuracy: 0.5745 - val_loss: 1.3688

Epoch 36/50
 33/33 2s 55ms/step -
 accuracy: 0.6546 - loss: 1.1764 - val_accuracy: 0.5673 - val_loss: 1.4578

Epoch 37/50
 33/33 2s 50ms/step -
 accuracy: 0.6684 - loss: 1.1202 - val_accuracy: 0.6189 - val_loss: 1.3412

```

Epoch 38/50
33/33          2s 53ms/step -
accuracy: 0.7068 - loss: 1.0492 - val_accuracy: 0.6232 - val_loss: 1.2689
Epoch 39/50
33/33          2s 51ms/step -
accuracy: 0.6788 - loss: 1.0312 - val_accuracy: 0.6461 - val_loss: 1.2609
Epoch 40/50
33/33          2s 53ms/step -
accuracy: 0.6981 - loss: 0.9774 - val_accuracy: 0.5989 - val_loss: 1.3219
Epoch 41/50
33/33          2s 53ms/step -
accuracy: 0.7071 - loss: 0.9732 - val_accuracy: 0.6361 - val_loss: 1.2340
Epoch 42/50
33/33          2s 53ms/step -
accuracy: 0.7175 - loss: 0.9505 - val_accuracy: 0.6275 - val_loss: 1.2091
Epoch 43/50
33/33          2s 53ms/step -
accuracy: 0.7206 - loss: 0.9661 - val_accuracy: 0.6920 - val_loss: 1.0771
Epoch 44/50
33/33          2s 51ms/step -
accuracy: 0.7210 - loss: 0.9566 - val_accuracy: 0.6304 - val_loss: 1.2521
Epoch 45/50
33/33          2s 52ms/step -
accuracy: 0.6993 - loss: 0.9491 - val_accuracy: 0.6218 - val_loss: 1.2147
Epoch 46/50
33/33          2s 55ms/step -
accuracy: 0.7266 - loss: 0.8977 - val_accuracy: 0.6633 - val_loss: 1.1671
Epoch 47/50
33/33          2s 52ms/step -
accuracy: 0.7521 - loss: 0.8862 - val_accuracy: 0.6619 - val_loss: 1.1203
Epoch 48/50
33/33          2s 51ms/step -
accuracy: 0.7667 - loss: 0.8230 - val_accuracy: 0.6246 - val_loss: 1.2619
Epoch 49/50
33/33          2s 55ms/step -
accuracy: 0.7548 - loss: 0.8559 - val_accuracy: 0.6189 - val_loss: 1.2847
Epoch 50/50
33/33          2s 53ms/step -
accuracy: 0.7470 - loss: 0.8342 - val_accuracy: 0.6633 - val_loss: 1.1711

```

```

[14]: test_loss, test_accuracy = model.evaluate(test_ds)
      print(f"Test Accuracy: {test_accuracy:.4f}")
      print(f"Test Loss: {test_loss:.4f}")

```


```

11/11          0s 9ms/step -
accuracy: 0.6181 - loss: 1.2693
Test Accuracy: 0.6304
Test Loss: 1.2593

```

```
[15]: import matplotlib.pyplot as plt
from sklearn.metrics import classification_report, confusion_matrix
import seaborn as sns
```

```
[16]: plt.figure(figsize=(12, 4))
plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')
# Plot training & validation loss values
plt.subplot(1, 2, 2)
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()
```



```
[17]: y_true, y_pred = [], []
target_names = [label_map[i] for i in range(len(label_map))]
for X_batch, y_batch in test_ds:
    y_true.append(y_batch.numpy())

    batch_pred = model.predict(X_batch, verbose=0)
    y_pred.append(np.argmax(batch_pred, axis=1))

y_true = np.concatenate(y_true)
```

```

y_pred = np.concatenate(y_pred)

print(classification_report(
    y_true, y_pred,
    digits=3,
    target_names=target_names
))

cm = confusion_matrix(y_true, y_pred, labels=range(len(label_map)))
labels = [label_map[i] for i in range(len(label_map))]

plt.figure(figsize=(10, 8))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues",
            xticklabels=labels, yticklabels=labels)
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.title("Confusion Matrix - Test Set")
plt.show()

```

	precision	recall	f1-score	support
A	0.545	0.600	0.571	20
B	0.810	0.739	0.773	23
C	0.818	0.321	0.462	28
D	0.476	0.500	0.488	20
E	0.952	0.769	0.851	26
F	0.789	0.625	0.698	24
G	0.667	0.966	0.789	29
H	0.621	0.720	0.667	25
I	0.929	0.867	0.897	30
J	0.933	0.933	0.933	30
K	0.560	0.560	0.560	25
L	1.000	0.481	0.650	27
M	0.500	0.355	0.415	31
N	0.438	0.656	0.525	32
O	0.714	0.167	0.270	30
P	1.000	0.400	0.571	25
Q	0.512	0.733	0.603	30
R	0.773	0.630	0.694	27
S	0.455	0.667	0.541	30
T	0.524	0.423	0.468	26
U	0.512	0.724	0.600	29
V	1.000	0.880	0.936	25
W	0.773	0.630	0.694	27
X	1.000	0.391	0.562	23
Y	0.247	0.741	0.370	27
Z	1.000	0.828	0.906	29

accuracy			0.630	698
macro avg	0.713	0.627	0.634	698
weighted avg	0.711	0.630	0.634	698

