

COMPARISON_CNN

June 21, 2025

```
[1]: import tensorflow as tf
from tensorflow.keras.utils import to_categorical
import os
from PIL import Image, UnidentifiedImageError
import shutil

# Configuration
IMG_SIZE = (28, 28)
BATCH_SIZE = 32
VALIDATION_SPLIT = 0.4
SEED = 42
ROOT_PATH = ''
DATASET_PATH = os.path.join(ROOT_PATH, "raw_data")
CORRUPT_PATH = os.path.join(ROOT_PATH, "corrupt_images")
os.makedirs(CORRUPT_PATH, exist_ok=True)

for root, dirs, files in os.walk(DATASET_PATH):
    for file in files:
        ext = os.path.splitext(file)[1].lower()
        if ext in [".jpg", ".jpeg", ".png", ".bmp", ".gif"]:
            path = os.path.join(root, file)
            try:
                with Image.open(path) as img:
                    img.verify() # Check integrity
            except (UnidentifiedImageError, OSError, IOError) as e:
                # Move the corrupt image
                print(f"Corrupt image found: {path} - moving to {CORRUPT_PATH}")
                dest_path = os.path.join(CORRUPT_PATH, os.path.relpath(path,
↳DATASET_PATH))
                os.makedirs(os.path.dirname(dest_path), exist_ok=True)
                shutil.move(path, dest_path)

LANDMARK_DIR = os.path.join(ROOT_PATH, "data")
RAW_IMAGE_DIR = os.path.join(ROOT_PATH, "raw_data")
FILTERED_IMAGE_DIR = os.path.join(ROOT_PATH, "filtered_raw_data")
DATASET_PATH = FILTERED_IMAGE_DIR
# Supported image extensions
```

```

IMAGE_EXTENSIONS = ['.jpg', '.jpeg', '.png', '.bmp']

# Create filtered output structure
os.makedirs(FILTERED_IMAGE_DIR, exist_ok=True)

for class_name in os.listdir(LANDMARK_DIR):
    if class_name == 'debug':
        continue
    landmark_class_dir = os.path.join(LANDMARK_DIR, class_name)
    raw_class_dir = os.path.join(RAW_IMAGE_DIR, class_name)
    filtered_class_dir = os.path.join(FILTERED_IMAGE_DIR, class_name)
    os.makedirs(filtered_class_dir, exist_ok=True)

    for file in os.listdir(landmark_class_dir):
        if not file.endswith("_landmarks.json"):
            continue

        # Get base filename without "_landmarks.json"
        base_name = file.replace("_landmarks.json", "")

        # Look for corresponding image in raw directory
        for ext in IMAGE_EXTENSIONS:
            image_file = os.path.join(raw_class_dir, base_name + ext)
            if os.path.exists(image_file):
                # Copy to filtered folder
                shutil.copy(image_file, os.path.join(filtered_class_dir, os.
↳ path.basename(image_file)))
                break

# Load training dataset with validation split
train_ds = tf.keras.preprocessing.image_dataset_from_directory(
    DATASET_PATH,
    validation_split=VALIDATION_SPLIT,
    subset="training",
    seed=SEED,
    color_mode="grayscale",
    image_size=IMG_SIZE,
    batch_size=BATCH_SIZE
)
num_classes = len(train_ds.class_names)
label_map = train_ds.class_names

val_ds = tf.keras.preprocessing.image_dataset_from_directory(
    DATASET_PATH,
    validation_split=VALIDATION_SPLIT,
    subset="validation",
    seed=SEED,

```

```

        color_mode="grayscale",
        image_size=IMG_SIZE,
        batch_size=BATCH_SIZE
    )

test_ds = val_ds.shard(2,0)
val_ds = val_ds.shard(2,1)
# Normalize pixel values to [0, 1]
normalization_layer = tf.keras.layers.Rescaling(1./255)
train_ds = train_ds.map(lambda x, y: (normalization_layer(x), y))
val_ds = val_ds.map(lambda x, y: (normalization_layer(x), y))
test_ds = test_ds.map(lambda x, y: (normalization_layer(x), y))
# Cache and prefetch for performance
AUTOTUNE = tf.data.AUTOTUNE
train_ds = train_ds.cache().shuffle(1000).prefetch(buffer_size=AUTOTUNE)
val_ds = val_ds.cache().prefetch(buffer_size=AUTOTUNE)
test_ds = test_ds.cache().prefetch(buffer_size=AUTOTUNE)

```

Found 1691 files belonging to 26 classes.
Using 1015 files for training.
Using 1015 files for training.
Found 1691 files belonging to 26 classes.
Using 676 files for validation.

```

[2]: from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Dropout, BatchNormalization, Input
from tensorflow.keras.layers import Flatten, Dense, GlobalAveragePooling2D
from tensorflow.keras.optimizers import Adam

model = Sequential([
    Input((28, 28, 1)),
    Conv2D(16, (3, 3), activation='relu'),
    BatchNormalization(),
    MaxPooling2D(pool_size=(2, 2)),
    Dropout(0.1),

    Conv2D(32, (3, 3), activation='relu'),
    BatchNormalization(),
    MaxPooling2D(pool_size=(2, 2)),
    Dropout(0.2),

    GlobalAveragePooling2D(),
    Flatten(),

    Dense(128, activation='relu'),
    Dropout(0.2),

```

```

        Dense(num_classes, activation='softmax')
    ])

model.compile(optimizer=Adam(1e-3),
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

```

```
[3]: history = model.fit(train_ds, validation_data=val_ds, epochs=50)
```

```

Epoch 1/50
32/32          10s 110ms/step -
accuracy: 0.0646 - loss: 3.3258 - val_accuracy: 0.0586 - val_loss: 3.2438
Epoch 2/50
32/32          0s 8ms/step -
accuracy: 0.0779 - loss: 3.1777 - val_accuracy: 0.0432 - val_loss: 3.2305
Epoch 3/50
32/32          0s 7ms/step -
accuracy: 0.0779 - loss: 3.1342 - val_accuracy: 0.0432 - val_loss: 3.2131
Epoch 4/50
32/32          0s 7ms/step -
accuracy: 0.0853 - loss: 3.0912 - val_accuracy: 0.0432 - val_loss: 3.2172
Epoch 5/50
32/32          0s 8ms/step -
accuracy: 0.1127 - loss: 3.0734 - val_accuracy: 0.0432 - val_loss: 3.2369
Epoch 6/50
32/32          0s 7ms/step -
accuracy: 0.1025 - loss: 3.0758 - val_accuracy: 0.0432 - val_loss: 3.2703
Epoch 7/50
32/32          0s 7ms/step -
accuracy: 0.0977 - loss: 3.0274 - val_accuracy: 0.0463 - val_loss: 3.2901
Epoch 8/50
32/32          0s 8ms/step -
accuracy: 0.1109 - loss: 3.0025 - val_accuracy: 0.0648 - val_loss: 3.3376
Epoch 9/50
32/32          0s 9ms/step -
accuracy: 0.1070 - loss: 3.0063 - val_accuracy: 0.0463 - val_loss: 3.3455
Epoch 10/50
32/32          0s 9ms/step -
accuracy: 0.1312 - loss: 2.9960 - val_accuracy: 0.0463 - val_loss: 3.3476
Epoch 11/50
32/32          0s 9ms/step -
accuracy: 0.1305 - loss: 2.9760 - val_accuracy: 0.0556 - val_loss: 3.3604
Epoch 12/50
32/32          0s 8ms/step -
accuracy: 0.1403 - loss: 2.8979 - val_accuracy: 0.0432 - val_loss: 3.4240
Epoch 13/50
32/32          0s 8ms/step -
accuracy: 0.1282 - loss: 2.9230 - val_accuracy: 0.0525 - val_loss: 3.2952

```

Epoch 14/50
 32/32 0s 9ms/step -
 accuracy: 0.1391 - loss: 2.9310 - val_accuracy: 0.0556 - val_loss: 3.2832
 Epoch 15/50
 32/32 0s 9ms/step -
 accuracy: 0.1431 - loss: 2.8571 - val_accuracy: 0.0586 - val_loss: 3.3570
 Epoch 16/50
 32/32 0s 8ms/step -
 accuracy: 0.1266 - loss: 2.8947 - val_accuracy: 0.0432 - val_loss: 3.3887
 Epoch 17/50
 32/32 0s 8ms/step -
 accuracy: 0.1579 - loss: 2.8416 - val_accuracy: 0.0648 - val_loss: 3.3575
 Epoch 18/50
 32/32 0s 8ms/step -
 accuracy: 0.1555 - loss: 2.8437 - val_accuracy: 0.0802 - val_loss: 3.2834
 Epoch 19/50
 32/32 0s 9ms/step -
 accuracy: 0.1573 - loss: 2.8043 - val_accuracy: 0.0833 - val_loss: 3.3043
 Epoch 20/50
 32/32 0s 8ms/step -
 accuracy: 0.1690 - loss: 2.7500 - val_accuracy: 0.0772 - val_loss: 3.2085
 Epoch 21/50
 32/32 0s 8ms/step -
 accuracy: 0.1864 - loss: 2.7843 - val_accuracy: 0.0679 - val_loss: 3.1802
 Epoch 22/50
 32/32 0s 8ms/step -
 accuracy: 0.1929 - loss: 2.7064 - val_accuracy: 0.0679 - val_loss: 3.1480
 Epoch 23/50
 32/32 0s 9ms/step -
 accuracy: 0.1874 - loss: 2.7182 - val_accuracy: 0.0802 - val_loss: 3.1285
 Epoch 24/50
 32/32 0s 8ms/step -
 accuracy: 0.1820 - loss: 2.6816 - val_accuracy: 0.0833 - val_loss: 3.2085
 Epoch 25/50
 32/32 0s 8ms/step -
 accuracy: 0.2284 - loss: 2.6256 - val_accuracy: 0.1173 - val_loss: 3.1260
 Epoch 26/50
 32/32 0s 8ms/step -
 accuracy: 0.1869 - loss: 2.6635 - val_accuracy: 0.1049 - val_loss: 3.1225
 Epoch 27/50
 32/32 0s 7ms/step -
 accuracy: 0.2261 - loss: 2.5858 - val_accuracy: 0.0833 - val_loss: 3.1284
 Epoch 28/50
 32/32 0s 7ms/step -
 accuracy: 0.2214 - loss: 2.6058 - val_accuracy: 0.0741 - val_loss: 3.4201
 Epoch 29/50
 32/32 0s 7ms/step -
 accuracy: 0.2314 - loss: 2.5598 - val_accuracy: 0.1049 - val_loss: 3.1550

Epoch 30/50
 32/32 0s 7ms/step -
 accuracy: 0.2445 - loss: 2.5720 - val_accuracy: 0.1019 - val_loss: 3.2403
 Epoch 31/50
 32/32 0s 7ms/step -
 accuracy: 0.2229 - loss: 2.5674 - val_accuracy: 0.1019 - val_loss: 3.1437
 Epoch 32/50
 32/32 0s 7ms/step -
 accuracy: 0.2751 - loss: 2.5286 - val_accuracy: 0.1235 - val_loss: 3.1262
 Epoch 33/50
 32/32 0s 7ms/step -
 accuracy: 0.2657 - loss: 2.4702 - val_accuracy: 0.1111 - val_loss: 3.2315
 Epoch 34/50
 32/32 0s 7ms/step -
 accuracy: 0.2206 - loss: 2.5265 - val_accuracy: 0.0895 - val_loss: 3.1556
 Epoch 35/50
 32/32 0s 7ms/step -
 accuracy: 0.2655 - loss: 2.4060 - val_accuracy: 0.1481 - val_loss: 3.1818
 Epoch 36/50
 32/32 0s 7ms/step -
 accuracy: 0.2655 - loss: 2.4504 - val_accuracy: 0.1204 - val_loss: 3.0639
 Epoch 37/50
 32/32 0s 7ms/step -
 accuracy: 0.2426 - loss: 2.4265 - val_accuracy: 0.0957 - val_loss: 3.4497
 Epoch 38/50
 32/32 0s 7ms/step -
 accuracy: 0.2816 - loss: 2.3985 - val_accuracy: 0.1235 - val_loss: 3.0809
 Epoch 39/50
 32/32 0s 7ms/step -
 accuracy: 0.2877 - loss: 2.3591 - val_accuracy: 0.1235 - val_loss: 3.1293
 Epoch 40/50
 32/32 0s 7ms/step -
 accuracy: 0.2752 - loss: 2.3650 - val_accuracy: 0.0895 - val_loss: 3.2817
 Epoch 41/50
 32/32 0s 7ms/step -
 accuracy: 0.2780 - loss: 2.3536 - val_accuracy: 0.1265 - val_loss: 3.1639
 Epoch 42/50
 32/32 0s 7ms/step -
 accuracy: 0.2833 - loss: 2.3562 - val_accuracy: 0.1204 - val_loss: 3.1205
 Epoch 43/50
 32/32 0s 7ms/step -
 accuracy: 0.3205 - loss: 2.3018 - val_accuracy: 0.1327 - val_loss: 3.1374
 Epoch 44/50
 32/32 0s 8ms/step -
 accuracy: 0.3468 - loss: 2.2001 - val_accuracy: 0.1265 - val_loss: 3.2206
 Epoch 45/50
 32/32 0s 7ms/step -
 accuracy: 0.3397 - loss: 2.2323 - val_accuracy: 0.1049 - val_loss: 3.2535

```

Epoch 46/50
32/32          0s 8ms/step -
accuracy: 0.3421 - loss: 2.2213 - val_accuracy: 0.1204 - val_loss: 3.2052
Epoch 47/50
32/32          0s 7ms/step -
accuracy: 0.3281 - loss: 2.2188 - val_accuracy: 0.1296 - val_loss: 3.1699
Epoch 48/50
32/32          0s 8ms/step -
accuracy: 0.3459 - loss: 2.1450 - val_accuracy: 0.0926 - val_loss: 3.3231
Epoch 49/50
32/32          0s 8ms/step -
accuracy: 0.3374 - loss: 2.1833 - val_accuracy: 0.1574 - val_loss: 3.1307
Epoch 50/50
32/32          0s 8ms/step -
accuracy: 0.3332 - loss: 2.1773 - val_accuracy: 0.1512 - val_loss: 3.2231

```

```

[4]: test_loss, test_accuracy = model.evaluate(test_ds)
      print(f"Test Accuracy: {test_accuracy:.4f}")
      print(f"Test Loss: {test_loss:.4f}")

```

```

11/11          3s 269ms/step -
accuracy: 0.1087 - loss: 3.3143
Test Accuracy: 0.1108
Test Loss: 3.3241

```

```


[5]: import matplotlib.pyplot as plt
      from sklearn.metrics import classification_report, confusion_matrix
      import seaborn as sns
      import numpy as np

```

```

[6]: plt.figure(figsize=(12, 4))
      plt.subplot(1, 2, 1)
      plt.plot(history.history['accuracy'])
      plt.plot(history.history['val_accuracy'])
      plt.title('Model accuracy')
      plt.ylabel('Accuracy')
      plt.xlabel('Epoch')
      plt.legend(['Train', 'Validation'], loc='upper left')
      # Plot training & validation loss values
      plt.subplot(1, 2, 2)
      plt.plot(history.history['loss'])
      plt.plot(history.history['val_loss'])
      plt.title('Model loss')
      plt.ylabel('Loss')
      plt.xlabel('Epoch')
      plt.legend(['Train', 'Validation'], loc='upper left')
      plt.show()

```



```
[7]: y_true, y_pred = [], []
target_names = [label_map[i] for i in range(len(label_map))]
for X_batch, y_batch in test_ds:
    y_true.append(y_batch.numpy())

    batch_pred = model.predict(X_batch, verbose=0)
    y_pred.append(np.argmax(batch_pred, axis=1))

y_true = np.concatenate(y_true)
y_pred = np.concatenate(y_pred)

print(classification_report(
    y_true, y_pred,
    digits=3,
    target_names=target_names
))

cm = confusion_matrix(y_true, y_pred, labels=range(len(label_map)))
labels = [label_map[i] for i in range(len(label_map))]

plt.figure(figsize=(10, 8))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues",
            xticklabels=labels, yticklabels=labels)
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.title("Confusion Matrix - Test Set")
plt.show()
```

c:\Users\chris\.conda\envs\ASLR\Lib\site-packages\sklearn\metrics_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.


```
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
c:\Users\chris\.conda\envs\ASLR\Lib\site-
packages\sklearn\metrics\_classification.py:1565: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
c:\Users\chris\.conda\envs\ASLR\Lib\site-
packages\sklearn\metrics\_classification.py:1565: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

	precision	recall	f1-score	support			
A	0.000	0.000	0.000	12			
B	0.000	0.000	0.000	9			
C	0.065	0.125	0.085	16			
D	0.000	0.000	0.000	9			
E	0.116	0.417	0.182	12			
F	0.000	0.000	0.000	5			
G	0.333	0.125	0.182	8			
H	0.000	0.000	0.000	9			
I	0.000	0.000	0.000	20			
J	0.167	0.105	0.129	19			
K	0.000	0.000	0.000	12			
L	0.000	0.000	0.000	23			
M	0.333	0.250	0.286	8			
N	0.000	0.000	0.000	8			
O	0.028	0.059	0.038	17			
P	0.000	0.000	0.000	12			
Q	1.000	0.077	0.143	13			
R	0.083	0.050	0.062	20			
S	0.222	0.250	0.235	8			
T	0.136	0.143	0.140	21			
U	0.043	0.062	0.051	16			
V	0.106	0.312	0.159	16			
W	0.115	0.150	0.130	20			
X	0.000	0.000	0.000	8			
Y	0.000	0.000	0.000	12			
Z	0.400	0.526	0.455	19			
accuracy				0.111	352		
macro avg				0.121	0.102	0.088	352
weighted avg				0.122	0.111	0.093	352

