

COMPARISON_MediaPipe+CNN+LSTM

June 21, 2025

```
[1]: from modules.SignLanguageProcessor import load_and_preprocess_data, parse_frame
import os
```

```
[2]: ROOT_PATH = ''
sequences, labels, label_map = load_and_preprocess_data(os.path.
    ↪join(ROOT_PATH, 'data'))
```

```
[3]: num_classes = len(label_map)
```

```
[4]: len(labels)
```

```
[4]: 1691
```

```
[5]: sequences.shape
```

```
[5]: (1691, 3, 61, 3)
```

```
[6]: from sklearn.model_selection import train_test_split

X_train, X_temp, y_train, y_temp = train_test_split(
    sequences, labels, test_size=0.4, stratify=labels, random_state=42
)

X_val, X_test, y_val, y_test = train_test_split(
    X_temp, y_temp, test_size=0.5, stratify=y_temp, random_state=42
)
```

```
[7]: import numpy as np
def normalize_landmark_data(X):
    """
    Normalize the landmark features (x, y) to have zero mean and unit variance_
    ↪across the training set.
    Assumes X shape is (N, F, L, T), where F=3 (x, y, vis).
    """
    X = X.copy()
    # Flatten across all samples, landmarks, and frames
    x_vals = X[:, 0, :, :].flatten()
    y_vals = X[:, 1, :, :].flatten()
```

```

    # Compute mean and std
    x_mean, x_std = np.mean(x_vals), np.std(x_vals)
    y_mean, y_std = np.mean(y_vals), np.std(y_vals)

    # Normalize
    X[:, 0, :, :] = (X[:, 0, :, :] - x_mean) / x_std
    X[:, 1, :, :] = (X[:, 1, :, :] - y_mean) / y_std

    return X, (x_mean, x_std), (y_mean, y_std)

def apply_normalization(X, x_mean, x_std, y_mean, y_std):
    X = X.copy()
    X[:, 0, :, :] = (X[:, 0, :, :] - x_mean) / x_std
    X[:, 1, :, :] = (X[:, 1, :, :] - y_mean) / y_std
    return X

```

```

[8]: def reshape_frames_for_cnn(X, y):
      X = X.transpose(0, 3, 2, 1) # (N, T, L, F)
      X = X[..., np.newaxis]      # (N, T, L, F, 1)
      return X,y

```

```

[9]: X_train_norm, (x_mean, x_std), (y_mean, y_std) = ↵
      ↪normalize_landmark_data(X_train)
      X_val_norm = apply_normalization(X_val, x_mean, x_std, y_mean, y_std)
      X_test_norm = apply_normalization(X_test, x_mean, x_std, y_mean, y_std)

      X_train_cnn, y_train_cnn = reshape_frames_for_cnn(X_train_norm, y_train)
      X_val_cnn, y_val_cnn = reshape_frames_for_cnn(X_val_norm, y_val)
      X_test_cnn, y_test_cnn = reshape_frames_for_cnn(X_test_norm, y_test)

      print(X_train_cnn.shape)
      print(y_train_cnn.shape)

```

```

(1014, 3, 61, 3, 1)
(1014,)

```

```

[10]: input_shape = X_train_cnn.shape[1:]
      print(input_shape)

```

```

(3, 61, 3, 1)

```

```

[11]: import tensorflow as tf

      train_ds = tf.data.Dataset.from_tensor_slices((X_train_cnn, y_train_cnn))
      train_ds = train_ds.shuffle(buffer_size=1000).batch(64).prefetch(tf.data.
      ↪AUTOTUNE)

```

```

val_ds = tf.data.Dataset.from_tensor_slices((X_val_cnn, y_val_cnn))
val_ds = val_ds.batch(64).prefetch(tf.data.AUTOTUNE)

test_ds = tf.data.Dataset.from_tensor_slices((X_test_cnn, y_test_cnn))
test_ds = test_ds.batch(64).prefetch(tf.data.AUTOTUNE)

```

```

[12]: from tensorflow.keras.models import Sequential
      from tensorflow.keras.layers import TimeDistributed, Conv2D, MaxPooling2D, Flatten, Input
      from tensorflow.keras.layers import LSTM, Dropout, Dense, BatchNormalization

model = Sequential([
    Input((3, 61, 3, 1)),
    TimeDistributed(Conv2D(32, (3, 2), activation='relu', padding='same')),
    TimeDistributed(BatchNormalization()),
    TimeDistributed(MaxPooling2D(pool_size=(2, 1))),
    TimeDistributed(Dropout(0.25)),

    TimeDistributed(Conv2D(64, (3, 2), activation='relu', padding='same')),
    TimeDistributed(BatchNormalization()),
    TimeDistributed(MaxPooling2D(pool_size=(2, 1))),
    TimeDistributed(Flatten()),

    LSTM(128, return_sequences=False),
    Dropout(0.5),
    Dense(num_classes, activation='softmax')
])

model.compile(optimizer='adam', loss='sparse_categorical_crossentropy',
             metrics=['accuracy'])

```

```

[13]: history = model.fit(train_ds, validation_data=val_ds, epochs=50, batch_size=64)

```

```

Epoch 1/50
16/16          6s 88ms/step -
accuracy: 0.0745 - loss: 3.1710 - val_accuracy: 0.1036 - val_loss: 3.0451
Epoch 2/50
16/16          1s 53ms/step -
accuracy: 0.1071 - loss: 2.8903 - val_accuracy: 0.1006 - val_loss: 2.9353
Epoch 3/50
16/16          1s 53ms/step -
accuracy: 0.0954 - loss: 2.8460 - val_accuracy: 0.1361 - val_loss: 2.8889
Epoch 4/50
16/16          1s 52ms/step -
accuracy: 0.1195 - loss: 2.7836 - val_accuracy: 0.1036 - val_loss: 2.8704
Epoch 5/50
16/16          1s 50ms/step -
accuracy: 0.1575 - loss: 2.7642 - val_accuracy: 0.1065 - val_loss: 2.8336

```

Epoch 6/50
16/16 1s 58ms/step -
accuracy: 0.1530 - loss: 2.7059 - val_accuracy: 0.0858 - val_loss: 2.7923

Epoch 7/50
16/16 1s 55ms/step -
accuracy: 0.1703 - loss: 2.6484 - val_accuracy: 0.1124 - val_loss: 2.7238

Epoch 8/50
16/16 1s 52ms/step -
accuracy: 0.2210 - loss: 2.5734 - val_accuracy: 0.1272 - val_loss: 2.6760

Epoch 9/50
16/16 1s 50ms/step -
accuracy: 0.2127 - loss: 2.5619 - val_accuracy: 0.1243 - val_loss: 2.6858

Epoch 10/50
16/16 1s 50ms/step -
accuracy: 0.2772 - loss: 2.4206 - val_accuracy: 0.2426 - val_loss: 2.5392

Epoch 11/50
16/16 1s 51ms/step -
accuracy: 0.3120 - loss: 2.3746 - val_accuracy: 0.3047 - val_loss: 2.4440

Epoch 12/50
16/16 1s 50ms/step -
accuracy: 0.3578 - loss: 2.2967 - val_accuracy: 0.3343 - val_loss: 2.3281

Epoch 13/50
16/16 1s 51ms/step -
accuracy: 0.3902 - loss: 2.1895 - val_accuracy: 0.4852 - val_loss: 2.1701

Epoch 14/50
16/16 1s 48ms/step -
accuracy: 0.4225 - loss: 2.0491 - val_accuracy: 0.4941 - val_loss: 2.0390

Epoch 15/50
16/16 1s 47ms/step -
accuracy: 0.4730 - loss: 1.9566 - val_accuracy: 0.5296 - val_loss: 1.9408

Epoch 16/50
16/16 1s 47ms/step -
accuracy: 0.5299 - loss: 1.8429 - val_accuracy: 0.5858 - val_loss: 1.8587

Epoch 17/50
16/16 1s 48ms/step -
accuracy: 0.5488 - loss: 1.7402 - val_accuracy: 0.5533 - val_loss: 1.8296

Epoch 18/50
16/16 1s 48ms/step -
accuracy: 0.5811 - loss: 1.6671 - val_accuracy: 0.6124 - val_loss: 1.7394

Epoch 19/50
16/16 1s 47ms/step -
accuracy: 0.5831 - loss: 1.6190 - val_accuracy: 0.6036 - val_loss: 1.7051

Epoch 20/50
16/16 1s 47ms/step -
accuracy: 0.5946 - loss: 1.5596 - val_accuracy: 0.6095 - val_loss: 1.5964

Epoch 21/50
16/16 1s 47ms/step -
accuracy: 0.6252 - loss: 1.4615 - val_accuracy: 0.6598 - val_loss: 1.5475

Epoch 22/50
16/16 1s 47ms/step -
accuracy: 0.6579 - loss: 1.3494 - val_accuracy: 0.6657 - val_loss: 1.4490

Epoch 23/50
16/16 1s 48ms/step -
accuracy: 0.6758 - loss: 1.3400 - val_accuracy: 0.7041 - val_loss: 1.4354

Epoch 24/50
16/16 1s 47ms/step -
accuracy: 0.6648 - loss: 1.2927 - val_accuracy: 0.6982 - val_loss: 1.3785

Epoch 25/50
16/16 1s 50ms/step -
accuracy: 0.6837 - loss: 1.2203 - val_accuracy: 0.7071 - val_loss: 1.3459

Epoch 26/50
16/16 1s 47ms/step -
accuracy: 0.7259 - loss: 1.1556 - val_accuracy: 0.6982 - val_loss: 1.3418

Epoch 27/50
16/16 1s 47ms/step -
accuracy: 0.7278 - loss: 1.0756 - val_accuracy: 0.7337 - val_loss: 1.2258

Epoch 28/50
16/16 1s 48ms/step -
accuracy: 0.7287 - loss: 1.0332 - val_accuracy: 0.7426 - val_loss: 1.2811

Epoch 29/50
16/16 1s 47ms/step -
accuracy: 0.7304 - loss: 1.0409 - val_accuracy: 0.7456 - val_loss: 1.1965

Epoch 30/50
16/16 1s 47ms/step -
accuracy: 0.7602 - loss: 0.9565 - val_accuracy: 0.7337 - val_loss: 1.2489

Epoch 31/50
16/16 1s 48ms/step -
accuracy: 0.7599 - loss: 0.9259 - val_accuracy: 0.7663 - val_loss: 1.1297

Epoch 32/50
16/16 1s 47ms/step -
accuracy: 0.7937 - loss: 0.8706 - val_accuracy: 0.7515 - val_loss: 1.1258

Epoch 33/50
16/16 1s 47ms/step -
accuracy: 0.7872 - loss: 0.8456 - val_accuracy: 0.7485 - val_loss: 1.0986

Epoch 34/50
16/16 1s 48ms/step -
accuracy: 0.8140 - loss: 0.8400 - val_accuracy: 0.7278 - val_loss: 1.1280

Epoch 35/50
16/16 1s 48ms/step -
accuracy: 0.7886 - loss: 0.8255 - val_accuracy: 0.7751 - val_loss: 1.0405

Epoch 36/50
16/16 1s 47ms/step -
accuracy: 0.7889 - loss: 0.7891 - val_accuracy: 0.7663 - val_loss: 1.0391

Epoch 37/50
16/16 1s 47ms/step -
accuracy: 0.8137 - loss: 0.7576 - val_accuracy: 0.7722 - val_loss: 1.0000

```

Epoch 38/50
16/16          1s 47ms/step -
accuracy: 0.8301 - loss: 0.7096 - val_accuracy: 0.7692 - val_loss: 1.0137
Epoch 39/50
16/16          1s 47ms/step -
accuracy: 0.8383 - loss: 0.6953 - val_accuracy: 0.7781 - val_loss: 1.0482
Epoch 40/50
16/16          1s 52ms/step -
accuracy: 0.8430 - loss: 0.6941 - val_accuracy: 0.7574 - val_loss: 1.0049
Epoch 41/50
16/16          1s 47ms/step -
accuracy: 0.8420 - loss: 0.6595 - val_accuracy: 0.7870 - val_loss: 0.9614
Epoch 42/50
16/16          1s 47ms/step -
accuracy: 0.8561 - loss: 0.6551 - val_accuracy: 0.7781 - val_loss: 0.9219
Epoch 43/50
16/16          1s 48ms/step -
accuracy: 0.8615 - loss: 0.5901 - val_accuracy: 0.7722 - val_loss: 0.9490
Epoch 44/50
16/16          1s 48ms/step -
accuracy: 0.8643 - loss: 0.6077 - val_accuracy: 0.7604 - val_loss: 0.9769
Epoch 45/50
16/16          1s 47ms/step -
accuracy: 0.8799 - loss: 0.5942 - val_accuracy: 0.7692 - val_loss: 0.9981
Epoch 46/50
16/16          1s 47ms/step -
accuracy: 0.8807 - loss: 0.5415 - val_accuracy: 0.7663 - val_loss: 0.9183
Epoch 47/50
16/16          1s 48ms/step -
accuracy: 0.9026 - loss: 0.5172 - val_accuracy: 0.7870 - val_loss: 0.9016
Epoch 48/50
16/16          1s 48ms/step -
accuracy: 0.8831 - loss: 0.4932 - val_accuracy: 0.7751 - val_loss: 0.8911
Epoch 49/50
16/16          1s 47ms/step -
accuracy: 0.8864 - loss: 0.4707 - val_accuracy: 0.7811 - val_loss: 0.8741
Epoch 50/50
16/16          1s 47ms/step -
accuracy: 0.8986 - loss: 0.4725 - val_accuracy: 0.7899 - val_loss: 0.8613

```

```

[14]: test_loss, test_accuracy = model.evaluate(test_ds)
      print(f"Test Accuracy: {test_accuracy:.4f}")
      print(f"Test Loss: {test_loss:.4f}")

```


```

6/6          0s 10ms/step -
accuracy: 0.7523 - loss: 0.9464
Test Accuracy: 0.7611
Test Loss: 0.9174

```

```
[15]: import matplotlib.pyplot as plt
from sklearn.metrics import classification_report, confusion_matrix
import seaborn as sns
```

```
[16]: plt.figure(figsize=(12, 4))
plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')
# Plot training & validation loss values
plt.subplot(1, 2, 2)
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()
```



```
[17]: y_true, y_pred = [], []
target_names = [label_map[i] for i in range(len(label_map))]
for X_batch, y_batch in test_ds:
    y_true.append(y_batch.numpy())

    batch_pred = model.predict(X_batch, verbose=0)
    y_pred.append(np.argmax(batch_pred, axis=1))

y_true = np.concatenate(y_true)
```

```

y_pred = np.concatenate(y_pred)

print(classification_report(
    y_true, y_pred,
    digits=3,
    target_names=target_names
))

cm = confusion_matrix(y_true, y_pred, labels=range(len(label_map)))
labels = [label_map[i] for i in range(len(label_map))]

plt.figure(figsize=(10, 8))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues",
            xticklabels=labels, yticklabels=labels)
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.title("Confusion Matrix - Test Set")
plt.show()

```

	precision	recall	f1-score	support
A	0.583	0.875	0.700	8
B	1.000	0.600	0.750	10
C	0.875	0.778	0.824	18
D	0.235	0.444	0.308	9
E	0.944	0.944	0.944	18
F	0.500	0.667	0.571	6
G	0.889	0.889	0.889	9
H	0.833	0.556	0.667	9
I	0.952	0.909	0.930	22
J	0.656	1.000	0.792	21
K	0.800	0.364	0.500	11
L	1.000	0.737	0.848	19
M	1.000	0.286	0.444	7
N	0.333	0.333	0.333	6
O	0.741	0.909	0.816	22
P	0.667	0.222	0.333	9
Q	1.000	0.556	0.714	9
R	0.750	0.947	0.837	19
S	0.833	0.455	0.588	11
T	0.474	0.692	0.562	13
U	0.850	0.944	0.895	18
V	1.000	1.000	1.000	16
W	0.765	0.765	0.765	17
X	1.000	0.500	0.667	8
Y	1.000	0.400	0.571	5
Z	0.792	1.000	0.884	19

accuracy			0.761	339
macro avg	0.787	0.684	0.697	339
weighted avg	0.805	0.761	0.756	339

