

# COMPARISON\_MediaPipe+CNN

June 21, 2025

Paper Reference : <https://j-innovative.org/index.php/Innovative/article/download/15199/10372/26113>

```
[ ]: import os
    from modules.SignLanguageProcessor import load_and_preprocess_data, parse_frame
```

```
[ ]: ROOT_PATH = ''
    sequences, labels, label_map = load_and_preprocess_data(os.path.
        ↪join(ROOT_PATH, 'data'))
```

```
[16]: num_classes = len(label_map)
```

```
[17]: len(labels)
```

```
[17]: 3413
```

```
[18]: sequences.shape
```

```
[18]: (3413, 3, 61, 3)
```

```
[19]: from sklearn.model_selection import train_test_split

X_train, X_temp, y_train, y_temp = train_test_split(
    sequences, labels, test_size=0.4, stratify=labels, random_state=42
)

X_val, X_test, y_val, y_test = train_test_split(
    X_temp, y_temp, test_size=0.5, stratify=y_temp, random_state=42
)
```

```
[20]: import numpy as np
    def normalize_landmark_data(X):
        """
        Normalize the landmark features (x, y) to have zero mean and unit variance
        ↪across the training set.
        Assumes X shape is (N, F, L, T), where F=3 (x, y, vis).
        """
        X = X.copy()
        # Flatten across all samples, landmarks, and frames
```

```

x_vals = X[:, 0, :, :].flatten()
y_vals = X[:, 1, :, :].flatten()

# Compute mean and std
x_mean, x_std = np.mean(x_vals), np.std(x_vals)
y_mean, y_std = np.mean(y_vals), np.std(y_vals)

# Normalize
X[:, 0, :, :] = (X[:, 0, :, :] - x_mean) / x_std
X[:, 1, :, :] = (X[:, 1, :, :] - y_mean) / y_std

return X, (x_mean, x_std), (y_mean, y_std)

def apply_normalization(X, x_mean, x_std, y_mean, y_std):
    X = X.copy()
    X[:, 0, :, :] = (X[:, 0, :, :] - x_mean) / x_std
    X[:, 1, :, :] = (X[:, 1, :, :] - y_mean) / y_std
    return X

```

```

[21]: def reshape_frames_for_cnn(X, y):
    """
    Reshape a dataset of (N, F, L, T) into (N*T, L, F, 1) for Conv2D,
    where each frame becomes its own sample.

    Parameters:
    - X: np.ndarray of shape (N, F, L, T)
    - y: np.ndarray of shape (N,)

    Returns:
    - reshaped_X: np.ndarray of shape (N*T, L, F, 1)
    - reshaped_y: np.ndarray of shape (N*T,)
    """
    reshaped_X = []
    reshaped_y = []

    for sample, label in zip(X, y):
        T = sample.shape[-1]
        for t in range(T):
            frame = sample[:, :, t].T[..., np.newaxis]
            reshaped_X.append(frame)
            reshaped_y.append(label)

    reshaped_X = np.array(reshaped_X)
    reshaped_y = np.array(reshaped_y)
    return reshaped_X, reshaped_y

```

```
[22]: X_train_norm, (x_mean, x_std), (y_mean, y_std) = ↳
        ↳normalize_landmark_data(X_train)
X_val_norm = apply_normalization(X_val, x_mean, x_std, y_mean, y_std)
X_test_norm = apply_normalization(X_test, x_mean, x_std, y_mean, y_std)

X_train_cnn, y_train_cnn = reshape_frames_for_cnn(X_train_norm, y_train)
X_val_cnn, y_val_cnn      = reshape_frames_for_cnn(X_val_norm, y_val)
X_test_cnn, y_test_cnn    = reshape_frames_for_cnn(X_test_norm, y_test)

print(X_train_cnn.shape)
print(y_train_cnn.shape)

(6141, 61, 3, 1)
(6141,)
```

```
[23]: input_shape = X_train_cnn.shape[1:]
print(input_shape)

(61, 3, 1)
```

```
[24]: import tensorflow as tf

train_ds = tf.data.Dataset.from_tensor_slices((X_train_cnn, y_train_cnn))
train_ds = train_ds.shuffle(buffer_size=1000).batch(64).prefetch(tf.data.
↳AUTOTUNE)

val_ds = tf.data.Dataset.from_tensor_slices((X_val_cnn, y_val_cnn))
val_ds = val_ds.batch(64).prefetch(tf.data.AUTOTUNE)

test_ds = tf.data.Dataset.from_tensor_slices((X_test_cnn, y_test_cnn))
test_ds = test_ds.batch(64).prefetch(tf.data.AUTOTUNE)
```

```
[25]: from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Dropout, Flatten, ↳
↳Dense, BatchNormalization, Input

cnn_model = Sequential([
    Input(input_shape),
    Conv2D(32, (3, 2), activation='relu', padding='same'),
    MaxPooling2D((2, 1)),
    Dropout(0.25),
    Conv2D(64, (3, 2), activation='relu', padding='same'),
    MaxPooling2D(pool_size=(2, 1)),
    Dropout(0.25),
    Flatten(),
    Dense(128, activation='relu'),
    Dropout(0.2),
    Dense(num_classes, activation='softmax')
```

```
])
cnn_model.compile(optimizer='adam', loss='sparse_categorical_crossentropy',
↳metrics=['accuracy'])
```

```
[26]: history = cnn_model.fit(train_ds,validation_data=val_ds, epochs=50,
↳batch_size=64)
```

```
Epoch 1/50
96/96          2s 10ms/step -
accuracy: 0.0670 - loss: 3.6065 - val_accuracy: 0.1215 - val_loss: 3.0120
Epoch 2/50
96/96          1s 8ms/step -
accuracy: 0.1005 - loss: 3.0360 - val_accuracy: 0.1396 - val_loss: 2.8006
Epoch 3/50
96/96          1s 8ms/step -
accuracy: 0.1372 - loss: 2.8652 - val_accuracy: 0.2016 - val_loss: 2.6945
Epoch 4/50
96/96          1s 8ms/step -
accuracy: 0.1682 - loss: 2.7492 - val_accuracy: 0.2484 - val_loss: 2.5616
Epoch 5/50
96/96          1s 8ms/step -
accuracy: 0.2075 - loss: 2.6354 - val_accuracy: 0.2596 - val_loss: 2.4914
Epoch 6/50
96/96          1s 8ms/step -
accuracy: 0.2305 - loss: 2.5260 - val_accuracy: 0.2879 - val_loss: 2.3693
Epoch 7/50
96/96          1s 8ms/step -
accuracy: 0.2644 - loss: 2.4390 - val_accuracy: 0.3011 - val_loss: 2.2906
Epoch 8/50
96/96          1s 8ms/step -
accuracy: 0.2873 - loss: 2.3636 - val_accuracy: 0.3490 - val_loss: 2.2312
Epoch 9/50
96/96          1s 8ms/step -
accuracy: 0.3036 - loss: 2.2794 - val_accuracy: 0.3631 - val_loss: 2.1365
Epoch 10/50
96/96          1s 8ms/step -
accuracy: 0.3142 - loss: 2.2113 - val_accuracy: 0.3875 - val_loss: 2.0848
Epoch 11/50
96/96          1s 8ms/step -
accuracy: 0.3441 - loss: 2.1542 - val_accuracy: 0.4109 - val_loss: 2.0155
Epoch 12/50
96/96          1s 8ms/step -
accuracy: 0.3581 - loss: 2.0818 - val_accuracy: 0.4178 - val_loss: 1.9533
Epoch 13/50
96/96          1s 8ms/step -
accuracy: 0.3731 - loss: 2.0186 - val_accuracy: 0.4378 - val_loss: 1.8896
Epoch 14/50
96/96          1s 8ms/step -
```

accuracy: 0.3940 - loss: 1.9732 - val\_accuracy: 0.4412 - val\_loss: 1.8445  
 Epoch 15/50  
 96/96 1s 8ms/step -  
 accuracy: 0.4118 - loss: 1.9094 - val\_accuracy: 0.4495 - val\_loss: 1.8477  
 Epoch 16/50  
 96/96 1s 8ms/step -  
 accuracy: 0.4218 - loss: 1.8669 - val\_accuracy: 0.4632 - val\_loss: 1.7870  
 Epoch 17/50  
 96/96 1s 8ms/step -  
 accuracy: 0.4271 - loss: 1.8450 - val\_accuracy: 0.4797 - val\_loss: 1.7529  
 Epoch 18/50  
 96/96 1s 8ms/step -  
 accuracy: 0.4397 - loss: 1.8074 - val\_accuracy: 0.4739 - val\_loss: 1.7555  
 Epoch 19/50  
 96/96 1s 8ms/step -  
 accuracy: 0.4487 - loss: 1.7875 - val\_accuracy: 0.4915 - val\_loss: 1.7247  
 Epoch 20/50  
 96/96 1s 8ms/step -  
 accuracy: 0.4578 - loss: 1.7494 - val\_accuracy: 0.4876 - val\_loss: 1.7083  
 Epoch 21/50  
 96/96 1s 8ms/step -  
 accuracy: 0.4743 - loss: 1.7076 - val\_accuracy: 0.4939 - val\_loss: 1.7087  
 Epoch 22/50  
 96/96 1s 8ms/step -  
 accuracy: 0.4591 - loss: 1.7266 - val\_accuracy: 0.5095 - val\_loss: 1.6672  
 Epoch 23/50  
 96/96 1s 8ms/step -  
 accuracy: 0.4692 - loss: 1.6908 - val\_accuracy: 0.5100 - val\_loss: 1.6503  
 Epoch 24/50  
 96/96 1s 8ms/step -  
 accuracy: 0.4728 - loss: 1.6675 - val\_accuracy: 0.5154 - val\_loss: 1.6285  
 Epoch 25/50  
 96/96 1s 8ms/step -  
 accuracy: 0.4837 - loss: 1.6586 - val\_accuracy: 0.5295 - val\_loss: 1.6044  
 Epoch 26/50  
 96/96 1s 8ms/step -  
 accuracy: 0.4809 - loss: 1.6423 - val\_accuracy: 0.5305 - val\_loss: 1.6145  
 Epoch 27/50  
 96/96 1s 8ms/step -  
 accuracy: 0.5002 - loss: 1.5973 - val\_accuracy: 0.5281 - val\_loss: 1.6160  
 Epoch 28/50  
 96/96 1s 8ms/step -  
 accuracy: 0.4962 - loss: 1.5962 - val\_accuracy: 0.5378 - val\_loss: 1.5876  
 Epoch 29/50  
 96/96 1s 8ms/step -  
 accuracy: 0.4920 - loss: 1.5688 - val\_accuracy: 0.5417 - val\_loss: 1.5669  
 Epoch 30/50  
 96/96 1s 8ms/step -

accuracy: 0.5170 - loss: 1.5396 - val\_accuracy: 0.5437 - val\_loss: 1.5693  
 Epoch 31/50  
 96/96 1s 8ms/step -  
 accuracy: 0.5149 - loss: 1.5514 - val\_accuracy: 0.5520 - val\_loss: 1.5508  
 Epoch 32/50  
 96/96 1s 8ms/step -  
 accuracy: 0.5157 - loss: 1.5185 - val\_accuracy: 0.5481 - val\_loss: 1.5579  
 Epoch 33/50  
 96/96 1s 8ms/step -  
 accuracy: 0.5291 - loss: 1.4850 - val\_accuracy: 0.5373 - val\_loss: 1.5638  
 Epoch 34/50  
 96/96 1s 8ms/step -  
 accuracy: 0.5142 - loss: 1.5288 - val\_accuracy: 0.5583 - val\_loss: 1.5380  
 Epoch 35/50  
 96/96 1s 8ms/step -  
 accuracy: 0.5314 - loss: 1.4901 - val\_accuracy: 0.5539 - val\_loss: 1.5378  
 Epoch 36/50  
 96/96 1s 8ms/step -  
 accuracy: 0.5272 - loss: 1.4751 - val\_accuracy: 0.5554 - val\_loss: 1.5386  
 Epoch 37/50  
 96/96 1s 8ms/step -  
 accuracy: 0.5338 - loss: 1.4683 - val\_accuracy: 0.5612 - val\_loss: 1.5260  
 Epoch 38/50  
 96/96 1s 8ms/step -  
 accuracy: 0.5261 - loss: 1.4780 - val\_accuracy: 0.5569 - val\_loss: 1.5180  
 Epoch 39/50  
 96/96 1s 8ms/step -  
 accuracy: 0.5259 - loss: 1.4645 - val\_accuracy: 0.5632 - val\_loss: 1.5240  
 Epoch 40/50  
 96/96 1s 8ms/step -  
 accuracy: 0.5457 - loss: 1.4485 - val\_accuracy: 0.5476 - val\_loss: 1.5257  
 Epoch 41/50  
 96/96 1s 8ms/step -  
 accuracy: 0.5492 - loss: 1.4288 - val\_accuracy: 0.5661 - val\_loss: 1.5123  
 Epoch 42/50  
 96/96 1s 8ms/step -  
 accuracy: 0.5490 - loss: 1.4285 - val\_accuracy: 0.5676 - val\_loss: 1.5095  
 Epoch 43/50  
 96/96 1s 8ms/step -  
 accuracy: 0.5433 - loss: 1.4381 - val\_accuracy: 0.5656 - val\_loss: 1.5106  
 Epoch 44/50  
 96/96 1s 8ms/step -  
 accuracy: 0.5527 - loss: 1.3903 - val\_accuracy: 0.5661 - val\_loss: 1.4892  
 Epoch 45/50  
 96/96 1s 8ms/step -  
 accuracy: 0.5439 - loss: 1.4025 - val\_accuracy: 0.5632 - val\_loss: 1.5096  
 Epoch 46/50  
 96/96 1s 8ms/step -

```

accuracy: 0.5587 - loss: 1.3836 - val_accuracy: 0.5647 - val_loss: 1.5091
Epoch 47/50
96/96          1s 8ms/step -
accuracy: 0.5538 - loss: 1.3902 - val_accuracy: 0.5710 - val_loss: 1.4939
Epoch 48/50
96/96          1s 8ms/step -
accuracy: 0.5590 - loss: 1.3638 - val_accuracy: 0.5695 - val_loss: 1.5154
Epoch 49/50
96/96          1s 8ms/step -
accuracy: 0.5601 - loss: 1.3747 - val_accuracy: 0.5671 - val_loss: 1.5004
Epoch 50/50
96/96          1s 7ms/step -
accuracy: 0.5760 - loss: 1.3444 - val_accuracy: 0.5661 - val_loss: 1.5173

```

```

[27]: test_loss, test_accuracy = cnn_model.evaluate(test_ds)
      print(f"Test Accuracy: {test_accuracy:.4f}")
      print(f"Test Loss: {test_loss:.4f}")

```

```

33/33          0s 3ms/step -
accuracy: 0.5737 - loss: 1.4620
Test Accuracy: 0.5588
Test Loss: 1.5120

```

```


[28]: import matplotlib.pyplot as plt
      from sklearn.metrics import classification_report, confusion_matrix
      import seaborn as sns

```

```

[29]: plt.figure(figsize=(12, 4))
      plt.subplot(1, 2, 1)
      plt.plot(history.history['accuracy'])
      plt.plot(history.history['val_accuracy'])
      plt.title('Model accuracy')
      plt.ylabel('Accuracy')
      plt.xlabel('Epoch')
      plt.legend(['Train', 'Validation'], loc='upper left')
      # Plot training & validation loss values
      plt.subplot(1, 2, 2)
      plt.plot(history.history['loss'])
      plt.plot(history.history['val_loss'])
      plt.title('Model loss')
      plt.ylabel('Loss')
      plt.xlabel('Epoch')
      plt.legend(['Train', 'Validation'], loc='upper left')
      plt.show()

```



```
[30]: y_true, y_pred = [], []
target_names = [label_map[i] for i in range(len(label_map))]
for X_batch, y_batch in test_ds:
    y_true.append(y_batch.numpy())

    batch_pred = cnn_model.predict(X_batch, verbose=0)
    y_pred.append(np.argmax(batch_pred, axis=1))

y_true = np.concatenate(y_true)
y_pred = np.concatenate(y_pred)

print(classification_report(
    y_true, y_pred,
    digits=3,
    target_names=target_names
))

cm = confusion_matrix(y_true, y_pred, labels=range(len(label_map)))
labels = [label_map[i] for i in range(len(label_map))]

plt.figure(figsize=(10, 8))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues",
            xticklabels=labels, yticklabels=labels)
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.title("Confusion Matrix - Test Set")
plt.show()
```

|   | precision | recall | f1-score | support |
|---|-----------|--------|----------|---------|
| A | 0.812     | 0.542  | 0.650    | 24      |
| B | 0.708     | 0.567  | 0.630    | 30      |



|               |       |       |       |    |
|---------------|-------|-------|-------|----|
| C             | 0.651 | 0.519 | 0.577 | 54 |
| D             | 0.323 | 0.370 | 0.345 | 27 |
| E             | 0.811 | 0.588 | 0.682 | 51 |
| F             | 0.600 | 0.167 | 0.261 | 18 |
| G             | 0.750 | 0.444 | 0.558 | 27 |
| H             | 1.000 | 0.370 | 0.541 | 27 |
| I             | 0.745 | 0.556 | 0.636 | 63 |
| J             | 0.800 | 0.508 | 0.621 | 63 |
| K             | 0.667 | 0.303 | 0.417 | 33 |
| L             | 0.773 | 0.596 | 0.673 | 57 |
| M             | 0.667 | 0.381 | 0.485 | 21 |
| N             | 0.375 | 0.167 | 0.231 | 18 |
| O             | 0.222 | 0.970 | 0.362 | 66 |
| P             | 0.667 | 0.222 | 0.333 | 27 |
| Q             | 0.692 | 0.333 | 0.450 | 27 |
| R             | 0.667 | 0.491 | 0.566 | 57 |
| S             | 0.846 | 0.306 | 0.449 | 36 |
| T             | 0.531 | 0.436 | 0.479 | 39 |
| U             | 0.464 | 0.481 | 0.473 | 54 |
| V             | 0.718 | 0.583 | 0.644 | 48 |
| W             | 0.221 | 0.627 | 0.327 | 51 |
| X             | 0.714 | 0.417 | 0.526 | 24 |
| Y             | 0.875 | 0.467 | 0.609 | 15 |
| Z             | 0.660 | 0.579 | 0.617 | 57 |
| baca          | 0.818 | 0.692 | 0.750 | 39 |
| bantu         | 0.774 | 0.727 | 0.750 | 33 |
| bapak         | 0.419 | 0.462 | 0.439 | 39 |
| buangairkecil | 0.875 | 0.667 | 0.757 | 21 |
| buat          | 0.581 | 0.923 | 0.713 | 39 |
| halo          | 0.621 | 0.759 | 0.683 | 54 |
| ibu           | 0.750 | 0.250 | 0.375 | 12 |
| kamu          | 0.750 | 0.474 | 0.581 | 57 |
| maaf          | 0.635 | 0.611 | 0.623 | 54 |
| makan         | 0.769 | 0.238 | 0.364 | 42 |
| mau           | 0.549 | 0.765 | 0.639 | 51 |
| nama          | 0.732 | 0.556 | 0.632 | 54 |
| pagi          | 0.571 | 0.702 | 0.630 | 57 |
| paham         | 0.677 | 0.700 | 0.689 | 60 |
| sakit         | 1.000 | 0.667 | 0.800 | 9  |
| sama-sama     | 0.783 | 0.653 | 0.712 | 72 |
| saya          | 0.462 | 0.333 | 0.387 | 18 |
| selamat       | 0.621 | 0.706 | 0.661 | 51 |
| siapa         | 0.852 | 0.639 | 0.730 | 36 |
| tanya         | 0.769 | 0.392 | 0.519 | 51 |
| tempat        | 1.000 | 0.667 | 0.800 | 12 |
| terima-kasih  | 0.757 | 0.519 | 0.615 | 54 |
| terlambat     | 0.846 | 0.564 | 0.677 | 39 |
| tidak         | 0.323 | 0.762 | 0.454 | 42 |

|              |       |       |       |      |
|--------------|-------|-------|-------|------|
| tolong       | 0.531 | 0.436 | 0.479 | 39   |
| accuracy     |       |       | 0.559 | 2049 |
| macro avg    | 0.675 | 0.527 | 0.561 | 2049 |
| weighted avg | 0.659 | 0.559 | 0.572 | 2049 |

