

# COMPARISON\_MobileNetV2

June 21, 2025

```
[ ]: import tensorflow as tf
from tensorflow.keras.utils import to_categorical
import os
from PIL import Image, UnidentifiedImageError
import shutil

# Configuration
IMG_SIZE = (96, 96)
BATCH_SIZE = 32
VALIDATION_SPLIT = 0.4
SEED = 42
ROOT_PATH = ''
DATASET_PATH = os.path.join(ROOT_PATH, "raw_data")
CORRUPT_PATH = os.path.join(ROOT_PATH, "corrupt_images")
os.makedirs(CORRUPT_PATH, exist_ok=True)

for root, dirs, files in os.walk(DATASET_PATH):
    for file in files:
        ext = os.path.splitext(file)[1].lower()
        if ext in [".jpg", ".jpeg", ".png", ".bmp", ".gif"]:
            path = os.path.join(root, file)
            try:
                with Image.open(path) as img:
                    img.verify() # Check integrity
            except (UnidentifiedImageError, OSError, IOError) as e:
                # Move the corrupt image
                print(f"Corrupt image found: {path} - moving to {CORRUPT_PATH}")
                dest_path = os.path.join(CORRUPT_PATH, os.path.relpath(path,
↳DATASET_PATH))
                os.makedirs(os.path.dirname(dest_path), exist_ok=True)
                shutil.move(path, dest_path)

LANDMARK_DIR = os.path.join(ROOT_PATH, "data")
RAW_IMAGE_DIR = os.path.join(ROOT_PATH, "raw_data")
FILTERED_IMAGE_DIR = os.path.join(ROOT_PATH, "filtered_raw_data")
DATASET_PATH = FILTERED_IMAGE_DIR
# Supported image extensions
```

```

IMAGE_EXTENSIONS = ['.jpg', '.jpeg', '.png', '.bmp']

# Create filtered output structure
os.makedirs(FILTERED_IMAGE_DIR, exist_ok=True)

for class_name in os.listdir(LANDMARK_DIR):
    if class_name == 'debug':
        continue
    landmark_class_dir = os.path.join(LANDMARK_DIR, class_name)
    raw_class_dir = os.path.join(RAW_IMAGE_DIR, class_name)
    filtered_class_dir = os.path.join(FILTERED_IMAGE_DIR, class_name)
    os.makedirs(filtered_class_dir, exist_ok=True)

    for file in os.listdir(landmark_class_dir):
        if not file.endswith("_landmarks.json"):
            continue

        # Get base filename without "_landmarks.json"
        base_name = file.replace("_landmarks.json", "")

        # Look for corresponding image in raw directory
        for ext in IMAGE_EXTENSIONS:
            image_file = os.path.join(raw_class_dir, base_name + ext)
            if os.path.exists(image_file):
                # Copy to filtered folder
                shutil.copy(image_file, os.path.join(filtered_class_dir, os.
↳ path.basename(image_file)))
                break

# Load training dataset with validation split
train_ds = tf.keras.preprocessing.image_dataset_from_directory(
    DATASET_PATH,
    validation_split=VALIDATION_SPLIT,
    subset="training",
    seed=SEED,
    color_mode="rgb",
    image_size=IMG_SIZE,
    batch_size=BATCH_SIZE
)
num_classes = len(train_ds.class_names)
label_map = train_ds.class_names

val_ds = tf.keras.preprocessing.image_dataset_from_directory(
    DATASET_PATH,
    validation_split=VALIDATION_SPLIT,
    subset="validation",
    seed=SEED,

```

```

        color_mode="rgb",
        image_size=IMG_SIZE,
        batch_size=BATCH_SIZE
    )

    test_ds = val_ds.shard(2,0)
    val_ds = val_ds.shard(2,1)
    # Normalize pixel values to [0, 1]
    normalization_layer = tf.keras.layers.Rescaling(1./255)
    train_ds = train_ds.map(lambda x, y: (normalization_layer(x), y))
    val_ds = val_ds.map(lambda x, y: (normalization_layer(x), y))
    test_ds = test_ds.map(lambda x, y: (normalization_layer(x), y))
    # Cache and prefetch for performance
    AUTOTUNE = tf.data.AUTOTUNE
    train_ds = train_ds.cache().shuffle(1000).prefetch(buffer_size=AUTOTUNE)
    val_ds = val_ds.cache().prefetch(buffer_size=AUTOTUNE)
    test_ds = test_ds.cache().prefetch(buffer_size=AUTOTUNE)

```

Found 3413 files belonging to 51 classes.

Using 2048 files for training.

Found 3413 files belonging to 51 classes.

Using 1365 files for validation.

```

[5]: from tensorflow.keras.models import Sequential
    from tensorflow.keras.layers import Conv2D, MaxPooling2D, Dropout,
        BatchNormalization
    from tensorflow.keras.layers import Flatten, Dense, GlobalAveragePooling2D
    from tensorflow.keras.optimizers import Adam
    from tensorflow.keras.applications import MobileNetV2
    base_model = MobileNetV2(input_shape=(96, 96, 3), include_top=False,
        weights='imagenet')
    base_model.trainable = False

    model = Sequential([
        base_model,
        GlobalAveragePooling2D(),
        Dropout(0.3),
        Dense(128, activation='relu'),
        Dropout(0.3),
        Dense(num_classes, activation='softmax')
    ])

    model.compile(optimizer=Adam(1e-3),
        loss='sparse_categorical_crossentropy',
        metrics=['accuracy'])

```

```

[6]: history = model.fit(train_ds, validation_data=val_ds, epochs=50)

```

Epoch 1/50  
64/64 18s 132ms/step -  
accuracy: 0.0737 - loss: 4.0134 - val\_accuracy: 0.3348 - val\_loss: 2.6370  
Epoch 2/50  
64/64 4s 61ms/step -  
accuracy: 0.2878 - loss: 2.6495 - val\_accuracy: 0.4866 - val\_loss: 1.9093  
Epoch 3/50  
64/64 4s 60ms/step -  
accuracy: 0.4534 - loss: 1.9079 - val\_accuracy: 0.5417 - val\_loss: 1.5783  
Epoch 4/50  
64/64 4s 60ms/step -  
accuracy: 0.5672 - loss: 1.4972 - val\_accuracy: 0.6071 - val\_loss: 1.3757  
Epoch 5/50  
64/64 4s 60ms/step -  
accuracy: 0.6182 - loss: 1.2579 - val\_accuracy: 0.6205 - val\_loss: 1.2890  
Epoch 6/50  
64/64 4s 60ms/step -  
accuracy: 0.6848 - loss: 1.0517 - val\_accuracy: 0.6205 - val\_loss: 1.2737  
Epoch 7/50  
64/64 4s 60ms/step -  
accuracy: 0.7213 - loss: 0.9044 - val\_accuracy: 0.6235 - val\_loss: 1.2364  
Epoch 8/50  
64/64 4s 60ms/step -  
accuracy: 0.7767 - loss: 0.7365 - val\_accuracy: 0.6295 - val\_loss: 1.2137  
Epoch 9/50  
64/64 4s 60ms/step -  
accuracy: 0.7840 - loss: 0.7048 - val\_accuracy: 0.6533 - val\_loss: 1.1466  
Epoch 10/50  
64/64 4s 61ms/step -  
accuracy: 0.8021 - loss: 0.6284 - val\_accuracy: 0.6562 - val\_loss: 1.1585  
Epoch 11/50  
64/64 4s 60ms/step -  
accuracy: 0.8454 - loss: 0.5215 - val\_accuracy: 0.6592 - val\_loss: 1.1608  
Epoch 12/50  
64/64 4s 60ms/step -  
accuracy: 0.8459 - loss: 0.5075 - val\_accuracy: 0.6488 - val\_loss: 1.1600  
Epoch 13/50  
64/64 4s 60ms/step -  
accuracy: 0.8566 - loss: 0.4898 - val\_accuracy: 0.6622 - val\_loss: 1.1625  
Epoch 14/50  
64/64 4s 60ms/step -  
accuracy: 0.8676 - loss: 0.4579 - val\_accuracy: 0.6399 - val\_loss: 1.1999  
Epoch 15/50  
64/64 4s 61ms/step -  
accuracy: 0.8525 - loss: 0.4536 - val\_accuracy: 0.6726 - val\_loss: 1.1428  
Epoch 16/50  
64/64 4s 64ms/step -  
accuracy: 0.8711 - loss: 0.3842 - val\_accuracy: 0.6607 - val\_loss: 1.1278

Epoch 17/50  
64/64 4s 64ms/step -  
accuracy: 0.8895 - loss: 0.3503 - val\_accuracy: 0.6592 - val\_loss: 1.1862  
Epoch 18/50  
64/64 4s 63ms/step -  
accuracy: 0.9022 - loss: 0.3197 - val\_accuracy: 0.6637 - val\_loss: 1.1861  
Epoch 19/50  
64/64 4s 63ms/step -  
accuracy: 0.9059 - loss: 0.3081 - val\_accuracy: 0.6682 - val\_loss: 1.1808  
Epoch 20/50  
64/64 4s 63ms/step -  
accuracy: 0.9200 - loss: 0.2911 - val\_accuracy: 0.6622 - val\_loss: 1.2058  
Epoch 21/50  
64/64 4s 64ms/step -  
accuracy: 0.9169 - loss: 0.2797 - val\_accuracy: 0.6652 - val\_loss: 1.2206  
Epoch 22/50  
64/64 4s 63ms/step -  
accuracy: 0.9109 - loss: 0.2708 - val\_accuracy: 0.6577 - val\_loss: 1.2065  
Epoch 23/50  
64/64 4s 63ms/step -  
accuracy: 0.9284 - loss: 0.2414 - val\_accuracy: 0.6696 - val\_loss: 1.1733  
Epoch 24/50  
64/64 4s 64ms/step -  
accuracy: 0.9346 - loss: 0.2258 - val\_accuracy: 0.6771 - val\_loss: 1.2226  
Epoch 25/50  
64/64 4s 63ms/step -  
accuracy: 0.9218 - loss: 0.2604 - val\_accuracy: 0.6622 - val\_loss: 1.2074  
Epoch 26/50  
64/64 4s 63ms/step -  
accuracy: 0.9451 - loss: 0.2080 - val\_accuracy: 0.6652 - val\_loss: 1.2440  
Epoch 27/50  
64/64 4s 63ms/step -  
accuracy: 0.9360 - loss: 0.2120 - val\_accuracy: 0.6771 - val\_loss: 1.1871  
Epoch 28/50  
64/64 4s 62ms/step -  
accuracy: 0.9481 - loss: 0.1854 - val\_accuracy: 0.6592 - val\_loss: 1.2634  
Epoch 29/50  
64/64 4s 62ms/step -  
accuracy: 0.9457 - loss: 0.1882 - val\_accuracy: 0.6667 - val\_loss: 1.2261  
Epoch 30/50  
64/64 4s 64ms/step -  
accuracy: 0.9424 - loss: 0.1867 - val\_accuracy: 0.6845 - val\_loss: 1.2612  
Epoch 31/50  
64/64 4s 62ms/step -  
accuracy: 0.9516 - loss: 0.1662 - val\_accuracy: 0.6577 - val\_loss: 1.3016  
Epoch 32/50  
64/64 4s 62ms/step -  
accuracy: 0.9450 - loss: 0.1642 - val\_accuracy: 0.6667 - val\_loss: 1.2758

Epoch 33/50  
64/64 4s 62ms/step -  
accuracy: 0.9435 - loss: 0.1751 - val\_accuracy: 0.6667 - val\_loss: 1.2586  
Epoch 34/50  
64/64 4s 63ms/step -  
accuracy: 0.9500 - loss: 0.1688 - val\_accuracy: 0.6667 - val\_loss: 1.2970  
Epoch 35/50  
64/64 4s 62ms/step -  
accuracy: 0.9497 - loss: 0.1456 - val\_accuracy: 0.6771 - val\_loss: 1.2745  
Epoch 36/50  
64/64 4s 62ms/step -  
accuracy: 0.9612 - loss: 0.1431 - val\_accuracy: 0.6771 - val\_loss: 1.3573  
Epoch 37/50  
64/64 4s 63ms/step -  
accuracy: 0.9459 - loss: 0.1550 - val\_accuracy: 0.6741 - val\_loss: 1.3014  
Epoch 38/50  
64/64 4s 62ms/step -  
accuracy: 0.9541 - loss: 0.1509 - val\_accuracy: 0.6756 - val\_loss: 1.3280  
Epoch 39/50  
64/64 4s 63ms/step -  
accuracy: 0.9491 - loss: 0.1604 - val\_accuracy: 0.6741 - val\_loss: 1.3327  
Epoch 40/50  
64/64 4s 63ms/step -  
accuracy: 0.9643 - loss: 0.1189 - val\_accuracy: 0.6682 - val\_loss: 1.3403  
Epoch 41/50  
64/64 4s 64ms/step -  
accuracy: 0.9601 - loss: 0.1220 - val\_accuracy: 0.6741 - val\_loss: 1.3032  
Epoch 42/50  
64/64 4s 64ms/step -  
accuracy: 0.9578 - loss: 0.1246 - val\_accuracy: 0.6860 - val\_loss: 1.3289  
Epoch 43/50  
64/64 4s 64ms/step -  
accuracy: 0.9584 - loss: 0.1316 - val\_accuracy: 0.6815 - val\_loss: 1.3691  
Epoch 44/50  
64/64 4s 63ms/step -  
accuracy: 0.9554 - loss: 0.1309 - val\_accuracy: 0.6786 - val\_loss: 1.3708  
Epoch 45/50  
64/64 4s 64ms/step -  
accuracy: 0.9563 - loss: 0.1302 - val\_accuracy: 0.6935 - val\_loss: 1.3532  
Epoch 46/50  
64/64 4s 63ms/step -  
accuracy: 0.9596 - loss: 0.1284 - val\_accuracy: 0.6756 - val\_loss: 1.3537  
Epoch 47/50  
64/64 4s 63ms/step -  
accuracy: 0.9667 - loss: 0.1133 - val\_accuracy: 0.6786 - val\_loss: 1.4074  
Epoch 48/50  
64/64 4s 63ms/step -  
accuracy: 0.9715 - loss: 0.1120 - val\_accuracy: 0.6756 - val\_loss: 1.3911


```
Epoch 49/50
64/64          4s 64ms/step -
accuracy: 0.9687 - loss: 0.1138 - val_accuracy: 0.6726 - val_loss: 1.3987
Epoch 50/50
64/64          4s 64ms/step -
accuracy: 0.9553 - loss: 0.1447 - val_accuracy: 0.6860 - val_loss: 1.3996
```

```
[7]: test_loss, test_accuracy = model.evaluate(test_ds)
      print(f"Test Accuracy: {test_accuracy:.4f}")
      print(f"Test Loss: {test_loss:.4f}")
```

```
22/22          4s 188ms/step -
accuracy: 0.7031 - loss: 1.3048
Test Accuracy: 0.6797
Test Loss: 1.4072
```

```
[8]: import matplotlib.pyplot as plt
      from sklearn.metrics import classification_report, confusion_matrix
      import seaborn as sns
      import numpy as np
```

```
[9]: plt.figure(figsize=(12, 4))
      plt.subplot(1, 2, 1)
      plt.plot(history.history['accuracy'])
      plt.plot(history.history['val_accuracy'])
      plt.title('Model accuracy')
      plt.ylabel('Accuracy')
      plt.xlabel('Epoch')
      plt.legend(['Train', 'Validation'], loc='upper left')
      # Plot training & validation loss values
      plt.subplot(1, 2, 2)
      plt.plot(history.history['loss'])
      plt.plot(history.history['val_loss'])
      plt.title('Model loss')
      plt.ylabel('Loss')
      plt.xlabel('Epoch')
      plt.legend(['Train', 'Validation'], loc='upper left')
      plt.show()
```



```
[10]: y_true, y_pred = [], []
target_names = [label_map[i] for i in range(len(label_map))]
for X_batch, y_batch in test_ds:
    y_true.append(y_batch.numpy())

    batch_pred = model.predict(X_batch, verbose=0)
    y_pred.append(np.argmax(batch_pred, axis=1))

y_true = np.concatenate(y_true)
y_pred = np.concatenate(y_pred)

print(classification_report(
    y_true, y_pred,
    digits=3,
    target_names=target_names
))

cm = confusion_matrix(y_true, y_pred, labels=range(len(label_map)))
labels = [label_map[i] for i in range(len(label_map))]

plt.figure(figsize=(10, 8))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues",
            xticklabels=labels, yticklabels=labels)
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.title("Confusion Matrix - Test Set")
plt.show()
```

c:\Users\chris\.conda\envs\ASLR\Lib\site-packages\sklearn\metrics\\_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.



```
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
c:\Users\chris\.conda\envs\ASLR\Lib\site-
packages\sklearn\metrics\_classification.py:1565: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
c:\Users\chris\.conda\envs\ASLR\Lib\site-
packages\sklearn\metrics\_classification.py:1565: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

|               | precision | recall | f1-score | support |
|---------------|-----------|--------|----------|---------|
| A             | 1.000     | 0.286  | 0.444    | 7       |
| B             | 0.300     | 0.500  | 0.375    | 6       |
| C             | 0.444     | 0.400  | 0.421    | 20      |
| D             | 0.500     | 0.200  | 0.286    | 10      |
| E             | 0.375     | 0.600  | 0.462    | 10      |
| F             | 0.333     | 0.444  | 0.381    | 9       |
| G             | 0.625     | 0.455  | 0.526    | 11      |
| H             | 0.375     | 0.300  | 0.333    | 10      |
| I             | 0.229     | 0.400  | 0.291    | 20      |
| J             | 0.370     | 0.455  | 0.408    | 22      |
| K             | 0.143     | 0.100  | 0.118    | 10      |
| L             | 0.474     | 0.529  | 0.500    | 17      |
| M             | 0.500     | 0.833  | 0.625    | 6       |
| N             | 0.000     | 0.000  | 0.000    | 9       |
| O             | 0.400     | 0.421  | 0.410    | 19      |
| P             | 0.385     | 0.556  | 0.455    | 9       |
| Q             | 0.182     | 0.182  | 0.182    | 11      |
| R             | 0.300     | 0.391  | 0.340    | 23      |
| S             | 0.500     | 0.400  | 0.444    | 15      |
| T             | 0.462     | 0.600  | 0.522    | 10      |
| U             | 0.167     | 0.053  | 0.080    | 19      |
| V             | 0.375     | 0.286  | 0.324    | 21      |
| W             | 0.667     | 0.571  | 0.615    | 14      |
| X             | 0.500     | 0.333  | 0.400    | 6       |
| Y             | 0.333     | 0.200  | 0.250    | 5       |
| Z             | 0.720     | 0.818  | 0.766    | 22      |
| baca          | 0.917     | 1.000  | 0.957    | 11      |
| bantu         | 0.769     | 1.000  | 0.870    | 10      |
| bapak         | 0.923     | 0.923  | 0.923    | 13      |
| buangairkecil | 1.000     | 0.667  | 0.800    | 9       |
| buat          | 0.900     | 0.900  | 0.900    | 10      |
| halo          | 1.000     | 1.000  | 1.000    | 16      |
| ibu           | 1.000     | 1.000  | 1.000    | 6       |
| kamu          | 0.955     | 0.913  | 0.933    | 23      |

|              |       |       |       |     |
|--------------|-------|-------|-------|-----|
| maaf         | 1.000 | 0.944 | 0.971 | 18  |
| makan        | 0.882 | 1.000 | 0.938 | 15  |
| mau          | 1.000 | 1.000 | 1.000 | 12  |
| nama         | 1.000 | 0.833 | 0.909 | 18  |
| pagi         | 1.000 | 1.000 | 1.000 | 15  |
| paham        | 1.000 | 0.952 | 0.976 | 21  |
| sakit        | 1.000 | 1.000 | 1.000 | 2   |
| sama-sama    | 0.958 | 0.958 | 0.958 | 24  |
| saya         | 1.000 | 0.714 | 0.833 | 7   |
| selamat      | 0.935 | 0.967 | 0.951 | 30  |
| siapa        | 0.882 | 0.938 | 0.909 | 16  |
| tanya        | 1.000 | 0.933 | 0.966 | 15  |
| tempat       | 1.000 | 0.800 | 0.889 | 5   |
| terima-kasih | 0.880 | 1.000 | 0.936 | 22  |
| terlambat    | 0.923 | 1.000 | 0.960 | 12  |
| tidak        | 0.900 | 1.000 | 0.947 | 9   |
| tolong       | 1.000 | 1.000 | 1.000 | 13  |
| accuracy     |       |       | 0.680 | 693 |
| macro avg    | 0.676 | 0.662 | 0.657 | 693 |
| weighted avg | 0.682 | 0.680 | 0.672 | 693 |

