

COMPARISON_CNN

June 21, 2025

```
[1]: import tensorflow as tf
from tensorflow.keras.utils import to_categorical
import os
from PIL import Image, UnidentifiedImageError
import shutil

# Configuration
IMG_SIZE = (28, 28)
BATCH_SIZE = 32
VALIDATION_SPLIT = 0.4
SEED = 42
ROOT_PATH = ''
DATASET_PATH = os.path.join(ROOT_PATH, "raw_data")
CORRUPT_PATH = os.path.join(ROOT_PATH, "corrupt_images")
os.makedirs(CORRUPT_PATH, exist_ok=True)

for root, dirs, files in os.walk(DATASET_PATH):
    for file in files:
        ext = os.path.splitext(file)[1].lower()
        if ext in [".jpg", ".jpeg", ".png", ".bmp", ".gif"]:
            path = os.path.join(root, file)
            try:
                with Image.open(path) as img:
                    img.verify() # Check integrity
            except (UnidentifiedImageError, OSError, IOError) as e:
                # Move the corrupt image
                print(f"Corrupt image found: {path} - moving to {CORRUPT_PATH}")
                dest_path = os.path.join(CORRUPT_PATH, os.path.relpath(path,
↳DATASET_PATH))
                os.makedirs(os.path.dirname(dest_path), exist_ok=True)
                shutil.move(path, dest_path)

LANDMARK_DIR = os.path.join(ROOT_PATH, "data")
RAW_IMAGE_DIR = os.path.join(ROOT_PATH, "raw_data")
FILTERED_IMAGE_DIR = os.path.join(ROOT_PATH, "filtered_raw_data")
DATASET_PATH = FILTERED_IMAGE_DIR
# Supported image extensions
```

```

IMAGE_EXTENSIONS = ['.jpg', '.jpeg', '.png', '.bmp']

# Create filtered output structure
os.makedirs(FILTERED_IMAGE_DIR, exist_ok=True)

for class_name in os.listdir(LANDMARK_DIR):
    if class_name == 'debug':
        continue
    landmark_class_dir = os.path.join(LANDMARK_DIR, class_name)
    raw_class_dir = os.path.join(RAW_IMAGE_DIR, class_name)
    filtered_class_dir = os.path.join(FILTERED_IMAGE_DIR, class_name)
    os.makedirs(filtered_class_dir, exist_ok=True)

    for file in os.listdir(landmark_class_dir):
        if not file.endswith("_landmarks.json"):
            continue

        # Get base filename without "_landmarks.json"
        base_name = file.replace("_landmarks.json", "")

        # Look for corresponding image in raw directory
        for ext in IMAGE_EXTENSIONS:
            image_file = os.path.join(raw_class_dir, base_name + ext)
            if os.path.exists(image_file):
                # Copy to filtered folder
                shutil.copy(image_file, os.path.join(filtered_class_dir, os.
↳ path.basename(image_file)))
                break

# Load training dataset with validation split
train_ds = tf.keras.preprocessing.image_dataset_from_directory(
    DATASET_PATH,
    validation_split=VALIDATION_SPLIT,
    subset="training",
    seed=SEED,
    color_mode="grayscale",
    image_size=IMG_SIZE,
    batch_size=BATCH_SIZE
)
num_classes = len(train_ds.class_names)
label_map = train_ds.class_names

val_ds = tf.keras.preprocessing.image_dataset_from_directory(
    DATASET_PATH,
    validation_split=VALIDATION_SPLIT,
    subset="validation",
    seed=SEED,

```

```

        color_mode="grayscale",
        image_size=IMG_SIZE,
        batch_size=BATCH_SIZE
    )

    test_ds = val_ds.shard(2,0)
    val_ds = val_ds.shard(2,1)
    # Normalize pixel values to [0, 1]
    normalization_layer = tf.keras.layers.Rescaling(1./255)
    train_ds = train_ds.map(lambda x, y: (normalization_layer(x), y))
    val_ds = val_ds.map(lambda x, y: (normalization_layer(x), y))
    test_ds = test_ds.map(lambda x, y: (normalization_layer(x), y))
    # Cache and prefetch for performance
    AUTOTUNE = tf.data.AUTOTUNE
    train_ds = train_ds.cache().shuffle(1000).prefetch(buffer_size=AUTOTUNE)
    val_ds = val_ds.cache().prefetch(buffer_size=AUTOTUNE)
    test_ds = test_ds.cache().prefetch(buffer_size=AUTOTUNE)

```

Found 5643 files belonging to 51 classes.

Using 3386 files for training.

Found 5643 files belonging to 51 classes.

Using 2257 files for validation.

```

[2]: from tensorflow.keras.models import Sequential
    from tensorflow.keras.layers import Conv2D, MaxPooling2D, Dropout, BatchNormalization, Input
    from tensorflow.keras.layers import Flatten, Dense, GlobalAveragePooling2D
    from tensorflow.keras.optimizers import Adam

    model = Sequential([
        Input((28, 28, 1)),
        Conv2D(16, (3, 3), activation='relu'),
        BatchNormalization(),
        MaxPooling2D(pool_size=(2, 2)),
        Dropout(0.1),

        Conv2D(32, (3, 3), activation='relu'),
        BatchNormalization(),
        MaxPooling2D(pool_size=(2, 2)),
        Dropout(0.2),

        GlobalAveragePooling2D(),
        Flatten(),

        Dense(128, activation='relu'),
        Dropout(0.2),

        Dense(num_classes, activation='softmax')
    ])

```

```

])

model.compile(optimizer=Adam(1e-3),
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

```

```
[3]: history = model.fit(train_ds, validation_data=val_ds, epochs=50)
```

```

Epoch 1/50
106/106          19s 67ms/step -
accuracy: 0.0285 - loss: 3.9030 - val_accuracy: 0.0188 - val_loss: 3.9602
Epoch 2/50
106/106          1s 7ms/step -
accuracy: 0.0746 - loss: 3.4015 - val_accuracy: 0.0205 - val_loss: 4.2724
Epoch 3/50
106/106          1s 7ms/step -
accuracy: 0.0855 - loss: 3.2238 - val_accuracy: 0.0232 - val_loss: 4.3350
Epoch 4/50
106/106          1s 7ms/step -
accuracy: 0.0961 - loss: 3.1364 - val_accuracy: 0.0259 - val_loss: 3.9981
Epoch 5/50
106/106          1s 7ms/step -
accuracy: 0.1320 - loss: 3.0458 - val_accuracy: 0.0473 - val_loss: 3.6321
Epoch 6/50
106/106          1s 7ms/step -
accuracy: 0.1539 - loss: 2.9616 - val_accuracy: 0.1437 - val_loss: 3.1492
Epoch 7/50
106/106          1s 7ms/step -
accuracy: 0.1671 - loss: 2.8833 - val_accuracy: 0.1375 - val_loss: 3.1823
Epoch 8/50
106/106          1s 7ms/step -
accuracy: 0.2108 - loss: 2.8134 - val_accuracy: 0.1009 - val_loss: 3.2990
Epoch 9/50
106/106          1s 7ms/step -
accuracy: 0.2060 - loss: 2.7633 - val_accuracy: 0.1411 - val_loss: 3.0562
Epoch 10/50
106/106          1s 7ms/step -
accuracy: 0.2275 - loss: 2.6895 - val_accuracy: 0.1857 - val_loss: 2.7917
Epoch 11/50
106/106          1s 7ms/step -
accuracy: 0.2492 - loss: 2.6086 - val_accuracy: 0.2232 - val_loss: 2.7362
Epoch 12/50
106/106          1s 7ms/step -
accuracy: 0.2650 - loss: 2.5422 - val_accuracy: 0.1804 - val_loss: 2.8804
Epoch 13/50
106/106          1s 7ms/step -
accuracy: 0.2629 - loss: 2.5373 - val_accuracy: 0.2464 - val_loss: 2.6650
Epoch 14/50

```

106/106 1s 7ms/step -
 accuracy: 0.2880 - loss: 2.4425 - val_accuracy: 0.2482 - val_loss: 2.6002
 Epoch 15/50
 106/106 1s 7ms/step -
 accuracy: 0.2977 - loss: 2.4105 - val_accuracy: 0.2688 - val_loss: 2.6124
 Epoch 16/50
 106/106 1s 7ms/step -
 accuracy: 0.3190 - loss: 2.3837 - val_accuracy: 0.2438 - val_loss: 2.6162
 Epoch 17/50
 106/106 1s 7ms/step -
 accuracy: 0.3104 - loss: 2.3365 - val_accuracy: 0.2812 - val_loss: 2.4488
 Epoch 18/50
 106/106 1s 7ms/step -
 accuracy: 0.3286 - loss: 2.2932 - val_accuracy: 0.2759 - val_loss: 2.4608
 Epoch 19/50
 106/106 1s 7ms/step -
 accuracy: 0.3545 - loss: 2.2151 - val_accuracy: 0.1884 - val_loss: 2.8882
 Epoch 20/50
 106/106 1s 7ms/step -
 accuracy: 0.3476 - loss: 2.2272 - val_accuracy: 0.3071 - val_loss: 2.3932
 Epoch 21/50
 106/106 1s 7ms/step -
 accuracy: 0.3327 - loss: 2.2491 - val_accuracy: 0.1875 - val_loss: 2.8813
 Epoch 22/50
 106/106 1s 7ms/step -
 accuracy: 0.3645 - loss: 2.1698 - val_accuracy: 0.2536 - val_loss: 2.5397
 Epoch 23/50
 106/106 1s 7ms/step -
 accuracy: 0.3512 - loss: 2.1625 - val_accuracy: 0.2929 - val_loss: 2.4505
 Epoch 24/50
 106/106 1s 7ms/step -
 accuracy: 0.3633 - loss: 2.1475 - val_accuracy: 0.3170 - val_loss: 2.3849
 Epoch 25/50
 106/106 1s 7ms/step -
 accuracy: 0.3645 - loss: 2.1227 - val_accuracy: 0.1929 - val_loss: 2.8297
 Epoch 26/50
 106/106 1s 7ms/step -
 accuracy: 0.3727 - loss: 2.1300 - val_accuracy: 0.2536 - val_loss: 2.6489
 Epoch 27/50
 106/106 1s 7ms/step -
 accuracy: 0.3834 - loss: 2.1089 - val_accuracy: 0.2750 - val_loss: 2.4914
 Epoch 28/50
 106/106 1s 7ms/step -
 accuracy: 0.3833 - loss: 2.0938 - val_accuracy: 0.2812 - val_loss: 2.5172
 Epoch 29/50
 106/106 1s 7ms/step -
 accuracy: 0.3908 - loss: 2.0487 - val_accuracy: 0.3357 - val_loss: 2.2638
 Epoch 30/50

106/106 1s 7ms/step -
 accuracy: 0.4000 - loss: 2.0492 - val_accuracy: 0.3071 - val_loss: 2.4016
 Epoch 31/50
 106/106 1s 7ms/step -
 accuracy: 0.3992 - loss: 2.0285 - val_accuracy: 0.2973 - val_loss: 2.4034
 Epoch 32/50
 106/106 1s 7ms/step -
 accuracy: 0.3928 - loss: 1.9994 - val_accuracy: 0.2875 - val_loss: 2.4529
 Epoch 33/50
 106/106 1s 7ms/step -
 accuracy: 0.4004 - loss: 2.0200 - val_accuracy: 0.3366 - val_loss: 2.2722
 Epoch 34/50
 106/106 1s 7ms/step -
 accuracy: 0.3933 - loss: 2.0160 - val_accuracy: 0.3134 - val_loss: 2.3764
 Epoch 35/50
 106/106 1s 7ms/step -
 accuracy: 0.4115 - loss: 1.9637 - val_accuracy: 0.2937 - val_loss: 2.4489
 Epoch 36/50
 106/106 1s 7ms/step -
 accuracy: 0.4028 - loss: 1.9667 - val_accuracy: 0.3384 - val_loss: 2.2480
 Epoch 37/50
 106/106 1s 7ms/step -
 accuracy: 0.3932 - loss: 1.9684 - val_accuracy: 0.3170 - val_loss: 2.3287
 Epoch 38/50
 106/106 1s 7ms/step -
 accuracy: 0.4251 - loss: 1.9092 - val_accuracy: 0.2625 - val_loss: 2.6642
 Epoch 39/50
 106/106 1s 8ms/step -
 accuracy: 0.4244 - loss: 1.9183 - val_accuracy: 0.2589 - val_loss: 2.5868
 Epoch 40/50
 106/106 1s 8ms/step -
 accuracy: 0.4159 - loss: 1.9364 - val_accuracy: 0.3241 - val_loss: 2.2809
 Epoch 41/50
 106/106 1s 8ms/step -
 accuracy: 0.4319 - loss: 1.8823 - val_accuracy: 0.3446 - val_loss: 2.2653
 Epoch 42/50
 106/106 1s 8ms/step -
 accuracy: 0.4493 - loss: 1.8265 - val_accuracy: 0.3313 - val_loss: 2.2376
 Epoch 43/50
 106/106 1s 8ms/step -
 accuracy: 0.4392 - loss: 1.8844 - val_accuracy: 0.3509 - val_loss: 2.2370
 Epoch 44/50
 106/106 1s 7ms/step -
 accuracy: 0.4439 - loss: 1.8559 - val_accuracy: 0.3634 - val_loss: 2.2481
 Epoch 45/50
 106/106 1s 7ms/step -
 accuracy: 0.4469 - loss: 1.8406 - val_accuracy: 0.3482 - val_loss: 2.2282
 Epoch 46/50

```

106/106          1s 7ms/step -
accuracy: 0.4548 - loss: 1.8436 - val_accuracy: 0.3580 - val_loss: 2.2125
Epoch 47/50
106/106          1s 7ms/step -
accuracy: 0.4498 - loss: 1.8072 - val_accuracy: 0.3036 - val_loss: 2.4754
Epoch 48/50
106/106          1s 7ms/step -
accuracy: 0.4469 - loss: 1.8241 - val_accuracy: 0.3393 - val_loss: 2.2447
Epoch 49/50
106/106          1s 7ms/step -
accuracy: 0.4348 - loss: 1.8508 - val_accuracy: 0.2866 - val_loss: 2.5828
Epoch 50/50
106/106          1s 7ms/step -
accuracy: 0.4575 - loss: 1.7862 - val_accuracy: 0.3625 - val_loss: 2.2806

```

```

[4]: test_loss, test_accuracy = model.evaluate(test_ds)
      print(f"Test Accuracy: {test_accuracy:.4f}")
      print(f"Test Loss: {test_loss:.4f}")

```

```

36/36          5s 136ms/step -
accuracy: 0.3699 - loss: 2.2056
Test Accuracy: 0.3817
Test Loss: 2.2056

```

```


[5]: import matplotlib.pyplot as plt
      from sklearn.metrics import classification_report, confusion_matrix
      import seaborn as sns
      import numpy as np

```

```

[6]: plt.figure(figsize=(12, 4))
      plt.subplot(1, 2, 1)
      plt.plot(history.history['accuracy'])
      plt.plot(history.history['val_accuracy'])
      plt.title('Model accuracy')
      plt.ylabel('Accuracy')
      plt.xlabel('Epoch')
      plt.legend(['Train', 'Validation'], loc='upper left')
      # Plot training & validation loss values
      plt.subplot(1, 2, 2)
      plt.plot(history.history['loss'])
      plt.plot(history.history['val_loss'])
      plt.title('Model loss')
      plt.ylabel('Loss')
      plt.xlabel('Epoch')
      plt.legend(['Train', 'Validation'], loc='upper left')
      plt.show()

```



```
[7]: y_true, y_pred = [], []
target_names = [label_map[i] for i in range(len(label_map))]
for X_batch, y_batch in test_ds:
    y_true.append(y_batch.numpy())

    batch_pred = model.predict(X_batch, verbose=0)
    y_pred.append(np.argmax(batch_pred, axis=1))

y_true = np.concatenate(y_true)
y_pred = np.concatenate(y_pred)

print(classification_report(
    y_true, y_pred,
    digits=3,
    target_names=target_names
))

cm = confusion_matrix(y_true, y_pred, labels=range(len(label_map)))
labels = [label_map[i] for i in range(len(label_map))]

plt.figure(figsize=(10, 8))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues",
            xticklabels=labels, yticklabels=labels)
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.title("Confusion Matrix - Test Set")
plt.show()
```

c:\Users\chris\.conda\envs\ASLR\Lib\site-packages\sklearn\metrics_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.


```
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
c:\Users\chris\.conda\envs\ASLR\Lib\site-
packages\sklearn\metrics\_classification.py:1565: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
c:\Users\chris\.conda\envs\ASLR\Lib\site-
packages\sklearn\metrics\_classification.py:1565: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

	precision	recall	f1-score	support
A	0.083	0.053	0.065	19
B	0.091	0.037	0.053	27
C	0.050	0.043	0.047	23
D	0.000	0.000	0.000	21
E	0.361	0.433	0.394	30
F	0.188	0.120	0.146	25
G	0.105	0.129	0.116	31
H	0.286	0.138	0.186	29
I	0.083	0.227	0.122	22
J	0.162	0.250	0.197	24
K	0.000	0.000	0.000	32
L	0.116	0.192	0.145	26
M	0.333	0.116	0.172	43
N	0.067	0.071	0.069	28
O	0.048	0.038	0.043	26
P	0.000	0.000	0.000	24
Q	0.122	0.278	0.169	18
R	0.062	0.095	0.075	21
S	0.200	0.094	0.128	32
T	0.115	0.136	0.125	22
U	0.238	0.152	0.185	33
V	0.129	0.267	0.174	30
W	0.206	0.269	0.233	26
X	0.000	0.000	0.000	23
Y	0.000	0.000	0.000	29
Z	0.260	0.679	0.376	28
baca	0.750	0.273	0.400	11
bantu	1.000	0.400	0.571	15
bapak	0.810	0.895	0.850	19
buangairkecil	1.000	0.364	0.533	11
buat	0.556	0.833	0.667	12
halo	1.000	0.667	0.800	15
ibu	0.800	0.667	0.727	6
kamu	0.605	0.719	0.657	32

maaf	0.850	0.773	0.810	22
makan	0.875	0.350	0.500	20
mau	0.815	0.957	0.880	23
nama	0.944	0.739	0.829	23
pagi	0.800	0.909	0.851	22
paham	0.826	0.594	0.691	32
sakit	1.000	0.571	0.727	7
sama-sama	0.682	0.750	0.714	20
saya	0.643	0.818	0.720	11
selamat	0.733	0.917	0.815	24
siapa	0.533	0.727	0.615	22
tanya	0.913	0.875	0.894	24
tempat	0.875	0.700	0.778	10
terima-kasih	0.432	0.941	0.593	17
terlambat	0.640	0.941	0.762	17
tidak	0.667	0.600	0.632	10
tolong	0.833	1.000	0.909	20
accuracy			0.382	1137
macro avg	0.449	0.427	0.415	1137
weighted avg	0.388	0.382	0.365	1137

