

# COMPARISON\_MediaPipe+CNN+LSTM

June 21, 2025

```
[1]: from modules.SignLanguageProcessor import load_and_preprocess_data, parse_frame
import os
```

```
[2]: ROOT_PATH = ''
sequences, labels, label_map = load_and_preprocess_data(os.path.
    ↪join(ROOT_PATH, 'data'))
```

```
[3]: num_classes = len(label_map)
```

```
[4]: len(labels)
```

```
[4]: 5643
```

```
[5]: sequences.shape
```

```
[5]: (5643, 3, 61, 3)
```

```
[6]: from sklearn.model_selection import train_test_split

X_train, X_temp, y_train, y_temp = train_test_split(
    sequences, labels, test_size=0.4, stratify=labels, random_state=42
)

X_val, X_test, y_val, y_test = train_test_split(
    X_temp, y_temp, test_size=0.5, stratify=y_temp, random_state=42
)
```

```
[7]: import numpy as np
def normalize_landmark_data(X):
    """
    Normalize the landmark features (x, y) to have zero mean and unit variance_
    ↪across the training set.
    Assumes X shape is (N, F, L, T), where F=3 (x, y, vis).
    """
    X = X.copy()
    # Flatten across all samples, landmarks, and frames
    x_vals = X[:, 0, :, :].flatten()
    y_vals = X[:, 1, :, :].flatten()
```

```

    # Compute mean and std
    x_mean, x_std = np.mean(x_vals), np.std(x_vals)
    y_mean, y_std = np.mean(y_vals), np.std(y_vals)

    # Normalize
    X[:, 0, :, :] = (X[:, 0, :, :] - x_mean) / x_std
    X[:, 1, :, :] = (X[:, 1, :, :] - y_mean) / y_std

    return X, (x_mean, x_std), (y_mean, y_std)

def apply_normalization(X, x_mean, x_std, y_mean, y_std):
    X = X.copy()
    X[:, 0, :, :] = (X[:, 0, :, :] - x_mean) / x_std
    X[:, 1, :, :] = (X[:, 1, :, :] - y_mean) / y_std
    return X

```

```

[8]: def reshape_frames_for_cnn(X, y):
      X = X.transpose(0, 3, 2, 1) # (N, T, L, F)
      X = X[..., np.newaxis]      # (N, T, L, F, 1)
      return X,y

```

```

[9]: X_train_norm, (x_mean, x_std), (y_mean, y_std) = ↵
      ↪normalize_landmark_data(X_train)
      X_val_norm = apply_normalization(X_val, x_mean, x_std, y_mean, y_std)
      X_test_norm = apply_normalization(X_test, x_mean, x_std, y_mean, y_std)

      X_train_cnn, y_train_cnn = reshape_frames_for_cnn(X_train_norm, y_train)
      X_val_cnn, y_val_cnn = reshape_frames_for_cnn(X_val_norm, y_val)
      X_test_cnn, y_test_cnn = reshape_frames_for_cnn(X_test_norm, y_test)

      print(X_train_cnn.shape)
      print(y_train_cnn.shape)

```

```

(3385, 3, 61, 3, 1)
(3385,)

```

```

[10]: input_shape = X_train_cnn.shape[1:]
      print(input_shape)

```

```

(3, 61, 3, 1)

```

```

[11]: import tensorflow as tf

      train_ds = tf.data.Dataset.from_tensor_slices((X_train_cnn, y_train_cnn))
      train_ds = train_ds.shuffle(buffer_size=1000).batch(64).prefetch(tf.data.
      ↪AUTOTUNE)

```

```

val_ds = tf.data.Dataset.from_tensor_slices((X_val_cnn, y_val_cnn))
val_ds = val_ds.batch(64).prefetch(tf.data.AUTOTUNE)

test_ds = tf.data.Dataset.from_tensor_slices((X_test_cnn, y_test_cnn))
test_ds = test_ds.batch(64).prefetch(tf.data.AUTOTUNE)

```

```

[19]: from tensorflow.keras.models import Sequential
      from tensorflow.keras.layers import TimeDistributed, Conv2D, MaxPooling2D, Flatten, Input
      from tensorflow.keras.layers import LSTM, Dropout, Dense, BatchNormalization

model = Sequential([
    Input((3, 61, 3, 1)),
    TimeDistributed(Conv2D(32, (3, 2), activation='relu', padding='same')),
    TimeDistributed(BatchNormalization()),
    TimeDistributed(MaxPooling2D(pool_size=(2, 1))),
    TimeDistributed(Dropout(0.25)),

    TimeDistributed(Conv2D(64, (3, 2), activation='relu', padding='same')),
    TimeDistributed(BatchNormalization()),
    TimeDistributed(MaxPooling2D(pool_size=(2, 1))),
    TimeDistributed(Flatten()),

    LSTM(128, return_sequences=False),
    Dropout(0.5),
    Dense(num_classes, activation='softmax')
])

model.compile(optimizer='adam', loss='sparse_categorical_crossentropy',
             metrics=['accuracy'])

```

```

[20]: history = model.fit(train_ds, validation_data=val_ds, epochs=50, batch_size=64)

```

```

Epoch 1/50
53/53          7s 60ms/step -
accuracy: 0.0507 - loss: 3.6923 - val_accuracy: 0.0806 - val_loss: 3.4596
Epoch 2/50
53/53          3s 51ms/step -
accuracy: 0.1152 - loss: 3.2279 - val_accuracy: 0.0912 - val_loss: 3.1668
Epoch 3/50
53/53          3s 48ms/step -
accuracy: 0.1367 - loss: 3.0772 - val_accuracy: 0.1293 - val_loss: 3.0201
Epoch 4/50
53/53          3s 51ms/step -
accuracy: 0.1592 - loss: 2.9782 - val_accuracy: 0.1453 - val_loss: 2.9099
Epoch 5/50
53/53          3s 52ms/step -
accuracy: 0.1810 - loss: 2.8905 - val_accuracy: 0.2002 - val_loss: 2.8233

```

Epoch 6/50  
53/53 3s 48ms/step -  
accuracy: 0.2183 - loss: 2.7815 - val\_accuracy: 0.2117 - val\_loss: 2.7393  
Epoch 7/50  
53/53 3s 47ms/step -  
accuracy: 0.2265 - loss: 2.7112 - val\_accuracy: 0.2427 - val\_loss: 2.6644  
Epoch 8/50  
53/53 3s 54ms/step -  
accuracy: 0.2851 - loss: 2.5682 - val\_accuracy: 0.2799 - val\_loss: 2.6393  
Epoch 9/50  
53/53 3s 51ms/step -  
accuracy: 0.3132 - loss: 2.4698 - val\_accuracy: 0.3003 - val\_loss: 2.5923  
Epoch 10/50  
53/53 3s 51ms/step -  
accuracy: 0.3256 - loss: 2.4252 - val\_accuracy: 0.3020 - val\_loss: 2.4991  
Epoch 11/50  
53/53 3s 50ms/step -  
accuracy: 0.3545 - loss: 2.3281 - val\_accuracy: 0.3029 - val\_loss: 2.4149  
Epoch 12/50  
53/53 3s 54ms/step -  
accuracy: 0.3917 - loss: 2.1795 - val\_accuracy: 0.3339 - val\_loss: 2.3556  
Epoch 13/50  
53/53 3s 49ms/step -  
accuracy: 0.4134 - loss: 2.0929 - val\_accuracy: 0.4039 - val\_loss: 2.2474  
Epoch 14/50  
53/53 3s 51ms/step -  
accuracy: 0.4289 - loss: 2.0593 - val\_accuracy: 0.4172 - val\_loss: 2.1792  
Epoch 15/50  
53/53 3s 52ms/step -  
accuracy: 0.4478 - loss: 1.9414 - val\_accuracy: 0.4198 - val\_loss: 2.1114  
Epoch 16/50  
53/53 3s 54ms/step -  
accuracy: 0.4768 - loss: 1.8474 - val\_accuracy: 0.4057 - val\_loss: 2.1733  
Epoch 17/50  
53/53 3s 55ms/step -  
accuracy: 0.4839 - loss: 1.8162 - val\_accuracy: 0.4756 - val\_loss: 1.9977  
Epoch 18/50  
53/53 3s 57ms/step -  
accuracy: 0.5030 - loss: 1.7098 - val\_accuracy: 0.4774 - val\_loss: 1.9651  
Epoch 19/50  
53/53 3s 53ms/step -  
accuracy: 0.5434 - loss: 1.5988 - val\_accuracy: 0.5102 - val\_loss: 1.8605  
Epoch 20/50  
53/53 3s 61ms/step -  
accuracy: 0.5671 - loss: 1.5555 - val\_accuracy: 0.5261 - val\_loss: 1.8027  
Epoch 21/50  
53/53 3s 61ms/step -  
accuracy: 0.5879 - loss: 1.4889 - val\_accuracy: 0.5536 - val\_loss: 1.6867

Epoch 22/50  
53/53 3s 54ms/step -  
accuracy: 0.5861 - loss: 1.4140 - val\_accuracy: 0.5757 - val\_loss: 1.6500

Epoch 23/50  
53/53 3s 56ms/step -  
accuracy: 0.6328 - loss: 1.3359 - val\_accuracy: 0.6058 - val\_loss: 1.5774

Epoch 24/50  
53/53 4s 74ms/step -  
accuracy: 0.6420 - loss: 1.3187 - val\_accuracy: 0.5421 - val\_loss: 1.6943

Epoch 25/50  
53/53 4s 78ms/step -  
accuracy: 0.6433 - loss: 1.2622 - val\_accuracy: 0.5270 - val\_loss: 1.6817

Epoch 26/50  
53/53 4s 79ms/step -  
accuracy: 0.6593 - loss: 1.2080 - val\_accuracy: 0.5226 - val\_loss: 1.6613

Epoch 27/50  
53/53 4s 78ms/step -  
accuracy: 0.6716 - loss: 1.1873 - val\_accuracy: 0.6058 - val\_loss: 1.5274

Epoch 28/50  
53/53 4s 71ms/step -  
accuracy: 0.6944 - loss: 1.1506 - val\_accuracy: 0.5757 - val\_loss: 1.5156

Epoch 29/50  
53/53 3s 60ms/step -  
accuracy: 0.7052 - loss: 1.1038 - val\_accuracy: 0.5864 - val\_loss: 1.5028

Epoch 30/50  
53/53 3s 60ms/step -  
accuracy: 0.7004 - loss: 1.0738 - val\_accuracy: 0.6050 - val\_loss: 1.4674

Epoch 31/50  
53/53 3s 58ms/step -  
accuracy: 0.7274 - loss: 0.9806 - val\_accuracy: 0.6023 - val\_loss: 1.4380

Epoch 32/50  
53/53 3s 47ms/step -  
accuracy: 0.7023 - loss: 1.0166 - val\_accuracy: 0.5846 - val\_loss: 1.4293

Epoch 33/50  
53/53 3s 48ms/step -  
accuracy: 0.7274 - loss: 0.9633 - val\_accuracy: 0.6484 - val\_loss: 1.3223

Epoch 34/50  
53/53 3s 52ms/step -  
accuracy: 0.7472 - loss: 0.9066 - val\_accuracy: 0.6457 - val\_loss: 1.3168

Epoch 35/50  
53/53 3s 58ms/step -  
accuracy: 0.7527 - loss: 0.8753 - val\_accuracy: 0.6572 - val\_loss: 1.2775

Epoch 36/50  
53/53 3s 60ms/step -  
accuracy: 0.7483 - loss: 0.8848 - val\_accuracy: 0.6900 - val\_loss: 1.2040

Epoch 37/50  
53/53 4s 69ms/step -  
accuracy: 0.7688 - loss: 0.8282 - val\_accuracy: 0.6962 - val\_loss: 1.1737

```

Epoch 38/50
53/53          3s 60ms/step -
accuracy: 0.7735 - loss: 0.8184 - val_accuracy: 0.6811 - val_loss: 1.1797
Epoch 39/50
53/53          3s 61ms/step -
accuracy: 0.7750 - loss: 0.7760 - val_accuracy: 0.6829 - val_loss: 1.1854
Epoch 40/50
53/53          3s 58ms/step -
accuracy: 0.7845 - loss: 0.7461 - val_accuracy: 0.6643 - val_loss: 1.2882
Epoch 41/50
53/53          3s 59ms/step -
accuracy: 0.7848 - loss: 0.7711 - val_accuracy: 0.6528 - val_loss: 1.3329
Epoch 42/50
53/53          3s 61ms/step -
accuracy: 0.7865 - loss: 0.7324 - val_accuracy: 0.6696 - val_loss: 1.2521
Epoch 43/50
53/53          3s 60ms/step -
accuracy: 0.7881 - loss: 0.7324 - val_accuracy: 0.6501 - val_loss: 1.2915
Epoch 44/50
53/53          3s 58ms/step -
accuracy: 0.8081 - loss: 0.6850 - val_accuracy: 0.6776 - val_loss: 1.1597
Epoch 45/50
53/53          3s 55ms/step -
accuracy: 0.8023 - loss: 0.6634 - val_accuracy: 0.6980 - val_loss: 1.1118
Epoch 46/50
53/53          3s 63ms/step -
accuracy: 0.7997 - loss: 0.6798 - val_accuracy: 0.6599 - val_loss: 1.1908
Epoch 47/50
53/53          3s 57ms/step -
accuracy: 0.8149 - loss: 0.6753 - val_accuracy: 0.6634 - val_loss: 1.1857
Epoch 48/50
53/53          3s 57ms/step -
accuracy: 0.8166 - loss: 0.6374 - val_accuracy: 0.6475 - val_loss: 1.2377
Epoch 49/50
53/53          3s 51ms/step -
accuracy: 0.8062 - loss: 0.6589 - val_accuracy: 0.6492 - val_loss: 1.2095
Epoch 50/50
53/53          3s 56ms/step -
accuracy: 0.8419 - loss: 0.5885 - val_accuracy: 0.6767 - val_loss: 1.1418

```

```

[21]: test_loss, test_accuracy = model.evaluate(test_ds)
      print(f"Test Accuracy: {test_accuracy:.4f}")
      print(f"Test Loss: {test_loss:.4f}")

```

```

1/18          0s 17ms/step -
accuracy: 0.6094 - loss: 1.1868


18/18         0s 9ms/step -
accuracy: 0.6504 - loss: 1.1274

```

Test Accuracy: 0.6723  
Test Loss: 1.0921

```
[22]: import matplotlib.pyplot as plt
from sklearn.metrics import classification_report, confusion_matrix
import seaborn as sns
```

```
[23]: plt.figure(figsize=(12, 4))
plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')
# Plot training & validation loss values
plt.subplot(1, 2, 2)
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()
```



```
[24]: y_true, y_pred = [], []
target_names = [label_map[i] for i in range(len(label_map))]
for X_batch, y_batch in test_ds:
    y_true.append(y_batch.numpy())

    batch_pred = model.predict(X_batch, verbose=0)
    y_pred.append(np.argmax(batch_pred, axis=1))
```

```

y_true = np.concatenate(y_true)
y_pred = np.concatenate(y_pred)

print(classification_report(
    y_true, y_pred,
    digits=3,
    target_names=target_names
))

cm = confusion_matrix(y_true, y_pred, labels=range(len(label_map)))
labels = [label_map[i] for i in range(len(label_map))]

plt.figure(figsize=(10, 8))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues",
            xticklabels=labels, yticklabels=labels)
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.title("Confusion Matrix - Test Set")
plt.show()

```

```

c:\Users\chris\.conda\envs\ASLR\Lib\site-
packages\sklearn\metrics\_classification.py:1565: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
c:\Users\chris\.conda\envs\ASLR\Lib\site-
packages\sklearn\metrics\_classification.py:1565: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
c:\Users\chris\.conda\envs\ASLR\Lib\site-
packages\sklearn\metrics\_classification.py:1565: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))

```

	precision	recall	f1-score	support
A	0.591	0.650	0.619	20
B	0.458	0.478	0.468	23
C	0.750	0.444	0.558	27
D	0.409	0.429	0.419	21
E	0.783	0.692	0.735	26
F	0.773	0.708	0.739	24
G	0.492	1.000	0.659	29
H	0.654	0.680	0.667	25
I	0.818	0.900	0.857	30



J	0.957	0.733	0.830	30
K	0.519	0.538	0.528	26
L	1.000	0.296	0.457	27
M	0.426	0.935	0.586	31
N	0.000	0.000	0.000	32
O	1.000	0.032	0.062	31
P	0.406	0.520	0.456	25
Q	0.434	0.742	0.548	31
R	1.000	0.231	0.375	26
S	0.792	0.633	0.704	30
T	0.394	0.500	0.441	26
U	0.909	0.714	0.800	28
V	0.913	0.808	0.857	26
W	0.667	0.519	0.583	27
X	0.371	0.591	0.456	22
Y	0.290	0.692	0.409	26
Z	0.826	0.679	0.745	28
baca	1.000	0.562	0.720	16
bantu	1.000	0.857	0.923	14
bapak	0.846	0.733	0.786	15
buangairkecil	1.000	1.000	1.000	8
buat	1.000	0.938	0.968	16
halo	0.818	0.900	0.857	20
ibu	1.000	0.667	0.800	6
kamu	0.704	0.864	0.776	22
maaf	0.913	1.000	0.955	21
makan	0.769	0.588	0.667	17
mau	1.000	1.000	1.000	20
nama	0.697	0.920	0.793	25
pagi	0.950	0.792	0.864	24
paham	0.821	0.920	0.868	25
sakit	1.000	0.800	0.889	5
sama-sama	0.867	0.929	0.897	28
saya	0.500	0.769	0.606	13
selamat	0.882	0.714	0.789	21
siapa	0.769	0.625	0.690	16
tanya	0.824	0.700	0.757	20
tempat	0.778	0.875	0.824	8
terima-kasih	0.783	0.900	0.837	20
terlambat	0.895	1.000	0.944	17
tidak	1.000	0.588	0.741	17
tolong	0.833	0.556	0.667	18
accuracy			0.672	1129
macro avg	0.751	0.693	0.690	1129
weighted avg	0.725	0.672	0.660	1129

