

COMPARISON_MobileNetV2

June 21, 2025

```
[1]: import tensorflow as tf
from tensorflow.keras.utils import to_categorical
import os
from PIL import Image, UnidentifiedImageError
import shutil

# Configuration
IMG_SIZE = (96, 96)
BATCH_SIZE = 32
VALIDATION_SPLIT = 0.4
SEED = 42
ROOT_PATH = ''
DATASET_PATH = os.path.join(ROOT_PATH, "raw_data")
CORRUPT_PATH = os.path.join(ROOT_PATH, "corrupt_images")
os.makedirs(CORRUPT_PATH, exist_ok=True)

for root, dirs, files in os.walk(DATASET_PATH):
    for file in files:
        ext = os.path.splitext(file)[1].lower()
        if ext in [".jpg", ".jpeg", ".png", ".bmp", ".gif"]:
            path = os.path.join(root, file)
            try:
                with Image.open(path) as img:
                    img.verify() # Check integrity
            except (UnidentifiedImageError, OSError, IOError) as e:
                # Move the corrupt image
                print(f"Corrupt image found: {path} - moving to {CORRUPT_PATH}")
                dest_path = os.path.join(CORRUPT_PATH, os.path.relpath(path,
↳DATASET_PATH))
                os.makedirs(os.path.dirname(dest_path), exist_ok=True)
                shutil.move(path, dest_path)

LANDMARK_DIR = os.path.join(ROOT_PATH, "data")
RAW_IMAGE_DIR = os.path.join(ROOT_PATH, "raw_data")
FILTERED_IMAGE_DIR = os.path.join(ROOT_PATH, "filtered_raw_data")
DATASET_PATH = FILTERED_IMAGE_DIR
# Supported image extensions
```

```

IMAGE_EXTENSIONS = ['.jpg', '.jpeg', '.png', '.bmp']

# Create filtered output structure
os.makedirs(FILTERED_IMAGE_DIR, exist_ok=True)

for class_name in os.listdir(LANDMARK_DIR):
    if class_name == 'debug':
        continue
    landmark_class_dir = os.path.join(LANDMARK_DIR, class_name)
    raw_class_dir = os.path.join(RAW_IMAGE_DIR, class_name)
    filtered_class_dir = os.path.join(FILTERED_IMAGE_DIR, class_name)
    os.makedirs(filtered_class_dir, exist_ok=True)

    for file in os.listdir(landmark_class_dir):
        if not file.endswith("_landmarks.json"):
            continue

        # Get base filename without "_landmarks.json"
        base_name = file.replace("_landmarks.json", "")

        # Look for corresponding image in raw directory
        for ext in IMAGE_EXTENSIONS:
            image_file = os.path.join(raw_class_dir, base_name + ext)
            if os.path.exists(image_file):
                # Copy to filtered folder
                shutil.copy(image_file, os.path.join(filtered_class_dir, os.
↳ path.basename(image_file)))
                break

# Load training dataset with validation split
train_ds = tf.keras.preprocessing.image_dataset_from_directory(
    DATASET_PATH,
    validation_split=VALIDATION_SPLIT,
    subset="training",
    seed=SEED,
    color_mode="rgb",
    image_size=IMG_SIZE,
    batch_size=BATCH_SIZE
)
num_classes = len(train_ds.class_names)
label_map = train_ds.class_names

val_ds = tf.keras.preprocessing.image_dataset_from_directory(
    DATASET_PATH,
    validation_split=VALIDATION_SPLIT,
    subset="validation",
    seed=SEED,

```

```

        color_mode="rgb",
        image_size=IMG_SIZE,
        batch_size=BATCH_SIZE
    )

    test_ds = val_ds.shard(2,0)
    val_ds = val_ds.shard(2,1)
    # Normalize pixel values to [0, 1]
    normalization_layer = tf.keras.layers.Rescaling(1./255)
    train_ds = train_ds.map(lambda x, y: (normalization_layer(x), y))
    val_ds = val_ds.map(lambda x, y: (normalization_layer(x), y))
    test_ds = test_ds.map(lambda x, y: (normalization_layer(x), y))
    # Cache and prefetch for performance
    AUTOTUNE = tf.data.AUTOTUNE
    train_ds = train_ds.cache().shuffle(1000).prefetch(buffer_size=AUTOTUNE)
    val_ds = val_ds.cache().prefetch(buffer_size=AUTOTUNE)
    test_ds = test_ds.cache().prefetch(buffer_size=AUTOTUNE)

```

Found 5643 files belonging to 51 classes.

Using 3386 files for training.

Found 5643 files belonging to 51 classes.

Using 2257 files for validation.

```

[2]: from tensorflow.keras.models import Sequential
    from tensorflow.keras.layers import Conv2D, MaxPooling2D, Dropout,
        BatchNormalization
    from tensorflow.keras.layers import Flatten, Dense, GlobalAveragePooling2D
    from tensorflow.keras.optimizers import Adam
    from tensorflow.keras.applications import MobileNetV2
    base_model = MobileNetV2(input_shape=(96, 96, 3), include_top=False,
        weights='imagenet')
    base_model.trainable = False

    model = Sequential([
        base_model,
        GlobalAveragePooling2D(),
        Dropout(0.3),
        Dense(128, activation='relu'),
        Dropout(0.3),
        Dense(num_classes, activation='softmax')
    ])

    model.compile(optimizer=Adam(1e-3),
        loss='sparse_categorical_crossentropy',
        metrics=['accuracy'])

```

```

[3]: history = model.fit(train_ds, validation_data=val_ds, epochs=50)

```

Epoch 1/50
106/106 31s 145ms/step -
accuracy: 0.0821 - loss: 3.8709 - val_accuracy: 0.3357 - val_loss: 2.4915

Epoch 2/50
106/106 7s 62ms/step -
accuracy: 0.2959 - loss: 2.4460 - val_accuracy: 0.5071 - val_loss: 1.8254

Epoch 3/50
106/106 7s 63ms/step -
accuracy: 0.4433 - loss: 1.9017 - val_accuracy: 0.5750 - val_loss: 1.5179

Epoch 4/50
106/106 8s 74ms/step -
accuracy: 0.5172 - loss: 1.5514 - val_accuracy: 0.6045 - val_loss: 1.3667

Epoch 5/50
106/106 8s 75ms/step -
accuracy: 0.5785 - loss: 1.3306 - val_accuracy: 0.6205 - val_loss: 1.2773

Epoch 6/50
106/106 7s 70ms/step -
accuracy: 0.6324 - loss: 1.1819 - val_accuracy: 0.6143 - val_loss: 1.2361

Epoch 7/50
106/106 7s 69ms/step -
accuracy: 0.6704 - loss: 1.0707 - val_accuracy: 0.6313 - val_loss: 1.1976

Epoch 8/50
106/106 7s 69ms/step -
accuracy: 0.7010 - loss: 0.9367 - val_accuracy: 0.6313 - val_loss: 1.2032

Epoch 9/50
106/106 7s 69ms/step -
accuracy: 0.7164 - loss: 0.8524 - val_accuracy: 0.6375 - val_loss: 1.1630

Epoch 10/50
106/106 7s 68ms/step -
accuracy: 0.7279 - loss: 0.8470 - val_accuracy: 0.6304 - val_loss: 1.1502

Epoch 11/50
106/106 7s 68ms/step -
accuracy: 0.7580 - loss: 0.7494 - val_accuracy: 0.6491 - val_loss: 1.1470

Epoch 12/50
106/106 7s 68ms/step -
accuracy: 0.7682 - loss: 0.6955 - val_accuracy: 0.6429 - val_loss: 1.1309

Epoch 13/50
106/106 7s 70ms/step -
accuracy: 0.7842 - loss: 0.6664 - val_accuracy: 0.6464 - val_loss: 1.1419

Epoch 14/50
106/106 7s 68ms/step -
accuracy: 0.8032 - loss: 0.6217 - val_accuracy: 0.6455 - val_loss: 1.1317

Epoch 15/50
106/106 7s 67ms/step -
accuracy: 0.8145 - loss: 0.5845 - val_accuracy: 0.6321 - val_loss: 1.1303

Epoch 16/50
106/106 7s 69ms/step -
accuracy: 0.8300 - loss: 0.5458 - val_accuracy: 0.6420 - val_loss: 1.1488

Epoch 17/50
106/106 7s 69ms/step -
accuracy: 0.8207 - loss: 0.5450 - val_accuracy: 0.6357 - val_loss: 1.1385

Epoch 18/50
106/106 7s 68ms/step -
accuracy: 0.8334 - loss: 0.5418 - val_accuracy: 0.6339 - val_loss: 1.1635

Epoch 19/50
106/106 7s 68ms/step -
accuracy: 0.8348 - loss: 0.4920 - val_accuracy: 0.6429 - val_loss: 1.1648

Epoch 20/50
106/106 7s 68ms/step -
accuracy: 0.8400 - loss: 0.4762 - val_accuracy: 0.6411 - val_loss: 1.1972

Epoch 21/50
106/106 7s 68ms/step -
accuracy: 0.8398 - loss: 0.4677 - val_accuracy: 0.6455 - val_loss: 1.1394

Epoch 22/50
106/106 7s 68ms/step -
accuracy: 0.8455 - loss: 0.4510 - val_accuracy: 0.6491 - val_loss: 1.1743

Epoch 23/50
106/106 7s 69ms/step -
accuracy: 0.8602 - loss: 0.4289 - val_accuracy: 0.6509 - val_loss: 1.2067

Epoch 24/50
106/106 7s 70ms/step -
accuracy: 0.8722 - loss: 0.3744 - val_accuracy: 0.6446 - val_loss: 1.1939

Epoch 25/50
106/106 7s 70ms/step -
accuracy: 0.8728 - loss: 0.3862 - val_accuracy: 0.6509 - val_loss: 1.2045

Epoch 26/50
106/106 7s 69ms/step -
accuracy: 0.8809 - loss: 0.3705 - val_accuracy: 0.6580 - val_loss: 1.2087

Epoch 27/50
106/106 7s 69ms/step -
accuracy: 0.8793 - loss: 0.3573 - val_accuracy: 0.6491 - val_loss: 1.2118

Epoch 28/50
106/106 7s 69ms/step -
accuracy: 0.8769 - loss: 0.3834 - val_accuracy: 0.6429 - val_loss: 1.2569

Epoch 29/50
106/106 7s 69ms/step -
accuracy: 0.8884 - loss: 0.3356 - val_accuracy: 0.6402 - val_loss: 1.2485

Epoch 30/50
106/106 7s 69ms/step -
accuracy: 0.8828 - loss: 0.3550 - val_accuracy: 0.6554 - val_loss: 1.2400

Epoch 31/50
106/106 7s 69ms/step -
accuracy: 0.9039 - loss: 0.2936 - val_accuracy: 0.6268 - val_loss: 1.2811

Epoch 32/50
106/106 7s 69ms/step -
accuracy: 0.8928 - loss: 0.3340 - val_accuracy: 0.6384 - val_loss: 1.2495

Epoch 33/50
106/106 7s 69ms/step -
accuracy: 0.9081 - loss: 0.2934 - val_accuracy: 0.6384 - val_loss: 1.2868

Epoch 34/50
106/106 7s 69ms/step -
accuracy: 0.8912 - loss: 0.3095 - val_accuracy: 0.6357 - val_loss: 1.2883

Epoch 35/50
106/106 7s 69ms/step -
accuracy: 0.8916 - loss: 0.3126 - val_accuracy: 0.6313 - val_loss: 1.3358

Epoch 36/50
106/106 7s 69ms/step -
accuracy: 0.8985 - loss: 0.3034 - val_accuracy: 0.6348 - val_loss: 1.3427

Epoch 37/50
106/106 7s 69ms/step -
accuracy: 0.9015 - loss: 0.2791 - val_accuracy: 0.6339 - val_loss: 1.3538

Epoch 38/50
106/106 7s 71ms/step -
accuracy: 0.8889 - loss: 0.3075 - val_accuracy: 0.6384 - val_loss: 1.3442

Epoch 39/50
106/106 7s 69ms/step -
accuracy: 0.9178 - loss: 0.2520 - val_accuracy: 0.6321 - val_loss: 1.4135

Epoch 40/50
106/106 7s 69ms/step -
accuracy: 0.9062 - loss: 0.2875 - val_accuracy: 0.6393 - val_loss: 1.3826

Epoch 41/50
106/106 7s 69ms/step -
accuracy: 0.9207 - loss: 0.2517 - val_accuracy: 0.6286 - val_loss: 1.4341

Epoch 42/50
106/106 7s 70ms/step -
accuracy: 0.9042 - loss: 0.2847 - val_accuracy: 0.6411 - val_loss: 1.3822

Epoch 43/50
106/106 8s 76ms/step -
accuracy: 0.9171 - loss: 0.2554 - val_accuracy: 0.6259 - val_loss: 1.4259

Epoch 44/50
106/106 8s 73ms/step -
accuracy: 0.9032 - loss: 0.2738 - val_accuracy: 0.6518 - val_loss: 1.3415

Epoch 45/50
106/106 8s 74ms/step -
accuracy: 0.9165 - loss: 0.2473 - val_accuracy: 0.6509 - val_loss: 1.3767

Epoch 46/50
106/106 8s 76ms/step -
accuracy: 0.9134 - loss: 0.2555 - val_accuracy: 0.6411 - val_loss: 1.4024

Epoch 47/50
106/106 8s 74ms/step -
accuracy: 0.9362 - loss: 0.2144 - val_accuracy: 0.6375 - val_loss: 1.4269

Epoch 48/50
106/106 8s 75ms/step -
accuracy: 0.9284 - loss: 0.2237 - val_accuracy: 0.6339 - val_loss: 1.4466


```
Epoch 49/50
106/106      8s 74ms/step -
accuracy: 0.9265 - loss: 0.2347 - val_accuracy: 0.6357 - val_loss: 1.4796
Epoch 50/50
106/106      8s 73ms/step -
accuracy: 0.9199 - loss: 0.2294 - val_accuracy: 0.6420 - val_loss: 1.4395
```

```
[4]: test_loss, test_accuracy = model.evaluate(test_ds)
      print(f"Test Accuracy: {test_accuracy:.4f}")
      print(f"Test Loss: {test_loss:.4f}")
```

```
36/36      8s 226ms/step -
accuracy: 0.6434 - loss: 1.4337
Test Accuracy: 0.6376
Test Loss: 1.4552
```

```
[5]: import matplotlib.pyplot as plt
      from sklearn.metrics import classification_report, confusion_matrix
      import seaborn as sns
      import numpy as np
```

```
[6]: plt.figure(figsize=(12, 4))
      plt.subplot(1, 2, 1)
      plt.plot(history.history['accuracy'])
      plt.plot(history.history['val_accuracy'])
      plt.title('Model accuracy')
      plt.ylabel('Accuracy')
      plt.xlabel('Epoch')
      plt.legend(['Train', 'Validation'], loc='upper left')
      # Plot training & validation loss values
      plt.subplot(1, 2, 2)
      plt.plot(history.history['loss'])
      plt.plot(history.history['val_loss'])
      plt.title('Model loss')
      plt.ylabel('Loss')
      plt.xlabel('Epoch')
      plt.legend(['Train', 'Validation'], loc='upper left')
      plt.show()
```



```
[7]: y_true, y_pred = [], []
target_names = [label_map[i] for i in range(len(label_map))]
for X_batch, y_batch in test_ds:
    y_true.append(y_batch.numpy())

    batch_pred = model.predict(X_batch, verbose=0)
    y_pred.append(np.argmax(batch_pred, axis=1))

y_true = np.concatenate(y_true)
y_pred = np.concatenate(y_pred)

print(classification_report(
    y_true, y_pred,
    digits=3,
    target_names=target_names
))

cm = confusion_matrix(y_true, y_pred, labels=range(len(label_map)))
labels = [label_map[i] for i in range(len(label_map))]

plt.figure(figsize=(10, 8))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues",
            xticklabels=labels, yticklabels=labels)
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.title("Confusion Matrix - Test Set")
plt.show()
```

	precision	recall	f1-score	support
A	0.667	0.632	0.649	19
B	0.389	0.259	0.311	27

C	0.275	0.609	0.378	23
D	0.257	0.429	0.321	21
E	0.643	0.600	0.621	30
F	0.381	0.320	0.348	25
G	0.585	0.774	0.667	31
H	0.654	0.586	0.618	29
I	0.286	0.273	0.279	22
J	0.400	0.500	0.444	24
K	0.611	0.344	0.440	32
L	0.571	0.615	0.593	26
M	0.355	0.256	0.297	43
N	0.179	0.250	0.209	28
O	0.476	0.385	0.426	26
P	0.333	0.208	0.256	24
Q	0.348	0.444	0.390	18
R	0.471	0.381	0.421	21
S	0.667	0.500	0.571	32
T	0.379	0.500	0.431	22
U	0.286	0.182	0.222	33
V	0.452	0.633	0.528	30
W	0.476	0.385	0.426	26
X	0.500	0.478	0.489	23
Y	0.478	0.379	0.423	29
Z	0.828	0.857	0.842	28
baca	1.000	0.818	0.900	11
bantu	0.933	0.933	0.933	15
bapak	0.950	1.000	0.974	19
buangairkecil	1.000	0.727	0.842	11
buat	0.800	1.000	0.889	12
halo	1.000	1.000	1.000	15
ibu	0.857	1.000	0.923	6
kamu	1.000	0.781	0.877	32
maaf	1.000	1.000	1.000	22
makan	1.000	0.900	0.947	20
mau	1.000	1.000	1.000	23
nama	0.952	0.870	0.909	23
pagi	0.778	0.955	0.857	22
paham	1.000	1.000	1.000	32
sakit	1.000	1.000	1.000	7
sama-sama	0.950	0.950	0.950	20
saya	0.909	0.909	0.909	11
selamat	0.957	0.917	0.936	24
siapa	0.769	0.909	0.833	22
tanya	0.920	0.958	0.939	24
tempat	1.000	0.900	0.947	10
terima-kasih	0.810	1.000	0.895	17
terlambat	0.944	1.000	0.971	17
tidak	0.900	0.900	0.900	10

tolong	0.944	0.850	0.895	20
accuracy			0.638	1137
macro avg	0.693	0.687	0.683	1137
weighted avg	0.650	0.638	0.636	1137

