

TestTrain

May 9, 2025

```
[1]: import torch
import torch.nn as nn
import torch.nn.functional as F
import numpy as np
from torch.utils.data import Dataset, DataLoader
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split
import os
from sklearn.metrics import classification_report, confusion_matrix
import seaborn as sns
import json
import matplotlib.pyplot as plt
import gc
from tqdm import tqdm

# Constants
NUM_POSE_LANDMARKS = 19
NUM_HAND_LANDMARKS = 21
NUM_NODES = NUM_POSE_LANDMARKS + NUM_HAND_LANDMARKS*2 # Total nodes in the
↳graph
FEATURE_DIM = 3 # x, y coordinates + visibility

class GCNLayer(nn.Module):
    def __init__(self, in_features, out_features, dropout=0.5):
        super(GCNLayer, self).__init__()
        # self.linear = nn.Linear(in_features, out_features, bias=True)
        self.conv = nn.Conv1d(in_features, out_features, kernel_size=1,
↳bias=True)
        self.dropout = nn.Dropout(dropout)

    def forward(self, x, adj):
        # x: (batch_size * seq_len, num_nodes, in_features)
        # adj: (num_nodes, num_nodes) - shared across all samples

        # Ensure adj is on the same device as x
        if adj.device != x.device:
            adj = adj.to(x.device)
```

```

        # Expand adj to match batch dimension
        batch_size_seq = x.size(0)
        adj_expanded = adj.unsqueeze(0).expand(batch_size_seq, -1, -1)

        # Graph convolution: first aggregate neighborhood features
        x = torch.bmm(adj_expanded, x) # (batch_size * seq_len, num_nodes,
        ↪ in_features)

        # For Conv1d: input needs to be (batch, channels, length)
        # So we permute from (batch, nodes, features) to (batch, features,
        ↪ nodes)
        x = x.permute(0, 2, 1) # -> (batch_size * seq_len, in_features,
        ↪ num_nodes)

        # Apply convolution
        x = self.conv(x) # (batch_size * seq_len, out_features, num_nodes)

        # Permute back to original format
        x = x.permute(0, 2, 1) # -> (batch_size * seq_len, num_nodes,
        ↪ out_features)

        x = F.gelu(x)
        x = self.dropout(x)
        return x

class GCNBiLSTM(nn.Module):
    def __init__(self, num_nodes=NUM_NODES, in_features=FEATURE_DIM,
                  gc_hidden=64, lstm_hidden=128, num_classes=10,
                  num_gcn_layers=2, dropout=0.5, label_map=None):
        super(GCNBiLSTM, self).__init__()

        # Create multiple GCN layers
        self.gcn_layers = nn.ModuleList()
        self.gcn_layers.append(GCNLayer(in_features, gc_hidden, dropout))

        for _ in range(num_gcn_layers - 1):
            self.gcn_layers.append(GCNLayer(gc_hidden, gc_hidden, dropout))

        # Bidirectional LSTM layer
        self.lstm = nn.LSTM(
            input_size=num_nodes * gc_hidden,
            hidden_size=lstm_hidden,
            num_layers=2,
            batch_first=True,
            bidirectional=True,
            dropout=dropout if num_gcn_layers > 1 else 0

```

```

)

# Attention mechanism
self.attention = nn.Sequential(
    nn.Linear(lstm_hidden * 2, 64),
    nn.Tanh(),
    nn.Linear(64, 1)
)

# Output classification layers
self.classifier = nn.Sequential(
    nn.Linear(lstm_hidden * 2, lstm_hidden),
    nn.ReLU(),
    nn.Dropout(dropout),
    nn.Linear(lstm_hidden, num_classes)
)

self.dropout = nn.Dropout(dropout)
self.label_map = label_map
self.num_nodes = num_nodes
self.gcn_hidden = gcn_hidden

def forward(self, x, adj):
    # x shape: (batch_size, seq_len, num_nodes * in_features)
    # Reshape to (batch_size, seq_len, num_nodes, in_features)
    batch_size, seq_len, _ = x.size()
    x = x.view(batch_size, seq_len, self.num_nodes, -1)

    # Process each time step through GCN
    gcn_outputs = []
    for t in range(seq_len):
        # Get current time step data
        curr_x = x[:, t, :, :] # (batch_size, num_nodes, in_features)

        # Process through GCN layers
        for gcn_layer in self.gcn_layers:
            curr_x = gcn_layer(curr_x, adj)
            curr_x = self.dropout(curr_x)

        # Flatten node features
        # curr_x = curr_x.view(batch_size, -1) # (batch_size, num_nodes *
↪gcn_hidden)
        curr_x = curr_x.contiguous().view(batch_size, -1)
        gcn_outputs.append(curr_x)

    # Stack outputs to (batch_size, seq_len, num_nodes * gcn_hidden)
    gcn_out = torch.stack(gcn_outputs, dim=1)

```

```

        # Process through BiLSTM
        lstm_out, _ = self.lstm(gcn_out) # (batch_size, seq_len, lstm_hidden * 2)

        # Apply attention mechanism
        attn_weights = self.attention(lstm_out).squeeze(-1) # (batch_size, seq_len)
        attn_weights = F.softmax(attn_weights, dim=1).unsqueeze(-1) # (batch_size, seq_len, 1)

        # Weighted sum of LSTM outputs
        context = torch.sum(lstm_out * attn_weights, dim=1) # (batch_size, lstm_hidden * 2)

        # Final classification
        output = self.classifier(context)

        return output

    def predict_label(self, x, adj):
        self.eval()
        with torch.no_grad():
            logits = self.forward(x, adj) # Forward pass
            pred_classes = torch.argmax(logits, dim=1) # Get the predicted class (index)

            if self.label_map is not None:
                pred_labels = [self.label_map[int(idx)] for idx in pred_classes.cpu().numpy()]
                return pred_labels
            else:
                return pred_classes

class GraphSequenceDataset(Dataset):
    def __init__(self, sequences, labels):
        self.sequences = torch.tensor(sequences, dtype=torch.float32)
        self.labels = torch.tensor(labels, dtype=torch.long)

    def __len__(self):
        return len(self.sequences)

    def __getitem__(self, idx):
        return self.sequences[idx], self.labels[idx]

```

```

def parse_frame(frame,):
    keypoints = []
    for part in ['pose', 'left_hand', 'right_hand']:
        for landmark in frame.get(part, []):
            keypoints.extend([landmark['x'], landmark['y'],
↪landmark['visibility']])
    return keypoints

def load_and_preprocess_data(data_dir, sequence_length=15):
    """
    Load and preprocess the JSON files into sequences.

    Args:
        data_dir: Directory containing the data
        sequence_length: Number of frames in each sequence

    Returns:
        sequences: array of shape (num_sequences, sequence_length, num_nodes *
↪features)
        sequence_labels: array of class labels
        label_encoder: fitted LabelEncoder
    """
    frame_data = []
    raw_labels = []

    # Step 1: Collect all labels
    for root, dirs, files in os.walk(data_dir):
        for file in files:
            if file.endswith(".json"):
                label = os.path.basename(os.path.dirname(os.path.join(root,
↪file)))
                raw_labels.append(label)

    # Step 2: Fit label encoder
    encoder = LabelEncoder()
    encoder.fit(raw_labels)
    label_map = {label: int(encoder.transform([label])[0]) for label in
↪set(raw_labels)}

    sequences = []
    sequence_labels = []

    # Step 3: Parse frames and assign encoded labels
    for root, dirs, files in os.walk(data_dir):

```

```

        for file in files:
            if file.endswith(".json"):
                label = os.path.basename(os.path.dirname(os.path.join(root,
↪file)))

                encoded_label = label_map[label]
                with open(os.path.join(root, file), 'r') as f:
                    frames = json.load(f)
                    features=[]
                    for frame in frames:
                        features.append(parse_frame(frame))
                        # frame_data.append([np.array(features), encoded_label])
                    sequences.append(np.stack(features))
                    sequence_labels.append(encoded_label)
idx_to_label = {v: k for k, v in label_map.items()}
label_map = idx_to_label
del idx_to_label
gc.collect()

return np.array(sequences), np.array(sequence_labels), label_map

def create_adjacency_matrix():
    """Create the adjacency matrix for the graph."""

    pose_connections = [
        # Mouth
        (9,10),
        # Left Eyes
        (1,2),(2,3),(3,7),
        # Right Eyes
        (4,5),(5,6),(6,8),
        # Nose
        (0,4),(0,1),
        # Shoulders
        (11, 12),
        # Connect shoulders to hip
        (11, 17), (12, 18),
        # Connect hip points
        (17, 18),
        # Left arm
        (11, 13), (13, 15),
        # Right arm
        (12, 14), (14, 16)
    ]

    hand_connections = [
        # Thumb

```

```

(0, 1), (1, 2), (2, 3), (3, 4),
# Index finger
(0, 5), (5, 6), (6, 7), (7, 8),
# Middle finger
(0, 9), (9, 10), (10, 11), (11, 12),
# Ring finger
(0, 13), (13, 14), (14, 15), (15, 16),
# Pinky
(0, 17), (17, 18), (18, 19), (19, 20),
# Palm connections
(5, 9), (9, 13), (13, 17)
]

def create_adj_matrix(num_nodes, connections):
    adj_matrix = np.zeros((num_nodes, num_nodes))
    for i, j in connections:
        adj_matrix[i, j] = 1
        adj_matrix[j, i] = 1
    # Add self-loops
    for i in range(num_nodes):
        adj_matrix[i, i] = 1
    return adj_matrix

pose_adj_matrix = create_adj_matrix(NUM_POSE_LANDMARKS, pose_connections)
left_hand_adj_matrix = create_adj_matrix(NUM_HAND_LANDMARKS,
↪hand_connections)
right_hand_adj_matrix = create_adj_matrix(NUM_HAND_LANDMARKS,
↪hand_connections)

# Calculate the total number of nodes
total_nodes = NUM_POSE_LANDMARKS + NUM_HAND_LANDMARKS + NUM_HAND_LANDMARKS

# Initialize a global adjacency matrix
global_adj_matrix = np.zeros((total_nodes, total_nodes))

# start_pose = NUM_FACE_LANDMARKS
start_pose=0
end_pose = start_pose + NUM_POSE_LANDMARKS
global_adj_matrix[start_pose:end_pose, start_pose:end_pose] =
↪pose_adj_matrix

start_lh = end_pose
end_lh = start_lh + NUM_HAND_LANDMARKS
global_adj_matrix[start_lh:end_lh, start_lh:end_lh] = left_hand_adj_matrix

start_rh = end_lh
end_rh = start_rh + NUM_HAND_LANDMARKS

```

```

global_adj_matrix[start_rh:end_rh, start_rh:end_rh] = right_hand_adj_matrix

# Connect pose to hands
pose_hand_connections = [
    (start_pose + 15, start_lh), # Left hand wrist to left hand base
    (start_pose + 16, start_rh), # Right hand wrist to right hand base
]
for i, j in pose_hand_connections:
    global_adj_matrix[i, j] = 1
    global_adj_matrix[j, i] = 1

# Normalize adjacency matrix ( $D^{-0.5} * A * D^{-0.5}$ )
# Add identity matrix to include self-connections
adj_matrix = global_adj_matrix + np.eye(total_nodes)

# Calculate degree matrix
degree_matrix = np.diag(np.sum(adj_matrix, axis=1))

#  $D^{-0.5}$ 
deg_inv_sqrt = np.linalg.inv(np.sqrt(degree_matrix))

# Normalized adjacency matrix
normalized_adj_matrix = deg_inv_sqrt @ adj_matrix @ deg_inv_sqrt

return torch.FloatTensor(normalized_adj_matrix)

def train_model(model, train_loader, val_loader, adj_matrix, num_epochs=50,
    lr=0.001,
    weight_decay=1e-5, patience=10, model_save_path='best_model.pt'):
    """
    Train the GCNBiLSTM model

    Args:
        model: GCNBiLSTM model
        train_loader: DataLoader for training data
        val_loader: DataLoader for validation data
        adj_matrix: Normalized adjacency matrix
        num_epochs: Number of training epochs
        lr: Learning rate
        weight_decay: Weight decay factor
        patience: Early stopping patience
        model_save_path: Path to save best model

    Returns:
        model: Trained model
        train_losses: List of training losses

```



```

    val_losses: List of validation losses
    train_accs: List of training accuracies
    val_accs: List of validation accuracies
    """
    device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
    print(f"Using device: {device}")

    model = model.to(device)
    adj_matrix = adj_matrix.to(device)

    criterion = nn.CrossEntropyLoss()
    optimizer = torch.optim.Adam(model.parameters(), lr=lr,
    ↪weight_decay=weight_decay)

    scheduler = torch.optim.lr_scheduler.ReduceLROnPlateau(optimizer, 'min',
    ↪patience=5, factor=0.5)

    train_losses = []
    val_losses = []
    train_accs = []
    val_accs = []

    best_val_loss = float('inf')
    early_stop_counter = 0

    for epoch in range(num_epochs):
        # Training phase
        model.train()
        train_loss = 0.0
        correct = 0
        total = 0

        progress_bar = tqdm(train_loader, desc=f"Epoch {epoch+1}/{num_epochs}",
    ↪[Train])
        for batch_idx, (inputs, targets) in enumerate(progress_bar):
            inputs, targets = inputs.to(device), targets.to(device)

            optimizer.zero_grad()
            outputs = model(inputs, adj_matrix)

            loss = criterion(outputs, targets)
            loss.backward()

            # Optional: gradient clipping
            torch.nn.utils.clip_grad_norm_(model.parameters(), max_norm=5.0)

```

```

optimizer.step()

train_loss += loss.item()
_, predicted = outputs.max(1)
total += targets.size(0)
correct += predicted.eq(targets).sum().item()

progress_bar.set_postfix({
    'loss': train_loss/(batch_idx+1),
    'acc': 100.*correct/total
})

train_loss = train_loss / len(train_loader)
train_acc = 100. * correct / total
train_losses.append(train_loss)
train_accs.append(train_acc)

# Validation phase
model.eval()
val_loss = 0.0
correct = 0
total = 0

with torch.no_grad():
    for batch_idx, (inputs, targets) in enumerate(val_loader):
        inputs, targets = inputs.to(device), targets.to(device)
        outputs = model(inputs, adj_matrix)

        loss = criterion(outputs, targets)
        val_loss += loss.item()

        _, predicted = outputs.max(1)
        total += targets.size(0)
        correct += predicted.eq(targets).sum().item()

val_loss = val_loss / len(val_loader)
val_acc = 100. * correct / total
val_losses.append(val_loss)
val_accs.append(val_acc)

# Update learning rate
scheduler.step(val_loss)

print(f"Epoch {epoch+1}/{num_epochs} - "
      f"Train Loss: {train_loss:.4f}, Train Acc: {train_acc:.2f}% - "
      f"Val Loss: {val_loss:.4f}, Val Acc: {val_acc:.2f}%")

```

```

    # Save best model
    if val_loss < best_val_loss:
        best_val_loss = val_loss
        # Save only the model state dict
        torch.save({'model_state_dict':model.state_dict(), 'label_map':model.
↪label_map}, model_save_path)

        early_stop_counter = 0
        print(f"Saved best model to {model_save_path}")
    else:
        early_stop_counter += 1

    # Early stopping
    if early_stop_counter >= patience:
        print(f"Early stopping triggered after {epoch+1} epochs")
        break

    # Load best model
    checkpoint = torch.load(model_save_path)
    model.load_state_dict(checkpoint['model_state_dict'])
    model.label_map = checkpoint['label_map']
    return model, train_losses, val_losses, train_accs, val_accs

def evaluate_model(model, test_loader, adj_matrix):
    """
    Evaluate the model on test data

    Args:
        model: Trained GCNBiLSTM model
        test_loader: DataLoader for test data
        adj_matrix: Normalized adjacency matrix

    Returns:
        test_acc: Test accuracy
        predictions: Predicted labels
        true_labels: True labels
    """
    device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
    model = model.to(device)
    adj_matrix = adj_matrix.to(device)

    model.eval()
    correct = 0
    total = 0

```

```

all_preds = []
all_targets = []

with torch.no_grad():
    for inputs, targets in test_loader:
        inputs, targets = inputs.to(device), targets.to(device)
        outputs = model(inputs, adj_matrix)

        _, predicted = outputs.max(1)
        total += targets.size(0)
        correct += predicted.eq(targets).sum().item()

        all_preds.extend(predicted.cpu().numpy())
        all_targets.extend(targets.cpu().numpy())

test_acc = 100. * correct / total
print(f"Test Accuracy: {test_acc:.2f}%")

return test_acc, np.array(all_preds), np.array(all_targets)

def plot_results(train_losses, val_losses, train_accs, val_accs):
    """
    Plot training and validation metrics

    Args:
        train_losses: List of training losses
        val_losses: List of validation losses
        train_accs: List of training accuracies
        val_accs: List of validation accuracies
    """
    plt.figure(figsize=(12, 5))

    plt.subplot(1, 2, 1)
    plt.plot(train_losses, label='Train Loss')
    plt.plot(val_losses, label='Val Loss')
    plt.xlabel('Epoch')
    plt.ylabel('Loss')
    plt.legend()
    plt.title('Training and Validation Loss')

    plt.subplot(1, 2, 2)
    plt.plot(train_accs, label='Train Accuracy')
    plt.plot(val_accs, label='Val Accuracy')
    plt.xlabel('Epoch')
    plt.ylabel('Accuracy (%)')
    plt.legend()

```

```

plt.title('Training and Validation Accuracy')

plt.tight_layout()
plt.savefig('training_history.png')
plt.show()

def main():
    torch.serialization.add_safe_globals([LabelEncoder])
    # Set random seeds for reproducibility
    SEED = 42
    torch.manual_seed(SEED)
    np.random.seed(SEED)

    # 1. Load and preprocess data
    data_dir = "data" # Update with your data directory
    test_dir = "test_data"
    sequences, sequence_labels, label_map = load_and_preprocess_data(data_dir)

    print(f"Loaded {len(sequences)} sequences with shape {sequences.shape}")
    print(f'Sequence Label {len(sequence_labels)}')
    print(f"Number of classes: {len(label_map)}")

    # 2. Create adjacency matrix

    adj_matrix = create_adjacency_matrix()
    print(f'Unique Label : {len(np.unique(sequence_labels))}')
    X_train, X_val, y_train, y_val = train_test_split(
        sequences, sequence_labels, test_size=0.4, random_state=SEED,
↪stratify=sequence_labels
    )
    print(f'Unique Train : {len(np.unique(y_train))}')
    print(f'Unique Val : {len(np.unique(y_val))}')
    X_val, X_test, y_val, y_test = train_test_split(
        X_val, y_val, test_size=0.5, random_state=SEED, stratify=y_val
    )
    print(f'Unique Train : {len(np.unique(y_train))}')
    print(f'Unique Test : {len(np.unique(y_test))}')
    # X_test, y_test, _ = load_and_preprocess_data(test_dir)

    # 4. Create datasets and dataloaders
    train_dataset = GraphSequenceDataset(X_train, y_train)
    val_dataset = GraphSequenceDataset(X_val, y_val)
    test_dataset = GraphSequenceDataset(X_test, y_test)

    train_loader = DataLoader(train_dataset, batch_size=32, shuffle=True)

```

```

val_loader = DataLoader(val_dataset, batch_size=32, shuffle=False)
test_loader = DataLoader(test_dataset, batch_size=32, shuffle=False)

# 5. Create and train the model
num_classes = len(label_map)
model = GCNBiLSTM(
    num_nodes=NUM_NODES,
    in_features=FEATURE_DIM,
    gcn_hidden=256,
    lstm_hidden=512,
    num_classes=num_classes,
    num_gcn_layers=2,
    dropout=0.3,
    label_map=label_map
)

# 6. Train the model
trained_model, train_losses, val_losses, train_accs, val_accs = train_model(
    model, train_loader, val_loader, adj_matrix,
    num_epochs=100, lr=0.001, weight_decay=5e-4,
    patience=15, model_save_path='best_gcn_bilstm_model.pt'
)

# 7. Evaluate the model
test_acc, predictions, true_labels = evaluate_model(trained_model,
↳test_loader, adj_matrix)

# 8. Plot results
plot_results(train_losses, val_losses, train_accs, val_accs)

# 9. Print classification report

print("\nClassification Report:")
actual_classes = np.unique(true_labels)
print(actual_classes)
class_names = [label_map[int(idx)] for idx in actual_classes]
print(classification_report(true_labels, predictions,
↳target_names=class_names))

# 10. Plot confusion matrix
plt.figure(figsize=(12, 10))
cm = confusion_matrix(true_labels, predictions)
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=class_names,
↳yticklabels=class_names)
plt.xlabel('Predicted')
plt.ylabel('True')

```

```
plt.title('Confusion Matrix')
plt.tight_layout()
plt.savefig('confusion_matrix.png')
plt.show()
```

```
[2]: main()
```

```
Loaded 5643 sequences with shape (5643, 3, 183)
Sequence Label 5643
Number of classes: 51
Unique Label : 51
Unique Train : 51
Unique Val : 51
Unique Train : 51
Unique Test : 51
Using device: cuda

Epoch 1/100 [Train]: 100%|      | 106/106 [00:04<00:00, 21.63it/s,
loss=3.24, acc=8.39]

Epoch 1/100 - Train Loss: 3.2447, Train Acc: 8.39% - Val Loss: 2.9506, Val Acc:
10.89%
Saved best model to best_gcn_bilstm_model.pt

Epoch 2/100 [Train]: 100%|      | 106/106 [00:04<00:00, 24.73it/s,
loss=2.94, acc=11.8]

Epoch 2/100 - Train Loss: 2.9434, Train Acc: 11.82% - Val Loss: 2.8766, Val Acc:
14.08%
Saved best model to best_gcn_bilstm_model.pt

Epoch 3/100 [Train]: 100%|      | 106/106 [00:04<00:00, 25.03it/s,
loss=2.87, acc=13.3]

Epoch 3/100 - Train Loss: 2.8734, Train Acc: 13.32% - Val Loss: 2.9031, Val Acc:
11.60%

Epoch 4/100 [Train]: 100%|      | 106/106 [00:04<00:00, 24.63it/s, loss=2.8,
acc=15.1]

Epoch 4/100 - Train Loss: 2.7979, Train Acc: 15.07% - Val Loss: 2.7328, Val Acc:
16.21%
Saved best model to best_gcn_bilstm_model.pt

Epoch 5/100 [Train]: 100%|      | 106/106 [00:04<00:00, 24.47it/s,
loss=2.61, acc=20.1]

Epoch 5/100 - Train Loss: 2.6057, Train Acc: 20.12% - Val Loss: 2.4283, Val Acc:
26.13%
Saved best model to best_gcn_bilstm_model.pt

Epoch 6/100 [Train]: 100%|      | 106/106 [00:04<00:00, 24.76it/s,
loss=2.45, acc=26]
```

Epoch 6/100 - Train Loss: 2.4539, Train Acc: 25.97% - Val Loss: 2.3292, Val Acc: 30.65%
 Saved best model to best_gcn_bilstm_model.pt

Epoch 7/100 [Train]: 100%| | 106/106 [00:04<00:00, 24.74it/s, loss=2.29, acc=30]

Epoch 7/100 - Train Loss: 2.2889, Train Acc: 29.99% - Val Loss: 2.0924, Val Acc: 34.01%
 Saved best model to best_gcn_bilstm_model.pt

Epoch 8/100 [Train]: 100%| | 106/106 [00:04<00:00, 23.79it/s, loss=2.08, acc=36.4]

Epoch 8/100 - Train Loss: 2.0797, Train Acc: 36.37% - Val Loss: 1.8916, Val Acc: 40.39%
 Saved best model to best_gcn_bilstm_model.pt

Epoch 9/100 [Train]: 100%| | 106/106 [00:04<00:00, 23.34it/s, loss=1.92, acc=39.7]

Epoch 9/100 - Train Loss: 1.9168, Train Acc: 39.68% - Val Loss: 1.7903, Val Acc: 44.46%
 Saved best model to best_gcn_bilstm_model.pt

Epoch 10/100 [Train]: 100%| | 106/106 [00:04<00:00, 23.83it/s, loss=1.75, acc=45.2]

Epoch 10/100 - Train Loss: 1.7506, Train Acc: 45.17% - Val Loss: 1.6826, Val Acc: 48.18%
 Saved best model to best_gcn_bilstm_model.pt

Epoch 11/100 [Train]: 100%| | 106/106 [00:04<00:00, 26.01it/s, loss=1.71, acc=45.5]

Epoch 11/100 - Train Loss: 1.7085, Train Acc: 45.49% - Val Loss: 1.7242, Val Acc: 44.64%

Epoch 12/100 [Train]: 100%| | 106/106 [00:04<00:00, 23.39it/s, loss=1.62, acc=48.9]

Epoch 12/100 - Train Loss: 1.6178, Train Acc: 48.89% - Val Loss: 1.5700, Val Acc: 51.02%
 Saved best model to best_gcn_bilstm_model.pt

Epoch 13/100 [Train]: 100%| | 106/106 [00:03<00:00, 27.28it/s, loss=1.57, acc=49.7]

Epoch 13/100 - Train Loss: 1.5671, Train Acc: 49.72% - Val Loss: 1.4661, Val Acc: 53.85%
 Saved best model to best_gcn_bilstm_model.pt

Epoch 14/100 [Train]: 100%| | 106/106 [00:03<00:00, 27.51it/s, loss=1.51, acc=51.5]

Epoch 14/100 - Train Loss: 1.5106, Train Acc: 51.52% - Val Loss: 1.4410, Val Acc: 55.80%
 Saved best model to best_gcn_bilstm_model.pt

Epoch 15/100 [Train]: 100%| | 106/106 [00:03<00:00, 26.63it/s, loss=1.45, acc=52.8]

Epoch 15/100 - Train Loss: 1.4452, Train Acc: 52.82% - Val Loss: 1.4643, Val Acc: 54.30%

Epoch 16/100 [Train]: 100%| | 106/106 [00:03<00:00, 27.25it/s, loss=1.38, acc=55]

Epoch 16/100 - Train Loss: 1.3826, Train Acc: 54.98% - Val Loss: 1.4156, Val Acc: 57.40%
 Saved best model to best_gcn_bilstm_model.pt

Epoch 17/100 [Train]: 100%| | 106/106 [00:03<00:00, 28.18it/s, loss=1.36, acc=56.6]

Epoch 17/100 - Train Loss: 1.3581, Train Acc: 56.60% - Val Loss: 1.3688, Val Acc: 57.66%
 Saved best model to best_gcn_bilstm_model.pt

Epoch 18/100 [Train]: 100%| | 106/106 [00:03<00:00, 27.73it/s, loss=1.31, acc=56.8]

Epoch 18/100 - Train Loss: 1.3122, Train Acc: 56.81% - Val Loss: 1.3383, Val Acc: 58.81%
 Saved best model to best_gcn_bilstm_model.pt

Epoch 19/100 [Train]: 100%| | 106/106 [00:03<00:00, 28.10it/s, loss=1.26, acc=58.8]

Epoch 19/100 - Train Loss: 1.2594, Train Acc: 58.79% - Val Loss: 1.3111, Val Acc: 60.76%
 Saved best model to best_gcn_bilstm_model.pt

Epoch 20/100 [Train]: 100%| | 106/106 [00:03<00:00, 26.97it/s, loss=1.22, acc=60]

Epoch 20/100 - Train Loss: 1.2156, Train Acc: 60.03% - Val Loss: 1.3677, Val Acc: 58.90%

Epoch 21/100 [Train]: 100%| | 106/106 [00:03<00:00, 27.15it/s, loss=1.2, acc=60.2]

Epoch 21/100 - Train Loss: 1.2008, Train Acc: 60.21% - Val Loss: 1.3543, Val Acc: 56.86%

Epoch 22/100 [Train]: 100%| | 106/106 [00:03<00:00, 27.23it/s, loss=1.17, acc=61.4]

Epoch 22/100 - Train Loss: 1.1728, Train Acc: 61.42% - Val Loss: 1.2542, Val Acc: 60.23%
 Saved best model to best_gcn_bilstm_model.pt

Epoch 23/100 [Train]: 100%| | 106/106 [00:03<00:00, 27.74it/s,
loss=1.13, acc=62.6]

Epoch 23/100 - Train Loss: 1.1329, Train Acc: 62.60% - Val Loss: 1.2982, Val
Acc: 61.65%

Epoch 24/100 [Train]: 100%| | 106/106 [00:03<00:00, 28.62it/s,
loss=1.08, acc=64.7]

Epoch 24/100 - Train Loss: 1.0817, Train Acc: 64.73% - Val Loss: 1.1734, Val
Acc: 64.48%

Saved best model to best_gcn_bilstm_model.pt

Epoch 25/100 [Train]: 100%| | 106/106 [00:03<00:00, 28.17it/s,
loss=1.08, acc=64.1]

Epoch 25/100 - Train Loss: 1.0775, Train Acc: 64.11% - Val Loss: 1.1945, Val
Acc: 64.57%

Epoch 26/100 [Train]: 100%| | 106/106 [00:03<00:00, 26.94it/s,
loss=1.06, acc=63.7]

Epoch 26/100 - Train Loss: 1.0634, Train Acc: 63.69% - Val Loss: 1.1430, Val
Acc: 63.68%

Saved best model to best_gcn_bilstm_model.pt

Epoch 27/100 [Train]: 100%| | 106/106 [00:03<00:00, 26.76it/s,
loss=1.05, acc=64.8]

Epoch 27/100 - Train Loss: 1.0518, Train Acc: 64.82% - Val Loss: 1.1720, Val
Acc: 64.48%

Epoch 28/100 [Train]: 100%| | 106/106 [00:03<00:00, 27.54it/s,
loss=1.02, acc=66.1]

Epoch 28/100 - Train Loss: 1.0159, Train Acc: 66.14% - Val Loss: 1.1750, Val
Acc: 64.04%

Epoch 29/100 [Train]: 100%| | 106/106 [00:03<00:00, 27.46it/s,
loss=0.992, acc=67.6]

Epoch 29/100 - Train Loss: 0.9924, Train Acc: 67.59% - Val Loss: 1.1648, Val
Acc: 66.52%

Epoch 30/100 [Train]: 100%| | 106/106 [00:03<00:00, 27.89it/s,
loss=0.981, acc=67.4]

Epoch 30/100 - Train Loss: 0.9810, Train Acc: 67.39% - Val Loss: 1.1160, Val
Acc: 67.32%

Saved best model to best_gcn_bilstm_model.pt

Epoch 31/100 [Train]: 100%| | 106/106 [00:03<00:00, 27.97it/s,
loss=0.965, acc=67.9]

Epoch 31/100 - Train Loss: 0.9653, Train Acc: 67.86% - Val Loss: 1.1381, Val
Acc: 65.63%

Epoch 32/100 [Train]: 100%| | 106/106 [00:03<00:00, 27.20it/s,
loss=0.95, acc=68.2]

Epoch 32/100 - Train Loss: 0.9501, Train Acc: 68.15% - Val Loss: 1.1037, Val
Acc: 66.70%

Saved best model to best_gcn_bilstm_model.pt

Epoch 33/100 [Train]: 100%| | 106/106 [00:03<00:00, 27.25it/s,
loss=0.903, acc=69.5]

Epoch 33/100 - Train Loss: 0.9027, Train Acc: 69.51% - Val Loss: 1.1151, Val
Acc: 67.05%

Epoch 34/100 [Train]: 100%| | 106/106 [00:03<00:00, 27.69it/s,
loss=0.904, acc=69.8]

Epoch 34/100 - Train Loss: 0.9043, Train Acc: 69.84% - Val Loss: 1.1096, Val
Acc: 66.78%

Epoch 35/100 [Train]: 100%| | 106/106 [00:03<00:00, 27.22it/s,
loss=0.893, acc=69.7]

Epoch 35/100 - Train Loss: 0.8932, Train Acc: 69.69% - Val Loss: 1.1229, Val
Acc: 66.25%

Epoch 36/100 [Train]: 100%| | 106/106 [00:03<00:00, 27.14it/s,
loss=0.908, acc=69.5]

Epoch 36/100 - Train Loss: 0.9080, Train Acc: 69.54% - Val Loss: 1.0642, Val
Acc: 68.64%

Saved best model to best_gcn_bilstm_model.pt

Epoch 37/100 [Train]: 100%| | 106/106 [00:03<00:00, 26.65it/s,
loss=0.851, acc=70.8]

Epoch 37/100 - Train Loss: 0.8512, Train Acc: 70.84% - Val Loss: 1.0972, Val
Acc: 68.56%

Epoch 38/100 [Train]: 100%| | 106/106 [00:03<00:00, 27.79it/s,
loss=0.85, acc=71.2]

Epoch 38/100 - Train Loss: 0.8500, Train Acc: 71.23% - Val Loss: 1.0428, Val
Acc: 70.68%

Saved best model to best_gcn_bilstm_model.pt

Epoch 39/100 [Train]: 100%| | 106/106 [00:03<00:00, 27.11it/s,
loss=0.839, acc=71]

Epoch 39/100 - Train Loss: 0.8392, Train Acc: 71.05% - Val Loss: 1.0875, Val
Acc: 67.49%

Epoch 40/100 [Train]: 100%| | 106/106 [00:03<00:00, 27.40it/s,
loss=0.854, acc=70.3]

Epoch 40/100 - Train Loss: 0.8536, Train Acc: 70.25% - Val Loss: 1.1043, Val
Acc: 66.34%

Epoch 41/100 [Train]: 100%| | 106/106 [00:03<00:00, 27.36it/s,
loss=0.839, acc=70.8]

Epoch 41/100 - Train Loss: 0.8386, Train Acc: 70.78% - Val Loss: 1.0732, Val
Acc: 69.18%

Epoch 42/100 [Train]: 100%| | 106/106 [00:04<00:00, 22.10it/s,
loss=0.839, acc=71.9]

Epoch 42/100 - Train Loss: 0.8386, Train Acc: 71.88% - Val Loss: 1.0698, Val
Acc: 68.47%

Epoch 43/100 [Train]: 100%| | 106/106 [00:04<00:00, 24.58it/s,
loss=0.808, acc=73.9]

Epoch 43/100 - Train Loss: 0.8076, Train Acc: 73.86% - Val Loss: 1.0390, Val
Acc: 69.71%

Saved best model to best_gcn_bilstm_model.pt

Epoch 44/100 [Train]: 100%| | 106/106 [00:03<00:00, 27.68it/s,
loss=0.775, acc=73.1]

Epoch 44/100 - Train Loss: 0.7751, Train Acc: 73.06% - Val Loss: 1.1167, Val
Acc: 67.58%

Epoch 45/100 [Train]: 100%| | 106/106 [00:03<00:00, 27.62it/s,
loss=0.766, acc=73.3]

Epoch 45/100 - Train Loss: 0.7655, Train Acc: 73.32% - Val Loss: 1.1265, Val
Acc: 68.29%

Epoch 46/100 [Train]: 100%| | 106/106 [00:03<00:00, 28.05it/s,
loss=0.754, acc=74.4]

Epoch 46/100 - Train Loss: 0.7539, Train Acc: 74.45% - Val Loss: 0.9989, Val
Acc: 70.86%

Saved best model to best_gcn_bilstm_model.pt

Epoch 47/100 [Train]: 100%| | 106/106 [00:04<00:00, 25.63it/s,
loss=0.74, acc=74]

Epoch 47/100 - Train Loss: 0.7404, Train Acc: 73.97% - Val Loss: 1.0373, Val
Acc: 67.85%

Epoch 48/100 [Train]: 100%| | 106/106 [00:03<00:00, 27.64it/s,
loss=0.75, acc=74.4]

Epoch 48/100 - Train Loss: 0.7501, Train Acc: 74.42% - Val Loss: 1.0045, Val
Acc: 70.06%

Epoch 49/100 [Train]: 100%| | 106/106 [00:03<00:00, 28.11it/s,
loss=0.731, acc=75.8]

Epoch 49/100 - Train Loss: 0.7307, Train Acc: 75.83% - Val Loss: 1.0088, Val
Acc: 71.12%

Epoch 50/100 [Train]: 100%| | 106/106 [00:04<00:00, 26.37it/s,
loss=0.696, acc=75.9]

Epoch 50/100 - Train Loss: 0.6961, Train Acc: 75.92% - Val Loss: 1.0327, Val
Acc: 69.53%

Epoch 51/100 [Train]: 100%| | 106/106 [00:03<00:00, 28.10it/s,
loss=0.71, acc=75.9]

Epoch 51/100 - Train Loss: 0.7103, Train Acc: 75.92% - Val Loss: 1.0503, Val
Acc: 70.50%

Epoch 52/100 [Train]: 100%| | 106/106 [00:03<00:00, 28.06it/s,
loss=0.714, acc=75]

Epoch 52/100 - Train Loss: 0.7142, Train Acc: 74.98% - Val Loss: 1.0294, Val
Acc: 71.30%

Epoch 53/100 [Train]: 100%| | 106/106 [00:03<00:00, 27.52it/s,
loss=0.576, acc=80]

Epoch 53/100 - Train Loss: 0.5762, Train Acc: 80.03% - Val Loss: 0.9211, Val
Acc: 73.78%

Saved best model to best_gcn_bilstm_model.pt

Epoch 54/100 [Train]: 100%| | 106/106 [00:03<00:00, 27.52it/s,
loss=0.535, acc=81.7]

Epoch 54/100 - Train Loss: 0.5348, Train Acc: 81.68% - Val Loss: 0.9625, Val
Acc: 73.78%

Epoch 55/100 [Train]: 100%| | 106/106 [00:04<00:00, 25.91it/s,
loss=0.528, acc=81.8]

Epoch 55/100 - Train Loss: 0.5283, Train Acc: 81.77% - Val Loss: 0.9218, Val
Acc: 73.69%

Epoch 56/100 [Train]: 100%| | 106/106 [00:03<00:00, 27.32it/s,
loss=0.526, acc=82]

Epoch 56/100 - Train Loss: 0.5255, Train Acc: 81.98% - Val Loss: 0.9268, Val
Acc: 74.49%

Epoch 57/100 [Train]: 100%| | 106/106 [00:03<00:00, 27.98it/s,
loss=0.494, acc=83.2]

Epoch 57/100 - Train Loss: 0.4940, Train Acc: 83.19% - Val Loss: 0.9834, Val
Acc: 73.96%

Epoch 58/100 [Train]: 100%| | 106/106 [00:04<00:00, 25.08it/s,
loss=0.517, acc=81.7]

Epoch 58/100 - Train Loss: 0.5171, Train Acc: 81.74% - Val Loss: 0.9688, Val
Acc: 74.14%

Epoch 59/100 [Train]: 100%| | 106/106 [00:03<00:00, 26.57it/s,
loss=0.523, acc=81.7]

Epoch 59/100 - Train Loss: 0.5233, Train Acc: 81.65% - Val Loss: 0.9537, Val Acc: 73.78%

Epoch 60/100 [Train]: 100%| | 106/106 [00:03<00:00, 26.80it/s,
loss=0.451, acc=84.1]

Epoch 60/100 - Train Loss: 0.4505, Train Acc: 84.14% - Val Loss: 0.8974, Val Acc: 75.47%

Saved best model to best_gcn_bilstm_model.pt

Epoch 61/100 [Train]: 100%| | 106/106 [00:04<00:00, 21.42it/s,
loss=0.425, acc=85.3]

Epoch 61/100 - Train Loss: 0.4250, Train Acc: 85.32% - Val Loss: 0.9072, Val Acc: 75.91%

Epoch 62/100 [Train]: 100%| | 106/106 [00:05<00:00, 21.06it/s,
loss=0.422, acc=85]

Epoch 62/100 - Train Loss: 0.4223, Train Acc: 84.96% - Val Loss: 0.9054, Val Acc: 76.26%

Epoch 63/100 [Train]: 100%| | 106/106 [00:04<00:00, 23.24it/s,
loss=0.398, acc=85.7]

Epoch 63/100 - Train Loss: 0.3984, Train Acc: 85.67% - Val Loss: 0.9141, Val Acc: 74.84%

Epoch 64/100 [Train]: 100%| | 106/106 [00:04<00:00, 23.50it/s,
loss=0.397, acc=86]

Epoch 64/100 - Train Loss: 0.3975, Train Acc: 86.00% - Val Loss: 0.9233, Val Acc: 75.29%

Epoch 65/100 [Train]: 100%| | 106/106 [00:04<00:00, 23.69it/s,
loss=0.399, acc=85.6]

Epoch 65/100 - Train Loss: 0.3994, Train Acc: 85.55% - Val Loss: 0.9193, Val Acc: 75.47%

Epoch 66/100 [Train]: 100%| | 106/106 [00:04<00:00, 23.41it/s,
loss=0.375, acc=87.6]

Epoch 66/100 - Train Loss: 0.3746, Train Acc: 87.59% - Val Loss: 0.9294, Val Acc: 75.20%

Epoch 67/100 [Train]: 100%| | 106/106 [00:04<00:00, 22.96it/s,
loss=0.365, acc=87.3]

Epoch 67/100 - Train Loss: 0.3646, Train Acc: 87.33% - Val Loss: 0.9115, Val Acc: 76.26%

Epoch 68/100 [Train]: 100%| | 106/106 [00:04<00:00, 23.35it/s,
loss=0.362, acc=87.3]

Epoch 68/100 - Train Loss: 0.3622, Train Acc: 87.27% - Val Loss: 0.8990, Val Acc: 76.26%

Epoch 69/100 [Train]: 100%| | 106/106 [00:04<00:00, 23.49it/s,
loss=0.349, acc=87.6]

Epoch 69/100 - Train Loss: 0.3493, Train Acc: 87.59% - Val Loss: 0.9075, Val
Acc: 76.17%

Epoch 70/100 [Train]: 100%| | 106/106 [00:04<00:00, 23.50it/s,
loss=0.348, acc=87.4]

Epoch 70/100 - Train Loss: 0.3484, Train Acc: 87.44% - Val Loss: 0.9114, Val
Acc: 76.09%

Epoch 71/100 [Train]: 100%| | 106/106 [00:04<00:00, 22.24it/s,
loss=0.328, acc=89.2]

Epoch 71/100 - Train Loss: 0.3277, Train Acc: 89.19% - Val Loss: 0.9436, Val
Acc: 76.09%

Epoch 72/100 [Train]: 100%| | 106/106 [00:04<00:00, 22.23it/s,
loss=0.35, acc=87.4]

Epoch 72/100 - Train Loss: 0.3502, Train Acc: 87.36% - Val Loss: 0.9318, Val
Acc: 75.20%

Epoch 73/100 [Train]: 100%| | 106/106 [00:04<00:00, 22.40it/s,
loss=0.321, acc=88.4]

Epoch 73/100 - Train Loss: 0.3212, Train Acc: 88.36% - Val Loss: 0.9162, Val
Acc: 76.17%


Epoch 74/100 [Train]: 100%| | 106/106 [00:04<00:00, 22.33it/s,
loss=0.321, acc=89.4]

Epoch 74/100 - Train Loss: 0.3209, Train Acc: 89.36% - Val Loss: 0.9140, Val
Acc: 76.88%

Epoch 75/100 [Train]: 100%| | 106/106 [00:04<00:00, 22.28it/s,
loss=0.317, acc=88.7]

Epoch 75/100 - Train Loss: 0.3173, Train Acc: 88.71% - Val Loss: 0.9203, Val
Acc: 75.73%

Early stopping triggered after 75 epochs
Test Accuracy: 77.59%



Classification Report:

```
[ 0  1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18 19 20 21 22 23
 24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40 41 42 43 44 45 46 47
 48 49 50]
```

| | precision | recall | f1-score | support |
|---|-----------|--------|----------|---------|
| A | 0.65 | 0.65 | 0.65 | 20 |
| B | 0.67 | 0.52 | 0.59 | 23 |
| C | 0.61 | 0.70 | 0.66 | 27 |
| D | 0.54 | 0.33 | 0.41 | 21 |
| E | 0.70 | 0.88 | 0.78 | 26 |
| F | 0.59 | 0.79 | 0.68 | 24 |
| G | 0.78 | 0.86 | 0.82 | 29 |
| H | 0.77 | 0.68 | 0.72 | 25 |
| I | 0.82 | 0.90 | 0.86 | 30 |
| J | 0.86 | 0.83 | 0.85 | 30 |
| K | 0.73 | 0.62 | 0.67 | 26 |
| L | 0.70 | 0.85 | 0.77 | 27 |
| M | 0.52 | 0.42 | 0.46 | 31 |
| N | 0.53 | 0.59 | 0.56 | 32 |
| O | 0.89 | 0.81 | 0.85 | 31 |
| P | 0.59 | 0.52 | 0.55 | 25 |
| Q | 0.75 | 0.58 | 0.65 | 31 |
| R | 0.73 | 0.62 | 0.67 | 26 |
| S | 0.69 | 0.80 | 0.74 | 30 |
| T | 0.62 | 0.50 | 0.55 | 26 |
| U | 0.83 | 0.86 | 0.84 | 28 |
| V | 0.93 | 1.00 | 0.96 | 26 |
| W | 0.78 | 0.78 | 0.78 | 27 |
| X | 0.59 | 0.59 | 0.59 | 22 |
| Y | 0.87 | 0.50 | 0.63 | 26 |

| | | | | |
|---------------|------|------|------|------|
| Z | 0.76 | 0.93 | 0.84 | 28 |
| baca | 0.93 | 0.81 | 0.87 | 16 |
| bantu | 0.85 | 0.79 | 0.81 | 14 |
| bapak | 1.00 | 1.00 | 1.00 | 15 |
| buangairkecil | 0.88 | 0.88 | 0.88 | 8 |
| buat | 0.93 | 0.88 | 0.90 | 16 |
| halo | 0.89 | 0.85 | 0.87 | 20 |
| ibu | 0.75 | 1.00 | 0.86 | 6 |
| kamu | 1.00 | 0.82 | 0.90 | 22 |
| maaf | 0.91 | 0.95 | 0.93 | 21 |
| makan | 0.74 | 0.82 | 0.78 | 17 |
| mau | 0.95 | 0.95 | 0.95 | 20 |
| nama | 0.88 | 0.88 | 0.88 | 25 |
| pagi | 0.91 | 0.83 | 0.87 | 24 |
| paham | 0.96 | 1.00 | 0.98 | 25 |
| sakit | 1.00 | 0.80 | 0.89 | 5 |
| sama-sama | 0.84 | 0.93 | 0.88 | 28 |
| saya | 0.71 | 0.92 | 0.80 | 13 |
| selamat | 0.86 | 0.90 | 0.88 | 21 |
| siapa | 0.80 | 0.75 | 0.77 | 16 |
| tanya | 0.95 | 0.90 | 0.92 | 20 |
| tempat | 0.89 | 1.00 | 0.94 | 8 |
| terima-kasih | 0.72 | 0.90 | 0.80 | 20 |
| terlambat | 0.89 | 1.00 | 0.94 | 17 |
| tidak | 0.78 | 0.82 | 0.80 | 17 |
| tolong | 0.85 | 0.94 | 0.89 | 18 |
| accuracy | | | 0.78 | 1129 |
| macro avg | 0.79 | 0.79 | 0.79 | 1129 |
| weighted avg | 0.78 | 0.78 | 0.77 | 1129 |

