

COMPARISON_CNN

June 21, 2025

```
[8]: import tensorflow as tf
from tensorflow.keras.utils import to_categorical
import os
from PIL import Image, UnidentifiedImageError
import shutil

# Configuration
IMG_SIZE = (28, 28)
BATCH_SIZE = 32
VALIDATION_SPLIT = 0.4
SEED = 42
ROOT_PATH = ''
DATASET_PATH = os.path.join(ROOT_PATH, "raw_data")
CORRUPT_PATH = os.path.join(ROOT_PATH, "corrupt_images")
os.makedirs(CORRUPT_PATH, exist_ok=True)

for root, dirs, files in os.walk(DATASET_PATH):
    for file in files:
        ext = os.path.splitext(file)[1].lower()
        if ext in [".jpg", ".jpeg", ".png", ".bmp", ".gif"]:
            path = os.path.join(root, file)
            try:
                with Image.open(path) as img:
                    img.verify() # Check integrity
            except (UnidentifiedImageError, OSError, IOError) as e:
                # Move the corrupt image
                print(f"Corrupt image found: {path} - moving to {CORRUPT_PATH}")
                dest_path = os.path.join(CORRUPT_PATH, os.path.relpath(path,
↳ DATASET_PATH))
                os.makedirs(os.path.dirname(dest_path), exist_ok=True)
                shutil.move(path, dest_path)

LANDMARK_DIR = os.path.join(ROOT_PATH, "data")
RAW_IMAGE_DIR = os.path.join(ROOT_PATH, "raw_data")
FILTERED_IMAGE_DIR = os.path.join(ROOT_PATH, "filtered_raw_data")
DATASET_PATH = FILTERED_IMAGE_DIR
# Supported image extensions
```

```

IMAGE_EXTENSIONS = ['.jpg', '.jpeg', '.png', '.bmp']

# Create filtered output structure
os.makedirs(FILTERED_IMAGE_DIR, exist_ok=True)

for class_name in os.listdir(LANDMARK_DIR):
    if class_name == 'debug':
        continue
    landmark_class_dir = os.path.join(LANDMARK_DIR, class_name)
    raw_class_dir = os.path.join(RAW_IMAGE_DIR, class_name)
    filtered_class_dir = os.path.join(FILTERED_IMAGE_DIR, class_name)
    os.makedirs(filtered_class_dir, exist_ok=True)

    for file in os.listdir(landmark_class_dir):
        if not file.endswith("_landmarks.json"):
            continue

        # Get base filename without "_landmarks.json"
        base_name = file.replace("_landmarks.json", "")

        # Look for corresponding image in raw directory
        for ext in IMAGE_EXTENSIONS:
            image_file = os.path.join(raw_class_dir, base_name + ext)
            if os.path.exists(image_file):
                # Copy to filtered folder
                shutil.copy(image_file, os.path.join(filtered_class_dir, os.
↳ path.basename(image_file)))
                break

# Load training dataset with validation split
train_ds = tf.keras.preprocessing.image_dataset_from_directory(
    DATASET_PATH,
    validation_split=VALIDATION_SPLIT,
    subset="training",
    seed=SEED,
    color_mode="grayscale",
    image_size=IMG_SIZE,
    batch_size=BATCH_SIZE
)
num_classes = len(train_ds.class_names)
label_map = train_ds.class_names

val_ds = tf.keras.preprocessing.image_dataset_from_directory(
    DATASET_PATH,
    validation_split=VALIDATION_SPLIT,
    subset="validation",
    seed=SEED,

```

```

        color_mode="grayscale",
        image_size=IMG_SIZE,
        batch_size=BATCH_SIZE
    )

test_ds = val_ds.shard(2,0)
val_ds = val_ds.shard(2,1)
# Normalize pixel values to [0, 1]
normalization_layer = tf.keras.layers.Rescaling(1./255)
train_ds = train_ds.map(lambda x, y: (normalization_layer(x), y))
val_ds = val_ds.map(lambda x, y: (normalization_layer(x), y))
test_ds = test_ds.map(lambda x, y: (normalization_layer(x), y))
# Cache and prefetch for performance
AUTOTUNE = tf.data.AUTOTUNE
train_ds = train_ds.cache().shuffle(1000).prefetch(buffer_size=AUTOTUNE)
val_ds = val_ds.cache().prefetch(buffer_size=AUTOTUNE)
test_ds = test_ds.cache().prefetch(buffer_size=AUTOTUNE)

```

Found 2155 files belonging to 25 classes.
Using 1293 files for training.
Using 1293 files for training.
Found 2155 files belonging to 25 classes.
Using 862 files for validation.

```

[10]: from tensorflow.keras.models import Sequential
      from tensorflow.keras.layers import Conv2D, MaxPooling2D, Dropout, BatchNormalization, Input
      from tensorflow.keras.layers import Flatten, Dense, GlobalAveragePooling2D
      from tensorflow.keras.optimizers import Adam

model = Sequential([
    Input((28, 28, 1)),
    Conv2D(16, (3, 3), activation='relu'),
    BatchNormalization(),
    MaxPooling2D(pool_size=(2, 2)),
    Dropout(0.1),

    Conv2D(32, (3, 3), activation='relu'),
    BatchNormalization(),
    MaxPooling2D(pool_size=(2, 2)),
    Dropout(0.2),

    GlobalAveragePooling2D(),
    Flatten(),

    Dense(128, activation='relu'),
    Dropout(0.2),

```

```

        Dense(num_classes, activation='softmax')
    ])

model.compile(optimizer=Adam(1e-3),
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

```

```
[11]: history = model.fit(train_ds, validation_data=val_ds, epochs=50)
```

```

Epoch 1/50
41/41          4s 30ms/step -
accuracy: 0.0792 - loss: 3.1791 - val_accuracy: 0.0553 - val_loss: 3.2039
Epoch 2/50
41/41          0s 8ms/step -
accuracy: 0.0990 - loss: 2.9955 - val_accuracy: 0.0721 - val_loss: 3.1999
Epoch 3/50
41/41          0s 8ms/step -
accuracy: 0.1329 - loss: 2.8954 - val_accuracy: 0.0697 - val_loss: 3.2144
Epoch 4/50
41/41          0s 7ms/step -
accuracy: 0.1812 - loss: 2.7469 - val_accuracy: 0.0457 - val_loss: 3.3359
Epoch 5/50
41/41          0s 8ms/step -
accuracy: 0.1987 - loss: 2.6381 - val_accuracy: 0.0577 - val_loss: 3.5013
Epoch 6/50
41/41          0s 7ms/step -
accuracy: 0.2533 - loss: 2.5510 - val_accuracy: 0.0385 - val_loss: 3.7691
Epoch 7/50
41/41          0s 7ms/step -
accuracy: 0.3091 - loss: 2.4013 - val_accuracy: 0.0457 - val_loss: 3.9707
Epoch 8/50
41/41          0s 8ms/step -
accuracy: 0.3289 - loss: 2.2726 - val_accuracy: 0.0962 - val_loss: 3.9763
Epoch 9/50
41/41          0s 8ms/step -
accuracy: 0.3697 - loss: 2.1622 - val_accuracy: 0.0793 - val_loss: 4.0191
Epoch 10/50
41/41          0s 8ms/step -
accuracy: 0.4231 - loss: 1.9860 - val_accuracy: 0.0697 - val_loss: 3.9028
Epoch 11/50
41/41          0s 8ms/step -
accuracy: 0.4787 - loss: 1.8349 - val_accuracy: 0.0601 - val_loss: 3.8484
Epoch 12/50
41/41          0s 7ms/step -
accuracy: 0.4719 - loss: 1.7966 - val_accuracy: 0.1010 - val_loss: 3.4286
Epoch 13/50
41/41          0s 8ms/step -
accuracy: 0.4923 - loss: 1.6918 - val_accuracy: 0.1202 - val_loss: 3.2569

```

Epoch 14/50
41/41 0s 8ms/step -
accuracy: 0.5286 - loss: 1.5895 - val_accuracy: 0.1538 - val_loss: 3.2706
Epoch 15/50
41/41 0s 8ms/step -
accuracy: 0.5598 - loss: 1.5081 - val_accuracy: 0.1923 - val_loss: 2.5698
Epoch 16/50
41/41 0s 7ms/step -
accuracy: 0.5595 - loss: 1.3816 - val_accuracy: 0.2404 - val_loss: 2.4932
Epoch 17/50
41/41 0s 8ms/step -
accuracy: 0.5971 - loss: 1.3521 - val_accuracy: 0.2404 - val_loss: 2.4065
Epoch 18/50
41/41 0s 8ms/step -
accuracy: 0.6317 - loss: 1.3178 - val_accuracy: 0.3245 - val_loss: 2.1631
Epoch 19/50
41/41 0s 8ms/step -
accuracy: 0.6310 - loss: 1.2264 - val_accuracy: 0.3438 - val_loss: 2.1272
Epoch 20/50
41/41 0s 7ms/step -
accuracy: 0.6476 - loss: 1.2063 - val_accuracy: 0.5264 - val_loss: 1.5468
Epoch 21/50
41/41 0s 8ms/step -
accuracy: 0.6589 - loss: 1.1364 - val_accuracy: 0.5457 - val_loss: 1.5659
Epoch 22/50
41/41 0s 8ms/step -
accuracy: 0.6765 - loss: 1.0622 - val_accuracy: 0.5673 - val_loss: 1.4684
Epoch 23/50
41/41 0s 8ms/step -
accuracy: 0.6960 - loss: 1.0273 - val_accuracy: 0.6322 - val_loss: 1.2309
Epoch 24/50
41/41 0s 8ms/step -
accuracy: 0.6858 - loss: 1.0117 - val_accuracy: 0.5553 - val_loss: 1.3881
Epoch 25/50
41/41 0s 8ms/step -
accuracy: 0.6939 - loss: 1.0023 - val_accuracy: 0.6635 - val_loss: 1.1033
Epoch 26/50
41/41 0s 8ms/step -
accuracy: 0.7328 - loss: 0.8993 - val_accuracy: 0.7212 - val_loss: 1.0240
Epoch 27/50
41/41 0s 8ms/step -
accuracy: 0.7280 - loss: 0.9257 - val_accuracy: 0.4808 - val_loss: 1.6999
Epoch 28/50
41/41 0s 7ms/step -
accuracy: 0.7392 - loss: 0.8670 - val_accuracy: 0.6562 - val_loss: 1.1325
Epoch 29/50
41/41 0s 8ms/step -
accuracy: 0.7125 - loss: 0.9107 - val_accuracy: 0.7476 - val_loss: 0.9960

Epoch 30/50
 41/41 0s 8ms/step -
 accuracy: 0.7633 - loss: 0.8272 - val_accuracy: 0.5457 - val_loss: 1.5624

Epoch 31/50
 41/41 0s 7ms/step -
 accuracy: 0.7518 - loss: 0.8066 - val_accuracy: 0.6971 - val_loss: 1.0020

Epoch 32/50
 41/41 0s 7ms/step -
 accuracy: 0.7725 - loss: 0.7367 - val_accuracy: 0.6442 - val_loss: 1.1853

Epoch 33/50
 41/41 0s 7ms/step -
 accuracy: 0.7559 - loss: 0.7928 - val_accuracy: 0.7043 - val_loss: 0.9951

Epoch 34/50
 41/41 0s 8ms/step -
 accuracy: 0.7950 - loss: 0.7116 - val_accuracy: 0.8005 - val_loss: 0.7544

Epoch 35/50
 41/41 0s 7ms/step -
 accuracy: 0.7948 - loss: 0.6844 - val_accuracy: 0.6827 - val_loss: 1.0278

Epoch 36/50
 41/41 0s 8ms/step -
 accuracy: 0.7684 - loss: 0.7328 - val_accuracy: 0.7548 - val_loss: 0.8494

Epoch 37/50
 41/41 0s 8ms/step -
 accuracy: 0.7618 - loss: 0.7166 - val_accuracy: 0.8413 - val_loss: 0.6279

Epoch 38/50
 41/41 0s 7ms/step -
 accuracy: 0.8095 - loss: 0.6516 - val_accuracy: 0.8101 - val_loss: 0.6800

Epoch 39/50
 41/41 0s 8ms/step -
 accuracy: 0.8033 - loss: 0.6262 - val_accuracy: 0.8029 - val_loss: 0.6571

Epoch 40/50
 41/41 0s 7ms/step -
 accuracy: 0.7766 - loss: 0.7060 - val_accuracy: 0.7692 - val_loss: 0.8074

Epoch 41/50
 41/41 0s 8ms/step -
 accuracy: 0.7863 - loss: 0.6773 - val_accuracy: 0.7981 - val_loss: 0.7020

Epoch 42/50
 41/41 0s 8ms/step -
 accuracy: 0.8257 - loss: 0.5982 - val_accuracy: 0.7091 - val_loss: 0.9816

Epoch 43/50
 41/41 0s 8ms/step -
 accuracy: 0.8237 - loss: 0.5878 - val_accuracy: 0.6442 - val_loss: 1.1126

Epoch 44/50
 41/41 0s 7ms/step -
 accuracy: 0.7964 - loss: 0.6137 - val_accuracy: 0.7308 - val_loss: 0.8835

Epoch 45/50
 41/41 0s 7ms/step -
 accuracy: 0.7883 - loss: 0.6493 - val_accuracy: 0.7885 - val_loss: 0.7244


```
Epoch 46/50
41/41          0s 7ms/step -
accuracy: 0.8074 - loss: 0.5850 - val_accuracy: 0.7260 - val_loss: 0.8995
Epoch 47/50
41/41          0s 7ms/step -
accuracy: 0.8247 - loss: 0.5645 - val_accuracy: 0.7620 - val_loss: 0.7885
Epoch 48/50
41/41          0s 7ms/step -
accuracy: 0.8261 - loss: 0.5454 - val_accuracy: 0.8125 - val_loss: 0.6782
Epoch 49/50
41/41          0s 7ms/step -
accuracy: 0.8410 - loss: 0.5211 - val_accuracy: 0.7163 - val_loss: 1.0502
Epoch 50/50
41/41          0s 7ms/step -
accuracy: 0.8434 - loss: 0.5401 - val_accuracy: 0.7308 - val_loss: 0.9555
```

```
[12]: test_loss, test_accuracy = model.evaluate(test_ds)
      print(f"Test Accuracy: {test_accuracy:.4f}")
      print(f"Test Loss: {test_loss:.4f}")
```

```
14/14          0s 16ms/step -
accuracy: 0.7293 - loss: 0.9152
Test Accuracy: 0.7399
Test Loss: 0.9218
```

```
[13]: import matplotlib.pyplot as plt
      from sklearn.metrics import classification_report, confusion_matrix
      import seaborn as sns
      import numpy as np
```

```
[14]: plt.figure(figsize=(12, 4))
      plt.subplot(1, 2, 1)
      plt.plot(history.history['accuracy'])
      plt.plot(history.history['val_accuracy'])
      plt.title('Model accuracy')
      plt.ylabel('Accuracy')
      plt.xlabel('Epoch')
      plt.legend(['Train', 'Validation'], loc='upper left')
      # Plot training & validation loss values
      plt.subplot(1, 2, 2)
      plt.plot(history.history['loss'])
      plt.plot(history.history['val_loss'])
      plt.title('Model loss')
      plt.ylabel('Loss')
      plt.xlabel('Epoch')
      plt.legend(['Train', 'Validation'], loc='upper left')
      plt.show()
```



```
[15]: y_true, y_pred = [], []
target_names = [label_map[i] for i in range(len(label_map))]
for X_batch, y_batch in test_ds:
    y_true.append(y_batch.numpy())

    batch_pred = model.predict(X_batch, verbose=0)
    y_pred.append(np.argmax(batch_pred, axis=1))

y_true = np.concatenate(y_true)
y_pred = np.concatenate(y_pred)

print(classification_report(
    y_true, y_pred,
    digits=3,
    target_names=target_names
))

cm = confusion_matrix(y_true, y_pred, labels=range(len(label_map)))
labels = [label_map[i] for i in range(len(label_map))]

plt.figure(figsize=(10, 8))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues",
            xticklabels=labels, yticklabels=labels)
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.title("Confusion Matrix - Test Set")
plt.show()
```

c:\Users\chris\.conda\envs\ASLR\Lib\site-packages\sklearn\metrics_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.


```
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
c:\Users\chris\.conda\envs\ASLR\Lib\site-
packages\sklearn\metrics\_classification.py:1565: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
c:\Users\chris\.conda\envs\ASLR\Lib\site-
packages\sklearn\metrics\_classification.py:1565: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

	precision	recall	f1-score	support
baca	0.333	1.000	0.500	15
bantu	0.571	0.444	0.500	9
bapak	1.000	0.600	0.750	15
buangairkecil	0.875	1.000	0.933	7
buat	0.533	0.800	0.640	10
halo	0.857	0.800	0.828	15
ibu	1.000	1.000	1.000	5
kamu	0.786	0.524	0.629	21
maaf	0.676	1.000	0.806	25
makan	0.625	0.938	0.750	16
mau	0.760	0.905	0.826	21
nama	1.000	0.290	0.450	31
pagi	0.839	0.963	0.897	27
paham	0.738	0.969	0.838	32
sakit	0.000	0.000	0.000	7
sama-sama	0.786	0.423	0.550	26
saya	0.579	0.688	0.629	16
selamat	1.000	0.444	0.615	18
siapa	0.857	0.600	0.706	20
tanya	0.818	0.947	0.878	19
tempat	0.750	1.000	0.857	9
terima-kasih	1.000	0.476	0.645	21
terlambat	0.929	0.765	0.839	17
tidak	0.778	0.955	0.857	22
tolong	1.000	0.955	0.977	22
accuracy			0.740	446
macro avg	0.764	0.739	0.716	446
weighted avg	0.794	0.740	0.725	446

