

COMPARISON_MediaPipe+CNN+LSTM

June 21, 2025

```
[1]: from modules.SignLanguageProcessor import load_and_preprocess_data, parse_frame
import os
```

```
[2]: ROOT_PATH = ''
sequences, labels, label_map = load_and_preprocess_data(os.path.
    ↪join(ROOT_PATH, 'data'))
```

```
[3]: num_classes = len(label_map)
```

```
[4]: len(labels)
```

```
[4]: 2155
```

```
[5]: sequences.shape
```

```
[5]: (2155, 3, 61, 3)
```

```
[6]: from sklearn.model_selection import train_test_split

X_train, X_temp, y_train, y_temp = train_test_split(
    sequences, labels, test_size=0.4, stratify=labels, random_state=42
)

X_val, X_test, y_val, y_test = train_test_split(
    X_temp, y_temp, test_size=0.5, stratify=y_temp, random_state=42
)
```

```
[7]: import numpy as np
def normalize_landmark_data(X):
    """
    Normalize the landmark features (x, y) to have zero mean and unit variance_
    ↪across the training set.
    Assumes X shape is (N, F, L, T), where F=3 (x, y, vis).
    """
    X = X.copy()
    # Flatten across all samples, landmarks, and frames
    x_vals = X[:, 0, :, :].flatten()
    y_vals = X[:, 1, :, :].flatten()
```

```

    # Compute mean and std
    x_mean, x_std = np.mean(x_vals), np.std(x_vals)
    y_mean, y_std = np.mean(y_vals), np.std(y_vals)

    # Normalize
    X[:, 0, :, :] = (X[:, 0, :, :] - x_mean) / x_std
    X[:, 1, :, :] = (X[:, 1, :, :] - y_mean) / y_std

    return X, (x_mean, x_std), (y_mean, y_std)

def apply_normalization(X, x_mean, x_std, y_mean, y_std):
    X = X.copy()
    X[:, 0, :, :] = (X[:, 0, :, :] - x_mean) / x_std
    X[:, 1, :, :] = (X[:, 1, :, :] - y_mean) / y_std
    return X

```

```

[8]: def reshape_frames_for_cnn(X, y):
      X = X.transpose(0, 3, 2, 1) # (N, T, L, F)
      X = X[..., np.newaxis]      # (N, T, L, F, 1)
      return X,y

```

```

[9]: X_train_norm, (x_mean, x_std), (y_mean, y_std) = ↵
      ↪normalize_landmark_data(X_train)
      X_val_norm = apply_normalization(X_val, x_mean, x_std, y_mean, y_std)
      X_test_norm = apply_normalization(X_test, x_mean, x_std, y_mean, y_std)

      X_train_cnn, y_train_cnn = reshape_frames_for_cnn(X_train_norm, y_train)
      X_val_cnn, y_val_cnn = reshape_frames_for_cnn(X_val_norm, y_val)
      X_test_cnn, y_test_cnn = reshape_frames_for_cnn(X_test_norm, y_test)

      print(X_train_cnn.shape)
      print(y_train_cnn.shape)

```

```

(1293, 3, 61, 3, 1)
(1293,)

```

```

[10]: input_shape = X_train_cnn.shape[1:]
      print(input_shape)

```

```

(3, 61, 3, 1)

```

```

[11]: import tensorflow as tf

      train_ds = tf.data.Dataset.from_tensor_slices((X_train_cnn, y_train_cnn))
      train_ds = train_ds.shuffle(buffer_size=1000).batch(64).prefetch(tf.data.
      ↪AUTOTUNE)

```

```

val_ds = tf.data.Dataset.from_tensor_slices((X_val_cnn, y_val_cnn))
val_ds = val_ds.batch(64).prefetch(tf.data.AUTOTUNE)

test_ds = tf.data.Dataset.from_tensor_slices((X_test_cnn, y_test_cnn))
test_ds = test_ds.batch(64).prefetch(tf.data.AUTOTUNE)

```

```

[12]: from tensorflow.keras.models import Sequential
      from tensorflow.keras.layers import TimeDistributed, Conv2D, MaxPooling2D, Flatten, Input
      from tensorflow.keras.layers import LSTM, Dropout, Dense, BatchNormalization

model = Sequential([
    Input((3, 61, 3, 1)),
    TimeDistributed(Conv2D(32, (3, 2), activation='relu', padding='same')),
    TimeDistributed(BatchNormalization()),
    TimeDistributed(MaxPooling2D(pool_size=(2, 1))),
    TimeDistributed(Dropout(0.25)),

    TimeDistributed(Conv2D(64, (3, 2), activation='relu', padding='same')),
    TimeDistributed(BatchNormalization()),
    TimeDistributed(MaxPooling2D(pool_size=(2, 1))),
    TimeDistributed(Flatten()),

    LSTM(128, return_sequences=False),
    Dropout(0.5),
    Dense(num_classes, activation='softmax')
])

model.compile(optimizer='adam', loss='sparse_categorical_crossentropy',
             metrics=['accuracy'])

```

```

[13]: history = model.fit(train_ds, validation_data=val_ds, epochs=50, batch_size=64)

```

```

Epoch 1/50
21/21          6s 68ms/step -
accuracy: 0.1066 - loss: 3.1065 - val_accuracy: 0.0464 - val_loss: 3.1254
Epoch 2/50
21/21          1s 46ms/step -
accuracy: 0.2142 - loss: 2.6277 - val_accuracy: 0.0464 - val_loss: 3.1366
Epoch 3/50
21/21          1s 43ms/step -
accuracy: 0.2999 - loss: 2.3772 - val_accuracy: 0.1230 - val_loss: 3.0154
Epoch 4/50
21/21          1s 44ms/step -
accuracy: 0.3226 - loss: 2.2264 - val_accuracy: 0.2158 - val_loss: 2.8296
Epoch 5/50
21/21          1s 43ms/step -
accuracy: 0.3770 - loss: 2.0676 - val_accuracy: 0.2019 - val_loss: 2.7027

```

Epoch 6/50
 21/21 1s 43ms/step -
 accuracy: 0.4285 - loss: 1.9289 - val_accuracy: 0.2854 - val_loss: 2.4966
 Epoch 7/50
 21/21 1s 55ms/step -
 accuracy: 0.4817 - loss: 1.8106 - val_accuracy: 0.3712 - val_loss: 2.3343
 Epoch 8/50
 21/21 1s 44ms/step -
 accuracy: 0.4966 - loss: 1.7196 - val_accuracy: 0.4362 - val_loss: 2.1730
 Epoch 9/50
 21/21 1s 44ms/step -
 accuracy: 0.5453 - loss: 1.6237 - val_accuracy: 0.4687 - val_loss: 2.0802
 Epoch 10/50
 21/21 1s 44ms/step -
 accuracy: 0.5712 - loss: 1.5956 - val_accuracy: 0.4872 - val_loss: 1.9338
 Epoch 11/50
 21/21 1s 43ms/step -
 accuracy: 0.5930 - loss: 1.4754 - val_accuracy: 0.4942 - val_loss: 1.8980
 Epoch 12/50
 21/21 1s 61ms/step -
 accuracy: 0.6227 - loss: 1.4098 - val_accuracy: 0.5220 - val_loss: 1.7979
 Epoch 13/50
 21/21 1s 49ms/step -
 accuracy: 0.6485 - loss: 1.3376 - val_accuracy: 0.5244 - val_loss: 1.7424
 Epoch 14/50
 21/21 1s 48ms/step -
 accuracy: 0.6548 - loss: 1.2770 - val_accuracy: 0.5383 - val_loss: 1.6904
 Epoch 15/50
 21/21 1s 43ms/step -
 accuracy: 0.7103 - loss: 1.1797 - val_accuracy: 0.5940 - val_loss: 1.6094
 Epoch 16/50
 21/21 1s 42ms/step -
 accuracy: 0.7115 - loss: 1.1876 - val_accuracy: 0.5824 - val_loss: 1.5863
 Epoch 17/50
 21/21 1s 57ms/step -
 accuracy: 0.6915 - loss: 1.1704 - val_accuracy: 0.6589 - val_loss: 1.4618
 Epoch 18/50
 21/21 1s 54ms/step -
 accuracy: 0.6988 - loss: 1.1164 - val_accuracy: 0.6450 - val_loss: 1.4769
 Epoch 19/50
 21/21 1s 48ms/step -
 accuracy: 0.7325 - loss: 1.0445 - val_accuracy: 0.6659 - val_loss: 1.3893
 Epoch 20/50
 21/21 1s 47ms/step -
 accuracy: 0.7659 - loss: 0.9465 - val_accuracy: 0.6705 - val_loss: 1.3289
 Epoch 21/50
 21/21 1s 51ms/step -
 accuracy: 0.7902 - loss: 0.9112 - val_accuracy: 0.6891 - val_loss: 1.3146

Epoch 22/50
 21/21 1s 65ms/step -
 accuracy: 0.7614 - loss: 0.9063 - val_accuracy: 0.7007 - val_loss: 1.2697

Epoch 23/50
 21/21 1s 49ms/step -
 accuracy: 0.7620 - loss: 0.9016 - val_accuracy: 0.7100 - val_loss: 1.2721

Epoch 24/50
 21/21 1s 47ms/step -
 accuracy: 0.7773 - loss: 0.8694 - val_accuracy: 0.7030 - val_loss: 1.2506

Epoch 25/50
 21/21 1s 44ms/step -
 accuracy: 0.7913 - loss: 0.8297 - val_accuracy: 0.7123 - val_loss: 1.1681

Epoch 26/50
 21/21 1s 46ms/step -
 accuracy: 0.8036 - loss: 0.7996 - val_accuracy: 0.7169 - val_loss: 1.1608

Epoch 27/50
 21/21 1s 45ms/step -
 accuracy: 0.8100 - loss: 0.7757 - val_accuracy: 0.7216 - val_loss: 1.1485

Epoch 28/50
 21/21 1s 45ms/step -
 accuracy: 0.8071 - loss: 0.7417 - val_accuracy: 0.7471 - val_loss: 1.0838

Epoch 29/50
 21/21 1s 44ms/step -
 accuracy: 0.8169 - loss: 0.7273 - val_accuracy: 0.7262 - val_loss: 1.1255

Epoch 30/50
 21/21 1s 43ms/step -
 accuracy: 0.8228 - loss: 0.7218 - val_accuracy: 0.7355 - val_loss: 1.1367

Epoch 31/50
 21/21 1s 43ms/step -
 accuracy: 0.8271 - loss: 0.6905 - val_accuracy: 0.7309 - val_loss: 1.0870

Epoch 32/50
 21/21 1s 43ms/step -
 accuracy: 0.8255 - loss: 0.6749 - val_accuracy: 0.7494 - val_loss: 1.0880

Epoch 33/50
 21/21 1s 44ms/step -
 accuracy: 0.8142 - loss: 0.7123 - val_accuracy: 0.7633 - val_loss: 1.0228

Epoch 34/50
 21/21 1s 44ms/step -
 accuracy: 0.8533 - loss: 0.6150 - val_accuracy: 0.7749 - val_loss: 0.9546

Epoch 35/50
 21/21 1s 43ms/step -
 accuracy: 0.8464 - loss: 0.6444 - val_accuracy: 0.7309 - val_loss: 1.0202

Epoch 36/50
 21/21 1s 45ms/step -
 accuracy: 0.8526 - loss: 0.5876 - val_accuracy: 0.7378 - val_loss: 1.0135

Epoch 37/50
 21/21 1s 43ms/step -
 accuracy: 0.8512 - loss: 0.6348 - val_accuracy: 0.7773 - val_loss: 0.9735

```

Epoch 38/50
21/21          1s 45ms/step -
accuracy: 0.8453 - loss: 0.5934 - val_accuracy: 0.7749 - val_loss: 0.9703
Epoch 39/50
21/21          1s 44ms/step -
accuracy: 0.8602 - loss: 0.5667 - val_accuracy: 0.7865 - val_loss: 0.9553
Epoch 40/50
21/21          1s 45ms/step -
accuracy: 0.8859 - loss: 0.5046 - val_accuracy: 0.7749 - val_loss: 0.9311
Epoch 41/50
21/21          1s 55ms/step -
accuracy: 0.8709 - loss: 0.5081 - val_accuracy: 0.7471 - val_loss: 0.9622
Epoch 42/50
21/21          1s 57ms/step -
accuracy: 0.8831 - loss: 0.4775 - val_accuracy: 0.7796 - val_loss: 0.9074
Epoch 43/50
21/21          1s 54ms/step -
accuracy: 0.8782 - loss: 0.4810 - val_accuracy: 0.7796 - val_loss: 0.9248
Epoch 44/50
21/21          1s 49ms/step -
accuracy: 0.8776 - loss: 0.4814 - val_accuracy: 0.7633 - val_loss: 0.8813
Epoch 45/50
21/21          1s 57ms/step -
accuracy: 0.8833 - loss: 0.4590 - val_accuracy: 0.7517 - val_loss: 0.9166
Epoch 46/50
21/21          1s 46ms/step -
accuracy: 0.8846 - loss: 0.4570 - val_accuracy: 0.7749 - val_loss: 0.9062
Epoch 47/50
21/21          1s 44ms/step -
accuracy: 0.8941 - loss: 0.4429 - val_accuracy: 0.7587 - val_loss: 0.9294
Epoch 48/50
21/21          1s 48ms/step -
accuracy: 0.8916 - loss: 0.4564 - val_accuracy: 0.7425 - val_loss: 0.9306
Epoch 49/50
21/21          1s 51ms/step -
accuracy: 0.8940 - loss: 0.4083 - val_accuracy: 0.7773 - val_loss: 0.8769
Epoch 50/50
21/21          1s 45ms/step -
accuracy: 0.8845 - loss: 0.4382 - val_accuracy: 0.7541 - val_loss: 0.9048

```

```

[14]: test_loss, test_accuracy = model.evaluate(test_ds)
      print(f"Test Accuracy: {test_accuracy:.4f}")
      print(f"Test Loss: {test_loss:.4f}")

```


```

7/7          0s 9ms/step -
accuracy: 0.7872 - loss: 0.8720
Test Accuracy: 0.7981
Test Loss: 0.8176

```

```
[15]: import matplotlib.pyplot as plt
from sklearn.metrics import classification_report, confusion_matrix
import seaborn as sns
```

```
[16]: plt.figure(figsize=(12, 4))
plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')
# Plot training & validation loss values
plt.subplot(1, 2, 2)
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()
```



```
[17]: y_true, y_pred = [], []
target_names = [label_map[i] for i in range(len(label_map))]
for X_batch, y_batch in test_ds:
    y_true.append(y_batch.numpy())

    batch_pred = model.predict(X_batch, verbose=0)
    y_pred.append(np.argmax(batch_pred, axis=1))

y_true = np.concatenate(y_true)
```

```

y_pred = np.concatenate(y_pred)

print(classification_report(
    y_true, y_pred,
    digits=3,
    target_names=target_names
))

cm = confusion_matrix(y_true, y_pred, labels=range(len(label_map)))
labels = [label_map[i] for i in range(len(label_map))]

plt.figure(figsize=(10, 8))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues",
            xticklabels=labels, yticklabels=labels)
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.title("Confusion Matrix - Test Set")
plt.show()

```

	precision	recall	f1-score	support
baca	0.933	0.875	0.903	16
bantu	0.800	0.923	0.857	13
bapak	0.778	0.467	0.583	15
buangairkecil	1.000	0.750	0.857	8
buat	0.944	1.000	0.971	17
halo	1.000	0.850	0.919	20
ibu	1.000	0.333	0.500	6
kamu	0.708	0.773	0.739	22
maaf	1.000	0.857	0.923	21
makan	0.923	0.706	0.800	17
mau	0.913	1.000	0.955	21
nama	0.700	0.840	0.764	25
pagi	0.909	0.833	0.870	24
paham	0.889	0.960	0.923	25
sakit	1.000	0.750	0.857	4
sama-sama	1.000	0.857	0.923	28
saya	0.533	0.615	0.571	13
selamat	0.833	0.714	0.769	21
siapa	0.325	0.812	0.464	16
tanya	0.778	0.700	0.737	20
tempat	0.500	0.375	0.429	8
terima-kasih	0.696	0.842	0.762	19
terlambat	0.722	0.765	0.743	17
tidak	1.000	0.647	0.786	17
tolong	1.000	0.889	0.941	18
accuracy			0.798	431

macro avg	0.835	0.765	0.782	431
weighted avg	0.839	0.798	0.806	431

