

# COMPARISON\_MediaPipe+CNN+LSTM

June 21, 2025

```
[1]: from modules.SignLanguageProcessor import load_and_preprocess_data, parse_frame
import os
```

```
[2]: ROOT_PATH = ''
sequences, labels, label_map = load_and_preprocess_data(os.path.
    ↪join(ROOT_PATH, 'data'))
```

```
[3]: num_classes = len(label_map)
```

```
[4]: len(labels)
```

```
[4]: 1722
```

```
[5]: sequences.shape
```

```
[5]: (1722, 3, 61, 3)
```

```
[6]: from sklearn.model_selection import train_test_split

X_train, X_temp, y_train, y_temp = train_test_split(
    sequences, labels, test_size=0.4, stratify=labels, random_state=42
)

X_val, X_test, y_val, y_test = train_test_split(
    X_temp, y_temp, test_size=0.5, stratify=y_temp, random_state=42
)
```

```
[7]: import numpy as np
def normalize_landmark_data(X):
    """
    Normalize the landmark features (x, y) to have zero mean and unit variance_
    ↪across the training set.
    Assumes X shape is (N, F, L, T), where F=3 (x, y, vis).
    """
    X = X.copy()
    # Flatten across all samples, landmarks, and frames
    x_vals = X[:, 0, :, :].flatten()
    y_vals = X[:, 1, :, :].flatten()
```

```

    # Compute mean and std
    x_mean, x_std = np.mean(x_vals), np.std(x_vals)
    y_mean, y_std = np.mean(y_vals), np.std(y_vals)

    # Normalize
    X[:, 0, :, :] = (X[:, 0, :, :] - x_mean) / x_std
    X[:, 1, :, :] = (X[:, 1, :, :] - y_mean) / y_std

    return X, (x_mean, x_std), (y_mean, y_std)

def apply_normalization(X, x_mean, x_std, y_mean, y_std):
    X = X.copy()
    X[:, 0, :, :] = (X[:, 0, :, :] - x_mean) / x_std
    X[:, 1, :, :] = (X[:, 1, :, :] - y_mean) / y_std
    return X

```

```

[8]: def reshape_frames_for_cnn(X, y):
      X = X.transpose(0, 3, 2, 1) # (N, T, L, F)
      X = X[..., np.newaxis]      # (N, T, L, F, 1)
      return X,y

```

```

[9]: X_train_norm, (x_mean, x_std), (y_mean, y_std) = ↵
      ↪normalize_landmark_data(X_train)
      X_val_norm = apply_normalization(X_val, x_mean, x_std, y_mean, y_std)
      X_test_norm = apply_normalization(X_test, x_mean, x_std, y_mean, y_std)

      X_train_cnn, y_train_cnn = reshape_frames_for_cnn(X_train_norm, y_train)
      X_val_cnn, y_val_cnn = reshape_frames_for_cnn(X_val_norm, y_val)
      X_test_cnn, y_test_cnn = reshape_frames_for_cnn(X_test_norm, y_test)

      print(X_train_cnn.shape)
      print(y_train_cnn.shape)

```

```

(1033, 3, 61, 3, 1)
(1033,)

```

```

[10]: input_shape = X_train_cnn.shape[1:]
      print(input_shape)

```

```

(3, 61, 3, 1)

```

```

[11]: import tensorflow as tf

      train_ds = tf.data.Dataset.from_tensor_slices((X_train_cnn, y_train_cnn))
      train_ds = train_ds.shuffle(buffer_size=1000).batch(64).prefetch(tf.data.
      ↪AUTOTUNE)

```

```

val_ds = tf.data.Dataset.from_tensor_slices((X_val_cnn, y_val_cnn))
val_ds = val_ds.batch(64).prefetch(tf.data.AUTOTUNE)

test_ds = tf.data.Dataset.from_tensor_slices((X_test_cnn, y_test_cnn))
test_ds = test_ds.batch(64).prefetch(tf.data.AUTOTUNE)

```

```

[12]: from tensorflow.keras.models import Sequential
      from tensorflow.keras.layers import TimeDistributed, Conv2D, MaxPooling2D, Flatten, Input
      from tensorflow.keras.layers import LSTM, Dropout, Dense, BatchNormalization

model = Sequential([
    Input((3, 61, 3, 1)),
    TimeDistributed(Conv2D(32, (3, 2), activation='relu', padding='same')),
    TimeDistributed(BatchNormalization()),
    TimeDistributed(MaxPooling2D(pool_size=(2, 1))),
    TimeDistributed(Dropout(0.25)),

    TimeDistributed(Conv2D(64, (3, 2), activation='relu', padding='same')),
    TimeDistributed(BatchNormalization()),
    TimeDistributed(MaxPooling2D(pool_size=(2, 1))),
    TimeDistributed(Flatten()),

    LSTM(128, return_sequences=False),
    Dropout(0.5),
    Dense(num_classes, activation='softmax')
])

model.compile(optimizer='adam', loss='sparse_categorical_crossentropy',
             metrics=['accuracy'])

```

```

[13]: history = model.fit(train_ds, validation_data=val_ds, epochs=50, batch_size=64)

```

```

Epoch 1/50
17/17          6s 91ms/step -
accuracy: 0.0948 - loss: 3.1329 - val_accuracy: 0.1192 - val_loss: 3.1134
Epoch 2/50
17/17          1s 58ms/step -
accuracy: 0.2185 - loss: 2.6554 - val_accuracy: 0.0872 - val_loss: 3.0680
Epoch 3/50
17/17          1s 55ms/step -
accuracy: 0.3175 - loss: 2.3707 - val_accuracy: 0.1163 - val_loss: 3.0046
Epoch 4/50
17/17          1s 57ms/step -
accuracy: 0.3634 - loss: 2.2721 - val_accuracy: 0.1453 - val_loss: 2.9339
Epoch 5/50
17/17          1s 54ms/step -
accuracy: 0.4012 - loss: 2.0538 - val_accuracy: 0.2645 - val_loss: 2.7263

```

Epoch 6/50  
17/17 1s 52ms/step -  
accuracy: 0.4447 - loss: 1.9395 - val\_accuracy: 0.2616 - val\_loss: 2.5388

Epoch 7/50  
17/17 1s 47ms/step -  
accuracy: 0.4373 - loss: 1.9222 - val\_accuracy: 0.3023 - val\_loss: 2.4330

Epoch 8/50  
17/17 1s 48ms/step -  
accuracy: 0.5558 - loss: 1.7145 - val\_accuracy: 0.4273 - val\_loss: 2.1990

Epoch 9/50  
17/17 1s 47ms/step -  
accuracy: 0.5358 - loss: 1.6715 - val\_accuracy: 0.4738 - val\_loss: 2.0136

Epoch 10/50  
17/17 1s 48ms/step -  
accuracy: 0.5771 - loss: 1.5514 - val\_accuracy: 0.5349 - val\_loss: 1.8477

Epoch 11/50  
17/17 1s 48ms/step -  
accuracy: 0.6159 - loss: 1.4467 - val\_accuracy: 0.5552 - val\_loss: 1.6849

Epoch 12/50  
17/17 1s 48ms/step -  
accuracy: 0.6478 - loss: 1.3760 - val\_accuracy: 0.5727 - val\_loss: 1.6199

Epoch 13/50  
17/17 1s 48ms/step -  
accuracy: 0.6681 - loss: 1.2819 - val\_accuracy: 0.5494 - val\_loss: 1.5836

Epoch 14/50  
17/17 1s 49ms/step -  
accuracy: 0.6684 - loss: 1.2353 - val\_accuracy: 0.5959 - val\_loss: 1.4854

Epoch 15/50  
17/17 1s 49ms/step -  
accuracy: 0.6957 - loss: 1.1325 - val\_accuracy: 0.6366 - val\_loss: 1.3885

Epoch 16/50  
17/17 1s 48ms/step -  
accuracy: 0.6981 - loss: 1.1730 - val\_accuracy: 0.6715 - val\_loss: 1.3355

Epoch 17/50  
17/17 1s 49ms/step -  
accuracy: 0.7269 - loss: 1.0604 - val\_accuracy: 0.6512 - val\_loss: 1.3351

Epoch 18/50  
17/17 1s 49ms/step -  
accuracy: 0.7616 - loss: 0.9893 - val\_accuracy: 0.6744 - val\_loss: 1.2629

Epoch 19/50  
17/17 1s 49ms/step -  
accuracy: 0.7735 - loss: 0.9548 - val\_accuracy: 0.6686 - val\_loss: 1.2639

Epoch 20/50  
17/17 1s 48ms/step -  
accuracy: 0.7561 - loss: 0.9183 - val\_accuracy: 0.6686 - val\_loss: 1.2440

Epoch 21/50  
17/17 1s 48ms/step -  
accuracy: 0.7898 - loss: 0.8567 - val\_accuracy: 0.7006 - val\_loss: 1.1422

Epoch 22/50  
17/17 1s 51ms/step -  
accuracy: 0.8096 - loss: 0.7984 - val\_accuracy: 0.7413 - val\_loss: 1.0513

Epoch 23/50  
17/17 1s 48ms/step -  
accuracy: 0.7992 - loss: 0.7533 - val\_accuracy: 0.7326 - val\_loss: 1.0464

Epoch 24/50  
17/17 1s 49ms/step -  
accuracy: 0.8064 - loss: 0.7515 - val\_accuracy: 0.7500 - val\_loss: 0.9901

Epoch 25/50  
17/17 1s 49ms/step -  
accuracy: 0.8258 - loss: 0.6996 - val\_accuracy: 0.7064 - val\_loss: 1.0419

Epoch 26/50  
17/17 1s 48ms/step -  
accuracy: 0.8576 - loss: 0.6416 - val\_accuracy: 0.7587 - val\_loss: 0.9356

Epoch 27/50  
17/17 1s 48ms/step -  
accuracy: 0.8602 - loss: 0.6108 - val\_accuracy: 0.7733 - val\_loss: 0.9348

Epoch 28/50  
17/17 1s 52ms/step -  
accuracy: 0.8455 - loss: 0.6148 - val\_accuracy: 0.7413 - val\_loss: 0.9779

Epoch 29/50  
17/17 1s 49ms/step -  
accuracy: 0.8537 - loss: 0.6062 - val\_accuracy: 0.7703 - val\_loss: 0.9444

Epoch 30/50  
17/17 1s 48ms/step -  
accuracy: 0.8626 - loss: 0.5710 - val\_accuracy: 0.8023 - val\_loss: 0.8876

Epoch 31/50  
17/17 1s 49ms/step -  
accuracy: 0.8774 - loss: 0.5386 - val\_accuracy: 0.7762 - val\_loss: 0.8669

Epoch 32/50  
17/17 1s 48ms/step -  
accuracy: 0.8631 - loss: 0.5756 - val\_accuracy: 0.7733 - val\_loss: 0.8590

Epoch 33/50  
17/17 1s 48ms/step -  
accuracy: 0.8748 - loss: 0.5268 - val\_accuracy: 0.7791 - val\_loss: 0.8600

Epoch 34/50  
17/17 1s 49ms/step -  
accuracy: 0.8776 - loss: 0.4866 - val\_accuracy: 0.7762 - val\_loss: 0.8685

Epoch 35/50  
17/17 1s 52ms/step -  
accuracy: 0.9087 - loss: 0.4569 - val\_accuracy: 0.7849 - val\_loss: 0.9124

Epoch 36/50  
17/17 1s 49ms/step -  
accuracy: 0.8865 - loss: 0.4913 - val\_accuracy: 0.8110 - val\_loss: 0.8460

Epoch 37/50  
17/17 1s 49ms/step -  
accuracy: 0.8830 - loss: 0.4608 - val\_accuracy: 0.8023 - val\_loss: 0.8428

```

Epoch 38/50
17/17          1s 48ms/step -
accuracy: 0.8596 - loss: 0.5185 - val_accuracy: 0.7994 - val_loss: 0.8429
Epoch 39/50
17/17          1s 48ms/step -
accuracy: 0.8957 - loss: 0.4660 - val_accuracy: 0.8052 - val_loss: 0.8251
Epoch 40/50
17/17          1s 48ms/step -
accuracy: 0.9164 - loss: 0.3995 - val_accuracy: 0.8140 - val_loss: 0.7643
Epoch 41/50
17/17          1s 48ms/step -
accuracy: 0.8954 - loss: 0.4456 - val_accuracy: 0.7733 - val_loss: 0.7977
Epoch 42/50
17/17          1s 49ms/step -
accuracy: 0.9109 - loss: 0.4084 - val_accuracy: 0.7849 - val_loss: 0.7672
Epoch 43/50
17/17          1s 48ms/step -
accuracy: 0.9080 - loss: 0.3954 - val_accuracy: 0.8140 - val_loss: 0.7641
Epoch 44/50
17/17          1s 51ms/step -
accuracy: 0.9062 - loss: 0.3786 - val_accuracy: 0.8256 - val_loss: 0.7286
Epoch 45/50
17/17          1s 55ms/step -
accuracy: 0.8937 - loss: 0.3991 - val_accuracy: 0.7965 - val_loss: 0.7984
Epoch 46/50
17/17          1s 62ms/step -
accuracy: 0.9230 - loss: 0.3566 - val_accuracy: 0.7849 - val_loss: 0.8290
Epoch 47/50
17/17          1s 50ms/step -
accuracy: 0.9388 - loss: 0.3292 - val_accuracy: 0.8052 - val_loss: 0.7480
Epoch 48/50
17/17          1s 53ms/step -
accuracy: 0.9307 - loss: 0.3336 - val_accuracy: 0.7762 - val_loss: 0.8068
Epoch 49/50
17/17          1s 53ms/step -
accuracy: 0.9117 - loss: 0.3713 - val_accuracy: 0.7936 - val_loss: 0.7892
Epoch 50/50
17/17          1s 52ms/step -
accuracy: 0.9248 - loss: 0.3455 - val_accuracy: 0.8198 - val_loss: 0.7747

```

```

[14]: test_loss, test_accuracy = model.evaluate(test_ds)
      print(f"Test Accuracy: {test_accuracy:.4f}")
      print(f"Test Loss: {test_loss:.4f}")

```


```

6/6          0s 9ms/step -
accuracy: 0.8481 - loss: 0.6277
Test Accuracy: 0.8522
Test Loss: 0.6277

```

```
[15]: import matplotlib.pyplot as plt
from sklearn.metrics import classification_report, confusion_matrix
import seaborn as sns
```

```
[16]: plt.figure(figsize=(12, 4))
plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')
# Plot training & validation loss values
plt.subplot(1, 2, 2)
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()
```



```
[17]: y_true, y_pred = [], []
target_names = [label_map[i] for i in range(len(label_map))]
for X_batch, y_batch in test_ds:
    y_true.append(y_batch.numpy())

    batch_pred = model.predict(X_batch, verbose=0)
    y_pred.append(np.argmax(batch_pred, axis=1))

y_true = np.concatenate(y_true)
```

```

y_pred = np.concatenate(y_pred)

print(classification_report(
    y_true, y_pred,
    digits=3,
    target_names=target_names
))

cm = confusion_matrix(y_true, y_pred, labels=range(len(label_map)))
labels = [label_map[i] for i in range(len(label_map))]

plt.figure(figsize=(10, 8))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues",
            xticklabels=labels, yticklabels=labels)
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.title("Confusion Matrix - Test Set")
plt.show()

```

	precision	recall	f1-score	support
baca	1.000	0.833	0.909	12
bantu	1.000	0.727	0.842	11
bapak	0.786	0.917	0.846	12
buangairkecil	1.000	1.000	1.000	6
buat	0.812	1.000	0.897	13
halo	0.900	1.000	0.947	18
ibu	1.000	0.750	0.857	4
kamu	0.682	0.789	0.732	19
maaf	1.000	1.000	1.000	18
makan	1.000	0.714	0.833	14
mau	0.933	0.824	0.875	17
nama	0.833	0.833	0.833	18
pagi	0.947	0.900	0.923	20
paham	0.950	0.950	0.950	20
sakit	1.000	0.667	0.800	3
sama-sama	0.885	0.920	0.902	25
saya	0.600	0.500	0.545	6
selamat	0.882	0.833	0.857	18
siapa	1.000	0.750	0.857	12
tanya	0.867	0.765	0.812	17
tempat	1.000	0.250	0.400	4
terima-kasih	0.773	0.944	0.850	18
terlambat	0.722	1.000	0.839	13
tidak	0.909	0.714	0.800	14
tolong	0.500	0.769	0.606	13
accuracy			0.852	345



macro avg	0.879	0.814	0.829	345
weighted avg	0.872	0.852	0.852	345

