

COMPARISON_MobileNetV2

June 21, 2025

```
[1]: import tensorflow as tf
from tensorflow.keras.utils import to_categorical
import os
from PIL import Image, UnidentifiedImageError
import shutil

# Configuration
IMG_SIZE = (96, 96)
BATCH_SIZE = 32
VALIDATION_SPLIT = 0.4
SEED = 42
ROOT_PATH = ''
DATASET_PATH = os.path.join(ROOT_PATH, "raw_data")
CORRUPT_PATH = os.path.join(ROOT_PATH, "corrupt_images")
os.makedirs(CORRUPT_PATH, exist_ok=True)

for root, dirs, files in os.walk(DATASET_PATH):
    for file in files:
        ext = os.path.splitext(file)[1].lower()
        if ext in [".jpg", ".jpeg", ".png", ".bmp", ".gif"]:
            path = os.path.join(root, file)
            try:
                with Image.open(path) as img:
                    img.verify() # Check integrity
            except (UnidentifiedImageError, OSError, IOError) as e:
                # Move the corrupt image
                print(f"Corrupt image found: {path} - moving to {CORRUPT_PATH}")
                dest_path = os.path.join(CORRUPT_PATH, os.path.relpath(path,
↳DATASET_PATH))
                os.makedirs(os.path.dirname(dest_path), exist_ok=True)
                shutil.move(path, dest_path)

LANDMARK_DIR = os.path.join(ROOT_PATH, "data")
RAW_IMAGE_DIR = os.path.join(ROOT_PATH, "raw_data")
FILTERED_IMAGE_DIR = os.path.join(ROOT_PATH, "filtered_raw_data")
DATASET_PATH = FILTERED_IMAGE_DIR
# Supported image extensions
```

```

IMAGE_EXTENSIONS = ['.jpg', '.jpeg', '.png', '.bmp']

# Create filtered output structure
os.makedirs(FILTERED_IMAGE_DIR, exist_ok=True)

for class_name in os.listdir(LANDMARK_DIR):
    if class_name == 'debug':
        continue
    landmark_class_dir = os.path.join(LANDMARK_DIR, class_name)
    raw_class_dir = os.path.join(RAW_IMAGE_DIR, class_name)
    filtered_class_dir = os.path.join(FILTERED_IMAGE_DIR, class_name)
    os.makedirs(filtered_class_dir, exist_ok=True)

    for file in os.listdir(landmark_class_dir):
        if not file.endswith("_landmarks.json"):
            continue

        # Get base filename without "_landmarks.json"
        base_name = file.replace("_landmarks.json", "")

        # Look for corresponding image in raw directory
        for ext in IMAGE_EXTENSIONS:
            image_file = os.path.join(raw_class_dir, base_name + ext)
            if os.path.exists(image_file):
                # Copy to filtered folder
                shutil.copy(image_file, os.path.join(filtered_class_dir, os.
↳ path.basename(image_file)))
                break

# Load training dataset with validation split
train_ds = tf.keras.preprocessing.image_dataset_from_directory(
    DATASET_PATH,
    validation_split=VALIDATION_SPLIT,
    subset="training",
    seed=SEED,
    color_mode="rgb",
    image_size=IMG_SIZE,
    batch_size=BATCH_SIZE
)
num_classes = len(train_ds.class_names)
label_map = train_ds.class_names

val_ds = tf.keras.preprocessing.image_dataset_from_directory(
    DATASET_PATH,
    validation_split=VALIDATION_SPLIT,
    subset="validation",
    seed=SEED,

```

```

        color_mode="rgb",
        image_size=IMG_SIZE,
        batch_size=BATCH_SIZE
    )

    test_ds = val_ds.shard(2,0)
    val_ds = val_ds.shard(2,1)
    # Normalize pixel values to [0, 1]
    normalization_layer = tf.keras.layers.Rescaling(1./255)
    train_ds = train_ds.map(lambda x, y: (normalization_layer(x), y))
    val_ds = val_ds.map(lambda x, y: (normalization_layer(x), y))
    test_ds = test_ds.map(lambda x, y: (normalization_layer(x), y))
    # Cache and prefetch for performance
    AUTOTUNE = tf.data.AUTOTUNE
    train_ds = train_ds.cache().shuffle(1000).prefetch(buffer_size=AUTOTUNE)
    val_ds = val_ds.cache().prefetch(buffer_size=AUTOTUNE)
    test_ds = test_ds.cache().prefetch(buffer_size=AUTOTUNE)

```

Found 1722 files belonging to 25 classes.

Using 1034 files for training.

Found 1722 files belonging to 25 classes.

Using 688 files for validation.

```

[2]: from tensorflow.keras.models import Sequential
    from tensorflow.keras.layers import Conv2D, MaxPooling2D, Dropout,
        BatchNormalization
    from tensorflow.keras.layers import Flatten, Dense, GlobalAveragePooling2D
    from tensorflow.keras.optimizers import Adam
    from tensorflow.keras.applications import MobileNetV2
    base_model = MobileNetV2(input_shape=(96, 96, 3), include_top=False,
        weights='imagenet')
    base_model.trainable = False

    model = Sequential([
        base_model,
        GlobalAveragePooling2D(),
        Dropout(0.3),
        Dense(128, activation='relu'),
        Dropout(0.3),
        Dense(num_classes, activation='softmax')
    ])

    model.compile(optimizer=Adam(1e-3),
                  loss='sparse_categorical_crossentropy',
                  metrics=['accuracy'])

```

```

[3]: history = model.fit(train_ds, validation_data=val_ds, epochs=50)

```

Epoch 1/50
 33/33 9s 120ms/step -
 accuracy: 0.1183 - loss: 3.4523 - val_accuracy: 0.5030 - val_loss: 1.9964

Epoch 2/50
 33/33 2s 60ms/step -
 accuracy: 0.4331 - loss: 1.8864 - val_accuracy: 0.7351 - val_loss: 1.1025

Epoch 3/50
 33/33 2s 58ms/step -
 accuracy: 0.6330 - loss: 1.2320 - val_accuracy: 0.8065 - val_loss: 0.7908

Epoch 4/50
 33/33 2s 58ms/step -
 accuracy: 0.7745 - loss: 0.8143 - val_accuracy: 0.8542 - val_loss: 0.6272

Epoch 5/50
 33/33 2s 57ms/step -
 accuracy: 0.8013 - loss: 0.6647 - val_accuracy: 0.8482 - val_loss: 0.5455

Epoch 6/50
 33/33 2s 57ms/step -
 accuracy: 0.8354 - loss: 0.5660 - val_accuracy: 0.8810 - val_loss: 0.4308

Epoch 7/50
 33/33 2s 57ms/step -
 accuracy: 0.8531 - loss: 0.5053 - val_accuracy: 0.8958 - val_loss: 0.3935

Epoch 8/50
 33/33 2s 58ms/step -
 accuracy: 0.8897 - loss: 0.3778 - val_accuracy: 0.8988 - val_loss: 0.3880

Epoch 9/50
 33/33 2s 58ms/step -
 accuracy: 0.9136 - loss: 0.2950 - val_accuracy: 0.9077 - val_loss: 0.3186

Epoch 10/50
 33/33 2s 57ms/step -
 accuracy: 0.9073 - loss: 0.2964 - val_accuracy: 0.9345 - val_loss: 0.2995

Epoch 11/50
 33/33 2s 57ms/step -
 accuracy: 0.9016 - loss: 0.3026 - val_accuracy: 0.9077 - val_loss: 0.2882

Epoch 12/50
 33/33 2s 57ms/step -
 accuracy: 0.9309 - loss: 0.2376 - val_accuracy: 0.9256 - val_loss: 0.2701

Epoch 13/50
 33/33 2s 57ms/step -
 accuracy: 0.9432 - loss: 0.1932 - val_accuracy: 0.9315 - val_loss: 0.2523

Epoch 14/50
 33/33 2s 57ms/step -
 accuracy: 0.9289 - loss: 0.2178 - val_accuracy: 0.9345 - val_loss: 0.2342

Epoch 15/50
 33/33 2s 58ms/step -
 accuracy: 0.9477 - loss: 0.1849 - val_accuracy: 0.9286 - val_loss: 0.2551

Epoch 16/50
 33/33 2s 57ms/step -
 accuracy: 0.9497 - loss: 0.1671 - val_accuracy: 0.9315 - val_loss: 0.2590

Epoch 17/50
33/33 2s 58ms/step -
accuracy: 0.9595 - loss: 0.1436 - val_accuracy: 0.9375 - val_loss: 0.2479

Epoch 18/50
33/33 2s 57ms/step -
accuracy: 0.9604 - loss: 0.1278 - val_accuracy: 0.9405 - val_loss: 0.2442

Epoch 19/50
33/33 2s 57ms/step -
accuracy: 0.9584 - loss: 0.1361 - val_accuracy: 0.9435 - val_loss: 0.2160

Epoch 20/50
33/33 2s 58ms/step -
accuracy: 0.9703 - loss: 0.1049 - val_accuracy: 0.9375 - val_loss: 0.2093

Epoch 21/50
33/33 2s 57ms/step -
accuracy: 0.9712 - loss: 0.0950 - val_accuracy: 0.9494 - val_loss: 0.2177

Epoch 22/50
33/33 2s 57ms/step -
accuracy: 0.9823 - loss: 0.0710 - val_accuracy: 0.9464 - val_loss: 0.2012

Epoch 23/50
33/33 2s 58ms/step -
accuracy: 0.9891 - loss: 0.0625 - val_accuracy: 0.9464 - val_loss: 0.2278

Epoch 24/50
33/33 2s 57ms/step -
accuracy: 0.9858 - loss: 0.0709 - val_accuracy: 0.9554 - val_loss: 0.2116

Epoch 25/50
33/33 2s 57ms/step -
accuracy: 0.9717 - loss: 0.0920 - val_accuracy: 0.9494 - val_loss: 0.2137

Epoch 26/50
33/33 2s 57ms/step -
accuracy: 0.9850 - loss: 0.0601 - val_accuracy: 0.9464 - val_loss: 0.2231

Epoch 27/50
33/33 2s 57ms/step -
accuracy: 0.9859 - loss: 0.0671 - val_accuracy: 0.9524 - val_loss: 0.2212

Epoch 28/50
33/33 2s 57ms/step -
accuracy: 0.9715 - loss: 0.0796 - val_accuracy: 0.9494 - val_loss: 0.1972

Epoch 29/50
33/33 2s 57ms/step -
accuracy: 0.9775 - loss: 0.0761 - val_accuracy: 0.9464 - val_loss: 0.2183

Epoch 30/50
33/33 2s 57ms/step -
accuracy: 0.9776 - loss: 0.0717 - val_accuracy: 0.9494 - val_loss: 0.2055

Epoch 31/50
33/33 2s 58ms/step -
accuracy: 0.9883 - loss: 0.0615 - val_accuracy: 0.9583 - val_loss: 0.1956

Epoch 32/50
33/33 2s 57ms/step -
accuracy: 0.9803 - loss: 0.0625 - val_accuracy: 0.9375 - val_loss: 0.2155

Epoch 33/50
 33/33 2s 58ms/step -
 accuracy: 0.9762 - loss: 0.0773 - val_accuracy: 0.9435 - val_loss: 0.2170

Epoch 34/50
 33/33 2s 58ms/step -
 accuracy: 0.9891 - loss: 0.0378 - val_accuracy: 0.9554 - val_loss: 0.2109

Epoch 35/50
 33/33 2s 58ms/step -
 accuracy: 0.9964 - loss: 0.0424 - val_accuracy: 0.9375 - val_loss: 0.2225

Epoch 36/50
 33/33 2s 57ms/step -
 accuracy: 0.9771 - loss: 0.0602 - val_accuracy: 0.9494 - val_loss: 0.1995

Epoch 37/50
 33/33 2s 57ms/step -
 accuracy: 0.9789 - loss: 0.0585 - val_accuracy: 0.9583 - val_loss: 0.2124

Epoch 38/50
 33/33 2s 58ms/step -
 accuracy: 0.9784 - loss: 0.0666 - val_accuracy: 0.9435 - val_loss: 0.2283

Epoch 39/50
 33/33 2s 57ms/step -
 accuracy: 0.9971 - loss: 0.0298 - val_accuracy: 0.9554 - val_loss: 0.2114

Epoch 40/50
 33/33 2s 62ms/step -
 accuracy: 0.9794 - loss: 0.0561 - val_accuracy: 0.9524 - val_loss: 0.2150

Epoch 41/50
 33/33 2s 59ms/step -
 accuracy: 0.9921 - loss: 0.0344 - val_accuracy: 0.9554 - val_loss: 0.2165

Epoch 42/50
 33/33 2s 61ms/step -
 accuracy: 0.9945 - loss: 0.0277 - val_accuracy: 0.9494 - val_loss: 0.2186

Epoch 43/50
 33/33 2s 59ms/step -
 accuracy: 0.9828 - loss: 0.0478 - val_accuracy: 0.9494 - val_loss: 0.2196

Epoch 44/50
 33/33 2s 60ms/step -
 accuracy: 0.9871 - loss: 0.0463 - val_accuracy: 0.9435 - val_loss: 0.2366

Epoch 45/50
 33/33 2s 59ms/step -
 accuracy: 0.9891 - loss: 0.0384 - val_accuracy: 0.9494 - val_loss: 0.2233

Epoch 46/50
 33/33 2s 61ms/step -
 accuracy: 0.9922 - loss: 0.0300 - val_accuracy: 0.9554 - val_loss: 0.2446

Epoch 47/50
 33/33 2s 60ms/step -
 accuracy: 0.9822 - loss: 0.0556 - val_accuracy: 0.9524 - val_loss: 0.2354

Epoch 48/50
 33/33 2s 61ms/step -
 accuracy: 0.9876 - loss: 0.0357 - val_accuracy: 0.9464 - val_loss: 0.2289


```
Epoch 49/50
33/33          2s 60ms/step -
accuracy: 0.9943 - loss: 0.0296 - val_accuracy: 0.9375 - val_loss: 0.2478
Epoch 50/50
33/33          2s 60ms/step -
accuracy: 0.9911 - loss: 0.0266 - val_accuracy: 0.9524 - val_loss: 0.2045
```

```
[4]: test_loss, test_accuracy = model.evaluate(test_ds)
      print(f"Test Accuracy: {test_accuracy:.4f}")
      print(f"Test Loss: {test_loss:.4f}")
```

```
11/11          1s 53ms/step -
accuracy: 0.9274 - loss: 0.2769
Test Accuracy: 0.9403
Test Loss: 0.2418
```

```
[5]: import matplotlib.pyplot as plt
      from sklearn.metrics import classification_report, confusion_matrix
      import seaborn as sns
      import numpy as np
```

```
[6]: plt.figure(figsize=(12, 4))
      plt.subplot(1, 2, 1)
      plt.plot(history.history['accuracy'])
      plt.plot(history.history['val_accuracy'])
      plt.title('Model accuracy')
      plt.ylabel('Accuracy')
      plt.xlabel('Epoch')
      plt.legend(['Train', 'Validation'], loc='upper left')
      # Plot training & validation loss values
      plt.subplot(1, 2, 2)
      plt.plot(history.history['loss'])
      plt.plot(history.history['val_loss'])
      plt.title('Model loss')
      plt.ylabel('Loss')
      plt.xlabel('Epoch')
      plt.legend(['Train', 'Validation'], loc='upper left')
      plt.show()
```



```
[7]: y_true, y_pred = [], []
target_names = [label_map[i] for i in range(len(label_map))]
for X_batch, y_batch in test_ds:
    y_true.append(y_batch.numpy())

    batch_pred = model.predict(X_batch, verbose=0)
    y_pred.append(np.argmax(batch_pred, axis=1))

y_true = np.concatenate(y_true)
y_pred = np.concatenate(y_pred)

print(classification_report(
    y_true, y_pred,
    digits=3,
    target_names=target_names
))

cm = confusion_matrix(y_true, y_pred, labels=range(len(label_map)))
labels = [label_map[i] for i in range(len(label_map))]

plt.figure(figsize=(10, 8))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues",
            xticklabels=labels, yticklabels=labels)
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.title("Confusion Matrix - Test Set")
plt.show()
```

	precision	recall	f1-score	support
baca	0.727	0.889	0.800	9
bantu	0.857	0.857	0.857	7

bapak	1.000	0.917	0.957	12
buangairkecil	0.833	1.000	0.909	5
buat	1.000	1.000	1.000	12
halo	0.846	1.000	0.917	11
ibu	1.000	1.000	1.000	3
kamu	0.955	0.913	0.933	23
maaf	1.000	0.944	0.971	18
makan	0.909	1.000	0.952	10
mau	1.000	0.810	0.895	21
nama	0.933	0.824	0.875	17
pagi	1.000	0.958	0.979	24
paham	1.000	1.000	1.000	21
sakit	1.000	0.667	0.800	6
sama-sama	0.862	0.893	0.877	28
saya	1.000	1.000	1.000	5
selamat	0.909	0.952	0.930	21
siapa	0.929	1.000	0.963	13
tanya	1.000	1.000	1.000	21
tempat	1.000	1.000	1.000	4
terima-kasih	0.920	0.958	0.939	24
terlambat	1.000	1.000	1.000	14
tidak	1.000	1.000	1.000	15
tolong	0.800	1.000	0.889	8
accuracy			0.940	352
macro avg	0.939	0.943	0.938	352
weighted avg	0.945	0.940	0.940	352

